

# Practice Exam 1 - Results

## Question 1 **Correct**

Which of the following statements are true regarding External Tables? Select all that apply.

**An external table supports internal stages.**

**An external table is another name for transient tables.**

**Your selection is correct**

**An external table supports only external stages.**

**Your selection is correct**

**The data for an external table is stored in cloud storage managed by the customer.**

Overall explanation

An external table is a metadata definition; that is, you register the definition of an external table, but the external table itself doesn't contain any data. Instead, the table metadata contains column definition, the name of the external stage from where the data for the external table is, and the file format which should be used to read that data. The external stage, in turn, points to object storage on the cloud, for example, an AWS bucket or Azure Blob storage, which contains the data for the external table. Note that an external table can only point to an external stage. An internal stage cannot be used to create an external table. <https://docs.snowflake.com/en/user-guide/tables-external-intro>

**Domain**

Data Loading and Unloading

## Question 2 **Correct**

True/False: Once created, micro-partitions are immutable and cannot be modified.

**Your answer is correct**

**True**

**False**

Overall explanation

Snowflake partitions are immutable, which means they cannot be changed once created. <https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions.html>

**Domain**

Micro partitions

## Question 3 **Correct**

You are required to share data from various tables in separate databases. What is the recommended approach to simplify the sharing process?

**Clone tables from the separate databases into a new database. Share the new database.**

**Copy all data from the various tables into new tables in a new database. Share the new database.**

**Create one Share per database.**

**Your answer is correct**

**Create secure views in a single database to consolidate the data from various databases into a new database. Share the new database.**

Overall explanation

You may create a secure view if you need to share data from many tables in separate databases. Because several databases cannot be added to a single share, Snowflake

suggests creating secure views within a single database and sharing that database.  
<https://docs.snowflake.com/en/user-guide/data-sharing-multiple-db>

### Domain

Data Sharing

### Question 4Correct

What permissions does the “OPERATE” task privilege provide to a role in Snowflake?

**Your answer is correct**

**A Snowflake role with OPERATE privilege can resume or suspend the tasks.**

### Explanation

With the “OPERATE” privilege for tasks, a role can suspend or resume them. This privilege allows the user to control the execution state of tasks, enabling them to suspend task execution or restart it as needed.

See the following link for more details:

<https://docs.snowflake.com/en/user-guide/tasks-intro#resuming-or-suspending-tasks>

<https://docs.snowflake.com/en/user-guide/security-access-control-privileges#:~:text=MANAGED%20ACCOUNTS.,OPERATE,-Warehouse%2C%20Task%2C%20Dynamic>

Warehouse%2C%20Task%2C%20Dynamic

**A Snowflake role with OPERATE privilege can modify and update task definitions.**

### Explanation

These actions require the MODIFY privilege or the CREATE TASK privilege, which allow altering task definitions.

**A Snowflake role with OPERATE privilege can prioritize or deprioritize task execution based on workload demands.**

### Explanation

Snowflake does not provide any option to increase or decrease the priority of a task execution.

**A Snowflake role with OPERATE privilege can view task status and execution history.**

### Explanation

Viewing and managing task status and execution history requires the MONITOR privilege.

### Domain

Tasks

### Question 5Correct

Large tables can have \_\_\_\_\_ micro-partitions.

Hundreds

Tens

**Your answer is correct**

**Millions or hundreds of millions**

Overall explanation

The number of micro-partitions for a given table depends mainly on the amount of data in that table. For a very large table, the number of micro-partitions can run into millions or hundreds of millions of micro-partitions. <https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions>

## Domain

Architecture

### Question 6Correct

Which of the following correctly describes materialized views?

Select two answers.

Materialized views need to be refreshed manually.

Your selection is correct

Materialized view refreshes are performed automatically.

Querying a materialized view is typically slower than the base table used in the materialized view.

Materialized views are created to improve the performance of all queries.

Materialized views definition can contain multiple tables and joins.

Your selection is correct

Materialized views are created to improve the performance of specific queries.

Materialized views are created to enable sharing of data.

Overall explanation

A materialized view is a view that pre-computes data based on a SELECT query. The query's results are pre-computed and physically stored to enhance performance for similar queries that are executed in the future. When the underlying table is updated, the materialized view refreshes automatically, requiring no additional maintenance. Snowflake-managed services perform the update in the background transparent to the user without interfering with the user's experience.

<https://docs.snowflake.com/en/user-guide/views-materialized>

## Domain

Performance Concepts

### Question 7Correct

What is the minimum Snowflake edition required to browse Snowflake Marketplace?

Business Critical

Enterprise

Virtual Private Snowflake

Standard

Your answer is correct

You don't need a Snowflake edition to browse the Snowflake marketplace listings

Overall explanation

Snowflake Marketplace can be browsed by non-Snowflake users as well, however they need to sign up to a Snowflake edition in order to consume data from the marketplace. Data Marketplace is supported in all Snowflake editions; thus, the minimum edition that

supports it is the Standard edition. Do note that VPS doesn't support Data Marketplace.  
<https://docs.snowflake.com/en/user-guide/intro-editions.html>

## Domain

Licensing & Features

### Question 8Correct

Which of the following are not supported by Search Optimization?

Select three answers.

Integer Columns

Your selection is correct

Cast on table columns

Your selection is correct

External Tables

Your selection is correct

Materialized Views

Date Columns

Overall explanation

Search optimization does not support

- External tables.
- Materialized views.
- Columns defined with a COLLATE clause.
- Column concatenation.
- Analytical expressions.
- Casts on table columns (except for fixed-point numbers cast to strings).

<https://docs.snowflake.com/en/user-guide/search-optimization-service#queries-not-supported-by-the-search-optimization-service>

## Domain

Performance Concepts

### Question 9Correct

Which of the following is true regarding Directory Tables?

Select all that apply.

Your selection is correct

To use a stream with a directory table, you must create the stream on the stage object.

Streams can NOT be used with directory tables.

Your selection is correct

Streams can be used with directory tables.

**To use a stream with a directory table, you must create the stream on the directory table object.**

Overall explanation

Streams can be used with directory tables to easily track which files have been added, removed, or changed. This is done by creating a stream on top of the stage object.

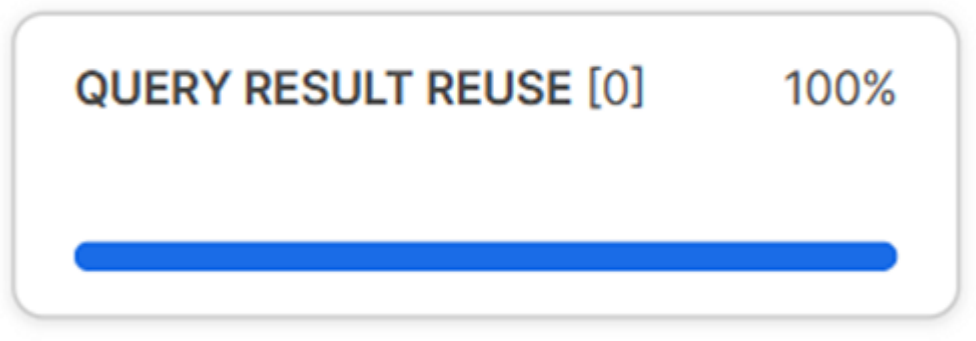
<https://docs.snowflake.com/en/user-guide/data-load-dirtables-manage#streams-on-directory-tables>

## Domain

Data Transformation

### Question 10Correct

**Which of the following correctly describes the query profile shown?**



Select all that apply.

**Your selection is correct**

The query profile indicates that results produced by a previous query were reused.

**Your selection is correct**

The query profile indicates that an active virtual warehouse was NOT required for this query.

The query profile indicates that the metadata cache was used.

The query profile indicates that the virtual warehouse cache was used.

Overall explanation

When Snowflake runs a query, it caches the results of that query for a predetermined amount of time. The stored query results are referred to as the Query Result Cache. The Query Result Cache can be used to fulfill future queries if they are like a previously executed query & there have been no changes to the data in the tables being queried.

<https://docs.snowflake.com/en/user-guide/querying-persisted-results>

## Domain

Performance Concepts

### Question 11Correct

**Which of the following Scaling Policy aims to preserve costs?**

Standard

**Your answer is correct**

Economy

Efficient

Fast

### Overall explanation

The Economy scaling policy attempts to conserve credits over performance and user experience. It doesn't spin up more virtual warehouses as soon as queuing is observed but instead applies additional criteria to ascertain whether or not to spin up new virtual warehouses. With the scaling policy set to Standard, Snowflake prefers to spin up extra virtual warehouses almost as soon as it detects that queries are starting to queue up.

The Standard scaling policy aims to prevent or minimize queuing.

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

### Domain

Performance Concepts

### Question 12Correct

True or False: Snowflake does not support the loading of semi-structured data.

True

Your answer is correct

False

### Overall explanation

Snowflake supports loading and processing semi-structured data. Snowflake provides the VARIANT data type, which can store any data and is appropriate for semi-structured data input and querying. SQL may be used to read and navigate JSON data once it has been loaded into a VARIANT column. <https://docs.snowflake.com/en/user-guide/semistructured-intro>

### Domain

Data Loading and Unloading

### Question 13Correct

Multi-factor authentication can be enabled for which of the following? Select all that apply.

Your selection is correct

Snowflake WebUI

Snowpipe

Your selection is correct

JDBC

Your selection is correct

SnowSQL

### Overall explanation

MFA is enabled by default for all Snowflake accounts and is available in all Snowflake editions. All Snowflake client tools, including the web interface, SnowSQL, and the various connectors and drivers, support MFA. Snowpipe is a snowflake-managed serverless service. A Snowflake user can not log into it; therefore, it doesn't require MFA. <https://docs.snowflake.com/en/user-guide/security-mfa>

### Domain

Security

### Question 14Correct

Which of the following strategies should be used to optimize the performance of a virtual warehouse?

Select two answers.

Configure memory spilling parameters to True.  
Increase the local disk cache size.  
Configure MAX\_CONCURRENCY\_LEVEL to a higher number  
Suspend the virtual warehouse.

**Your selection is correct**

**Reduce queuing.**

**Your selection is correct**

**Increase the virtual warehouse size.**

Overall explanation

The following strategies may be applied to improve the performance of a virtual warehouse.

1. Reduce queuing
2. Resolve memory spillage.
3. Increase warehouse size.
4. Try query acceleration.
5. Optimize the warehouse cache.
6. Limit concurrently running queries.

<https://docs.snowflake.com/en/user-guide/performance-query-warehouse>

## Domain

Performance Concepts

### Question 15 Correct

Which of the following statements accurately describes Snowflake's encryption for data at rest? Select all that apply.

**Your selection is correct**

**Snowflake rekeys encrypted data after 1 year**

**Your selection is correct**

**Snowflake manages encryption keys by default**

**Snowflake uses AES - 128-bit encryption to encrypt data at rest**

**Your selection is correct**

**Every 30 days, Snowflake rotates the keys used for encryption**

Overall explanation

By default, Snowflake manages encryption keys automatically, requiring no customer intervention. Snowflake-managed keys are rotated regularly (at 30-day intervals), and an annual rekeying process re-encrypts data with new keys. The data encryption and key management processes are entirely transparent to the users. Snowflake uses AES 256-bit encryption to encrypt data at rest. <https://docs.snowflake.com/en/user-guide/security-encryption-manage>

## Domain

Security

### Question 16Correct

Which of the following is correct regarding a directory table?

Select all that apply.

A directory table is a separate object.

Privileges can be assigned to a directory table.

Your selection is correct

A directory table is NOT a separate object.

Your selection is correct

Privileges can NOT be assigned directly to a directory table.

Overall explanation

A directory table is not a separate database object but is an implicit object available with a stage. You can enable the directory table for a stage while creating the stage or enable it afterward.

Since Directory Tables are not separate objects, you cannot provide privileges to them.

<https://docs.snowflake.com/en/user-guide/data-load-dirtables>

## Domain

Data Transformation

### Question 17Correct

True or False: If an IP address is in a network policy's block list and the allowed list, Snowflake applies the blocked list first.

Your answer is correct

True

False

Overall explanation

If both the allowed and blocked IP address lists are populated, Snowflake applies the block list first, followed by the allowed list. <https://docs.snowflake.com/en/user-guide/network-policies>

## Domain

Security

### Question 18Correct

Which of the following can create a new resource monitor?

Your answer is correct

An account administrator (i.e., a person with the ACCOUNTADMIN role)

A system administrator (i.e., a person with the SYSADMIN role).

A user who has MONITOR and MODIFY privilege on the resource monitor.

Any user of the system

Overall explanation

From a privilege perspective, only Account Administrators (users with ACCOUNTADMIN role) can create new resource monitors. However, account administrators can grant privileges to an existing resource monitor to allow other users



to view and modify the resource monitor configuration. The MONITOR and MODIFY privileges on a resource monitor allow other users to view and modify a specific resource monitor. <https://docs.snowflake.com/en/user-guide/resource-monitors#access-control-privileges-for-resource-monitors>

### Domain

Account Usage & Monitoring

### Question 19Correct

Which one of the following objects can NOT be cloned?

Schemas

External Stage

**Your answer is correct**

Named Internal Stage

Databases

Overall explanation

Named Internal Stages cannot be cloned. When a database or schema is cloned, any Snowpipe that points to a Named Internal Stage is not cloned. Named External Stages can be cloned. Since a table stage is associated with a table, it is automatically cloned when the table is cloned. Additionally, external tables cannot be cloned either. Databases, Schema, Tables, etc., can be cloned. <https://docs.snowflake.com/en/user-guide/object-clone#cloning-and-stages>

### Domain

Cloning

### Question 20Correct

True or False: The functions provided in INFORMATION\_SCHEMA can be used to view account-level information.

**Your answer is correct**

True

False

Overall explanation

The INFORMATION\_SCHEMA provides data on the objects in the parent database of the INFORMATION\_SCHEMA. It also provides data on account-level objects such as roles, warehouses, and databases. <https://docs.snowflake.com/en/sql-reference/info-schema#information-schema-views-and-table-functions>

### Domain

Account Usage & Monitoring

### Question 21Correct

Which of the following can be used to find the query ID of the 2nd most recent query executed in the current session?

SELECT LAST\_QUERY\_ID(-1);

**Your answer is correct**

SELECT LAST\_QUERY\_ID(-2);

SELECT LAST\_QUERY\_ID(1);

SELECT LAST\_QUERY\_ID(2);

Overall explanation

The LAST\_QUERY\_ID function returns the query ID of a specified query in the current session. The function takes a number as the parameter, which specifies the position of the query in the session.

The parameter can take positive or negative values. A negative value means you are attempting to fetch the most recent query in the session, where

-1 = most recent query

-2 = 2nd most recent query

, and so on. The function defaults to -1, so if no value is provided, it will return the query id of the most recent query.

A positive number returns the earliest queries in the session. i.e.

1 = first query

2 = 2nd query

[https://docs.snowflake.com/en/sql-reference/functions/last\\_query\\_id](https://docs.snowflake.com/en/sql-reference/functions/last_query_id)

## Domain

General

### Question 22 Correct

True or False: The COPY command in a Snowpipe definition supports the same transformation as provided by the typical COPY command.

**Your answer is correct**

**True**

**False**

Overall explanation

Snowpipe uses the COPY command in its definition and can support the same transformations available to the COPY command. <https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro#how-does-snowpipe-work>

## Domain

Data Loading and Unloading

### Question 23 Correct

Which of the following is true regarding data encryption when using PUT to upload data to a Snowflake internal stage? Select two.

**The PUT command does not support encryption.**

**Your selection is correct**

**Data is stored encrypted in the Snowflake internal stage.**

**Snowflake's internal stages do not support encryption.**

**Your selection is correct**

**Data is encrypted automatically at the client machine before being transmitted to the Snowflake internal stage.**

Overall explanation

Data is encrypted automatically at the client machine before being transmitted to the Snowflake internal stage. Once the data is in an internal stage, it is stored encrypted.

This is part of the end-to-end encryption managed by Snowflake.

<https://docs.snowflake.com/en/user-guide/security-encryption-end-to-end>

## Domain

Data Loading and Unloading

### Question 24Correct

Which type of database object does Snowflake provide to monitor changes made by DML commands (such as insert, update, and delete) to tables?

Your answer is correct

Stream

### Explanation

Snowflake Streams help you keep track of any changes made to a table, such as new data being added (inserts), existing data being modified (updates), or data being removed (deletes). They allow you to query and process only the changed data since the last offset. See the link for more details: <https://docs.snowflake.com/en/user-guide/streams-intro>

### Task

### Explanation

Tasks in Snowflake are automated, time-based processes that run SQL statements, including calling stored procedures, to perform operations such as data loading, transformation, and scheduled reporting. They can be set to run at specific intervals or triggered by other tasks, enabling complex, dependent workflows.

### Stored Procedure

### Explanation

Stored procedures allow complex operations and procedural logic to be executed. They can include control-flow statements and error handling and call SQL commands, enabling the automation and customization of database tasks.

### Snowpipe

### Explanation

Snowpipe facilitates continuous and real-time data ingestion into Snowflake, enabling uninterrupted streaming data integration.

## Domain

Streams

### Question 25Correct

Which of the following are Snowflake Data Integration partners? Select all that apply.

Your selection is correct

Talend

Your selection is correct

Matillion

Your selection is correct

AbInitio

Your selection is correct

Informatica

Your selection is correct

## IBM DataStage

Overall explanation

All of these are Data Integration partners of Snowflake. Please see <https://docs.Snowflake.com/en/user-guide/ecosystem.html>

### Domain

Partners

### Question 26Correct

What is the best way for a system administrator to determine the initial size of a new virtual warehouse?

Choose X-Small virtual warehouse to conserve costs.

**Your answer is correct**

Try different types of queries and warehouse sizes to find the optimum fit for your query needs and workload.

Choose 5X-Large virtual warehouse size to ensure performance.

Contact Snowflake to get help with determining the right size for your organization.

Overall explanation

Experiment with a defined set of queries against various warehouse sizes (e.g., X-Large, Large, Medium) warehouse sizes to determine the optimal combination for your specific query requirements and workload.

<https://docs.snowflake.com/en/user-guide/warehouses-considerations#selecting-an-initial-warehouse-size>

### Domain

Performance Concepts

### Question 27Incorrect

Which of the following is true for the DATABASE\_STORAGE\_USAGE\_HISTORY view in the ACCOUNT\_USAGE schema?

Select four answers.

**Your selection is correct**

This view contains information for the last 365 days.

This view does not show information for deleted databases.

**Correct selection**

This view shows the number of bytes of fail-safe storage used.

This view contains real-time information.

**Your selection is correct**

This view shows the number of bytes of database storage used, including Time Travel storage.

**Your selection is correct**

This view shows information for all databases, including deleted databases.

Overall explanation

The DATABASE\_STORAGE\_USAGE\_HISTORY view in the ACCOUNT\_USAGE schema shows the number of bytes of database storage used by each database, including information for data that is in Time Travel. The view also separately shows the number of bytes in fail-safe storage. Like other ACCOUNT\_USAGE views, the data for the last 365 days is shown; this view can have a latency of up to 3 hours (not real-time) and includes deleted objects.

[https://docs.snowflake.com/en/sql-reference/account-usage/database\\_storage\\_usage\\_history](https://docs.snowflake.com/en/sql-reference/account-usage/database_storage_usage_history)

### Domain

Account Usage & Monitoring

#### Question 28Correct

True or False: The ACCOUNT\_USAGE views contain information on objects that have been deleted.

**Your answer is correct**

**True**

**False**

Overall explanation

ACCOUNT\_USAGE views include information for all dropped objects. Many of these views include a DELETED column showing the dropped object's information.

INFORMATION\_SCHEMA does not include dropped objects. Furthermore, because objects can be dropped and recreated with the same name, the account use views include ID columns that display the internal IDs generated and assigned to each object by the system to differentiate amongst object records with the same name.

<https://docs.snowflake.com/en/sql-reference/account-usage#differences-between-account-usage-and-information-schema>

### Domain

Account Usage & Monitoring

#### Question 29Correct

What of the following Snowflake edition allows data clustering to improve query performance?

**Your selection is correct**

**Virtual Private Snowflake**

**Your selection is correct**

**Business Critical**

**Your selection is correct**

**Standard**

**Your selection is correct**

**Enterprise**

Overall explanation

All Snowflake editions support data clustering. <https://docs.snowflake.com/en/user-guide/intro-editions.html>

### Domain

Licensing & Features

#### Question 30Correct

**Snowflake recommends that any custom roles should be assigned to the pre-defined \_\_\_\_\_ role.**

USERADMIN

**Your answer is correct**

**SYSADMIN**

ORGADMIN

ACCOUNTADMIN

SECURITYADMIN

Overall explanation

Snowflake recommends establishing a hierarchy of custom roles, with the top custom role given to the pre-defined system role SYSADMIN. SYSADMIN can act as the owner of all securable objects in the system and can manage these objects.

<https://docs.snowflake.com/en/user-guide/security-access-control-overview#custom-roles>

**Domain**

Security

**Question 31**Correct

True or False: When unloading data, each exported file is 16MB, and this configuration cannot be changed.

True

**Your answer is correct**

**False**

Overall explanation

The default size of each output file is 16 MB but can be changed using the MAX\_FILE\_SIZE parameter. The maximum allowed size per file is 5GB if you export data to cloud storage. <https://docs.snowflake.com/en/user-guide/data-unload-considerations#unloading-to-a-single-file>

**Domain**

Data Loading and Unloading

**Question 32**Correct

**Which role must be granted to a Snowflake user to allow them to create new Snowflake accounts?**

SECURITYADMIN

GLOBALADMIN

SYSADMIN

ACCOUNTADMIN

**Your answer is correct**

**ORGADMIN**

Overall explanation

The ORGADMIN role performs organization-specific tasks like listing all accounts and creating new ones.

<https://docs.snowflake.com/en/user-guide/organizations-gs#enabling-the-orgadmin-role-in-an-account>

## Domain

Security

### Question 33Correct

Query Result Cache can be turned off at which levels? Select all that apply.

**Your selection is correct**

User

**Your selection is correct**

Account

**Your selection is correct**

Session

Table

Overall explanation

Query result cache is enabled by default but can be turned off at a session, user, or account level using the USE\_CACHED\_RESULT parameter.

<https://docs.snowflake.com/en/user-guide/querying-persisted-results>

## Domain

Performance Concepts

### Question 34Correct

Snowflake provides which of the following data protection features automatically?

Select all that apply.

**Your selection is correct**

Fail-Safe

**Your selection is correct**

Time Travel

Backups to on-premises

Incremental Backups

Tape Backups

Overall explanation

Snowflake provides Time Travel & Fail-Safe. Time Travel enables accessing, retrieving, and recovering past data stored in tables. The time travel can range between one and ninety days.

Failsafe storage retains data for an additional seven days after the time travel duration has expired. Failsafe adds another layer of protection against data loss; however, only Snowflake support can recover data from failsafe storage.

## Domain

Cloning

### Question 35Correct

Which of the following statements about micro-partitions is correct? Select all that apply.

**Column values can never overlap between micro-partitions.**

**Snowflake uses a row storage format to store columns in each micro-partition.**

**Your selection is correct**

**Micro-partitions are created and added to a table in the order that new data arrives.**

**Your selection is correct**

**Column values may overlap across micro-partitions.**

Overall explanation

Because micro-partitions are immutable and new or changed data must be added to a new micro-partition, similar values may not be in the same physical partition. When micro-partitions are added to a table, they are created in the order that the data came in. When more data is added to a table, another micro-partition or possibly many micro-partitions are created to store the new data. Unlike partitioning in many other databases, in Snowflake, values can overlap between different micro-partitions.

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions>

**Domain**

Architecture

**Question 36Correct**

Which of the following privileges allows a user in a consumer account to create a database from a share? Select two.

**SECURITY ADMIN role**

**Your selection is correct**

**ACCOUNTADMIN role**

**Your selection is correct**

**IMPORT SHARE privileges**

**SYSADMIN role**

Overall explanation

A user in a consumer account can create a database from the Share if they have the ACCOUNTADMIN role OR the IMPORT SHARE privileges

<https://docs.snowflake.com/en/user-guide/data-share-consumers>

**Domain**

Data Sharing

**Question 37Correct**

Through which of the following can users search for and consume third-party datasets?

**Partner Hub**

**Secure Data Sharing**

**Sharing Centre**

**Your answer is correct**

**Snowflake Marketplace**

Overall explanation

Snowflake Marketplace is the place to find and use third-party datasets that different organizations have made available. Users can search and consume third-party data sets using Snowflake Marketplace. <https://other-docs.snowflake.com/en/collaboration/collaboration-marketplace-about.html>

<https://other-docs.snowflake.com/en/collaboration/collaboration-marketplace-about.html>

**Domain**

Data Sharing

**Question 38Correct**



Which of the following statements is true regarding the SYSADMIN role? Select all that apply.

**Your selection is correct**

**A user with the SYSADMIN role can create a new database.**

**Your selection is correct**

**A user with the SYSADMIN role can create a new virtual warehouse.**

**A user with the SYSADMIN role can create new users.**

**A user with the SYSADMIN role can create new roles.**

Overall explanation

The SYSADMIN role can create and manage most Snowflake objects, including databases, tables, views, virtual warehouses, etc. However, the SYSADMIN role does not have the privileges to create new users or roles. <https://docs.snowflake.com/en/user-guide/security-access-control-overview#system-defined-roles>.

**Domain**

Security

**Question 39Correct**

Which of the following is a command line tool used to connect to Snowflake?

**Your answer is correct**

**SnowSQL**

**Snowsight**

**Snowmobile**

**Snowpipe**

Overall explanation

Command Line Client (CLI), also known as SnowSQL – is the method to connect to your Snowflake instance via a command-line interface. <https://docs.snowflake.com/en/user-guide/snowsql>

**Domain**

Tools & Interfaces

**Question 40Correct**

**True/False: Once the Time Travel retention period has ended for a transient table, historical data for that table can not be recovered by Snowflake support.**

**Your answer is correct**

**True**

**False**

Overall explanation

Transient and temporary tables don't have fail-safe functionality; therefore, data in such tables goes through zero days of fail-safe storage. Also, Transient and Temporary tables have a maximum of 1 day of Time Travel. Therefore, once the Time Travel period for these tables is complete, there is no way to recover historical data.

<https://docs.snowflake.com/en/user-guide/tables-temp-transient>

**Domain**

Data Protection

**Question 41Correct**

You are unloading data from a multi-gigabyte table to an external stage; which of the following statements regarding the exported file(s) are correct? Select all that apply.

**The data is exported to a single large file.**

**Your selection is correct**

**Each exported file is 16MB in size.**

**Your selection is correct**

**The exported file(s) are compressed.**

**Your selection is correct**

**The data is exported to multiple files.**

**The exported file(s) are NOT compressed.**

Overall explanation

When data is unloaded from Snowflake, it is automatically compressed using gzip compression. This is the default behavior; however, you can specify alternate compression methods or turn off compression entirely. The unloading process automatically exports to multiple files so that it can take advantage of the parallelism offered by Snowflake. However, if needed, you can set the SINGLE parameter to true to ensure the export goes to a single file. The default size of each output file is 16 MB but can be changed using the MAX\_FILE\_SIZE parameter. The maximum allowed size per file is 5GB if you export data to cloud storage. <https://docs.snowflake.com/en/user-guide/data-unload-considerations#unloading-to-a-single-file>

**Domain**

Data Loading and Unloading

**Question 42**Correct

The usage data provided through the INFORMATION SCHEMA has a retention of how many days?

**Your answer is correct**

**7 days - 6 months**

Forever

365 days

128 days

Overall explanation

The data in the INFORMATION\_SCHEMA views is retained for a shorter period. Typical data retention in INFORMATION\_SCHEMA is 14 days but can be seven days for specific views and up to 6 months for usage history views. Thus, these views have retention ranging from 7 days to a maximum of 6 months, depending on the view. So typically, the views in the INFORMATION\_SCHEMA can be used to find more recent information. <https://docs.snowflake.com/en/sql-reference/account-usage#differences-between-account-usage-and-information-schema>

**Domain**

Account Usage & Monitoring

**Question 43**Correct

**A business user is executing a complex query. Another user executed the same query less than 24 hours ago. Assuming that the underlying data for the query hasn't changed, what is the likely way this query will be executed?**

**Snowflake will return the query results from the metadata cache.**

**Your answer is correct**

**Snowflake will use the query result cache to fulfill the query because the query has previously been executed.**

**Snowflake will re-run the query, scanning the underlying tables and re-calculating the query results.**

Overall explanation

When Snowflake runs a query, it caches the results of that query for a predetermined amount of time. The stored query results are referred to as the Query Result Cache. The Query Result Cache can be used to fulfill future queries if they are similar to a previously executed query & there have been no changes to the data in the tables being

queried. Once a result cache is generated for a query stays valid for 24 hours. If another query that reuses the query result cache is executed within that 24-hour window, the result cache expiry is extended for another 24 hours from that point onwards. If the result cache for a query keeps getting used, it will stay valid for up to 31 days. After 31 days, the result cache for a query will be purged regardless of any other condition. <https://docs.snowflake.com/en/user-guide/querying-persisted-results>

## Domain

Performance Concepts

### Question 44Correct

Imagine an external stage named **FLIGHTS\_STAGE**.

Which of the following commands produce the same columns in the result set?

Select all that apply.

Your selection is correct

```
SELECT * FROM DIRECTORY(@FLIGHTS_STAGE);  
LIST @FLIGHTS_STAGE;
```

Your selection is correct

```
SELECT * FROM DIRECTORY(@FLIGHTS_STAGE) WHERE SIZE > 1000;
```

Overall explanation

The columns in the output obtained from querying a directory table differ from those in the output when listing a stage.

<https://docs.snowflake.com/en/user-guide/data-load-dirtables-manage#output>

<https://docs.snowflake.com/en/sql-reference/sql/list#output>

## Domain

Data Transformation

### Question 45Correct

True or False: When defining a clustering key, you should choose columns that have very high cardinality.

True

Your answer is correct

False

Overall explanation

When defining clustering keys, the initial candidate clustering columns are those columns that are frequently used in the WHERE clause or other selective filters. Additionally, columns that are used for joining can also be considered. Furthermore, the columns' cardinality (number of distinct values) is also important. It is crucial to choose a column with a high enough cardinality to allow effective partition pruning while having a low enough cardinality for Snowflake to group data into micro-partitions efficiently. A column with too few distinct values (e.g., gender) will result in minimal partition pruning. On the other hand, a column that has too many distinct values (e.g.,

customer id) will result in too much overhead when maintaining the partitions. When creating a multi-column cluster key, order the columns from the lowest cardinality to the higher cardinality; otherwise, the effectiveness of clustering will be reduced.  
<https://docs.snowflake.com/en/user-guide/tables-clustering-keys>

### Domain

Performance Concepts

### Question 46Correct

What is the minimum Snowflake edition that supports a dedicated metadata store?

Enterprise

Your answer is correct

Virtual Private Snowflake

Standard

Business Critical

Overall explanation

The VPS edition is meant to provide isolation from other customers; thus, each instance has its own metadata store and compute resources.

<https://docs.snowflake.com/en/user-guide/intro-editions.html>

### Domain

Licensing & Features

### Question 47Correct

How frequently does Snowflake release new software?

Daily

Fortnightly

Monthly

Yearly

Your answer is correct

Weekly

### Domain

General

### Question 48Correct

A virtual warehouse was started, used for 45 seconds, and shut down after that. The customer will be charged for how many seconds?

3600 seconds

315 seconds

45 seconds

Your answer is correct

60 seconds

Overall explanation

Snowflake credits are billed on a per-second usage basis. However, note that a minimum of 60 seconds of billing applies, so if a virtual warehouse were started and shut down within the first 1st minute, a minimum of 60-second credit usage would apply.

### Domain

Architecture

### Question 49Correct

Snowflake is compliant with which of the following standards? Select all that apply.

Your selection is correct

IRAP – Protected

**Your selection is correct**

**PCI-DSS**

**Your selection is correct**

**HIPAA**

**Your selection is correct**

**FedRAMP**

Overall explanation

Snowflake is compliant with the following security and financial standards.

- IRAP Protected
- ITAR
- FedRAMP Moderate
- GxP
- SOC 1 Type II
- SOC 2 Type II
- PCI-DSS
- HITRUST / HIPAA
- ISO/IEC 27001

<https://www.snowflake.com/snowflakes-security-compliance-reports/>

**Domain**

Security

**Question 50Correct**

A virtual warehouse is running and executing two queries. The virtual warehouse is resized to a smaller size. What best describes the resize operation?

**The queries are stopped, and the virtual warehouse is immediately resized to a smaller size.**

**Your answer is correct**

**The resize operation succeeds, but the node removal occurs once the active queries are finished.**

**The resize operation fails as a virtual warehouse cannot be resized while it is running queries**

**The queries are paused while the virtual warehouse is resized to a smaller size.**

Overall explanation

You can resize a virtual warehouse anytime, even when they are running queries. When resizing to a smaller size, nodes' removal occurs only when all active queries on those nodes have finished. <https://docs.snowflake.com/en/user-guide/warehouses-tasks#resizing-a-warehouse>

## Domain

### Architecture

#### Question 51 **Correct**

Assume a table with the structure below, which you have already loaded with JSON data.

```
CREATE TABLE my_json_table (  
  json_data VARIANT  
);
```

The JSON data looks like following

```
{  
  "data_set": "organizations",  
  "date_extracted": "2019-12-10",  
  "organizations": [  
    {  
      "Company": "Netus Et Malesuada Industries",  
      "State": "VIC",  
      "OrganisationCode": "36783603099"  
    },  
    {  
      "Company": "Amet Luctus PC",  
      "State": "NSW",  
      "OrganisationCode": "37908951399"  
    }  
  ]  
}
```

What is the correct way to access the "date\_extracted" value to be loaded into a relational table?

**SELECT**

**extract(date\_extracted from json\_data)**

**FROM my\_json\_table;**

**Your answer is correct**

**SELECT**

**json\_data:date\_extracted**

**FROM my\_json\_table;**

**There is no way to access this data as the data has been loaded into a VARIANT column.**

**Overall explanation**

Please see the following link on performing queries on JSON data.

<https://docs.snowflake.com/en/user-guide/semistructured-intro>

**Domain**

Data Loading and Unloading

**Question 52****Correct**

Which object types can be recovered using the UNDROP command after they have been dropped? Select all that apply.

**User**

**Your selection is correct**

**Database**

**Your selection is correct**

**Schema**

**Your selection is correct**

**Table**

**Role**

**Overall explanation**

The UNDROP functionality applies to tables, schemas, and databases. That means you can restore complete databases or schemas and their child objects.

<https://docs.snowflake.com/en/user-guide/data-time-travel#restoring-objects>

**Domain**

Time Travel

**Question 53****Correct**

True or False: An ACCOUNTADMIN can see the results of any query executed by any user in a Snowflake system.

**True**

**Your answer is correct**

**False**

**Overall explanation**

You can only view results for queries you have personally executed. For example, as an administrator, If you have permission to view queries run by another user, the Query Detail page displays the query's details but not the actual query result for data privacy reasons.

**Domain**

## Security

### Question 54Correct

True or False: A reader account can be used to share data with a non-Snowflake user or a non-Snowflake organization.

False

Your answer is correct

True

Overall explanation

Sharing data with a non-Snowflake user or organization is possible by creating a reader account. This reader account is created by the data provider solely for sharing purposes. <https://docs.snowflake.com/en/user-guide/data-sharing-reader-create>

Domain

Data Sharing

### Question 55Correct

Which of the following could be used as a remote service for an external function?

Your selection is correct

AWS Lambda Function

Your selection is correct

Microsoft Azure Function

Your selection is correct

Node.js running on an EC2 instance

Overall explanation

All of these are valid examples of how an external function could be implemented. <https://docs.snowflake.com/en/sql-reference/external-functions-introduction>

Domain

Extending Snowflake Functionality

### Question 56Correct

True or False: After a table has been cloned, any updates to the data in the source table will automatically update the data in the cloned table.

Your answer is correct

False

True

Overall explanation

Micro-partitions and metadata in the cloud services layer enable rapid and efficient zero-copy cloning because the cloned table's metadata references the existing micro-partitions. The source and cloned items are independent; thus, modifying data in one will not affect the other. For example, the source table can be dropped altogether, which doesn't affect the cloned table. <https://docs.snowflake.com/en/user-guide/tables-storage-considerations#label-cloning-tables>

Domain

Cloning

### Question 57Correct

True/False: It is possible to disable failsafe entirely for a Snowflake account.

Your answer is correct

False

True

Overall explanation

Once the Time Travel period ends, Snowflake keeps the data for a further 7-day period as further protection. This fail-safe can not be disabled or configured. You can NOT



change it for a Snowflake account, database, schema, or table. However, you can use Transient or Temporary tables, which have zero days of fail-safe storage.

<https://docs.snowflake.com/en/user-guide/data-failsafe>

<https://docs.snowflake.com/en/user-guide/tables-temp-transient>

### Domain

Fail-safe

### Question 58Correct

True/False: The storage capacity of your Snowflake-based data warehouse is virtually unlimited because it uses cloud storage as the underlying storage mechanism.

False

**Your answer is correct**

True

Overall explanation

Since Snowflake uses cloud-based storage like Amazon S3 or Azure Blob storage, the amount of space available to Snowflake is virtually unlimited.

### Domain

Architecture

### Question 59Correct

Snowflake can eliminate unneeded partitions while executing a query. What is the name given to this optimization technique?

**Predicate optimization**

**WHERE clause optimization**

**Retrieve needed data only (RNDO)**

**Your answer is correct**

**Partition pruning**

Overall explanation

Snowflake stores data in small partitions known as micro-partitions. Data in Snowflake tables is mapped to individual micro-partitions and structured in a columnar manner.

Micro-partitions are added to a table in the order in which the data is received.

Additional micro-partitions are produced when data is added to a table. Because the column values are scattered across numerous micro-partitions, Snowflake must keep track of what range of data is kept in which micro-partitions for each column. This metadata enables Snowflake to eliminate unnecessary micro-partitions when running queries, boosting performance. This process of eliminating micro-partitions is also known as partition pruning. <https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions#query-pruning>

### Domain

Performance Concepts

### Question 60Correct

Which of the following is an appropriate scenario for creating a Stored Procedure?

**Given two inputs, calculate the MAX of the two values.**

**Convert a character column into a numeric column**

**Your answer is correct**

**Execute one or more SQL statements that are assembled dynamically.**

Overall explanation

Stored procedures are often used to perform recurring administrative activities, e.g., in a particular organization setting up a new user on the system may require creating the user, granting them several roles, creating a private database from them, etc. These steps can easily be placed in a stored procedure, and then the stored procedure can be

called whenever there is a requirement to create a new user.

<https://docs.snowflake.com/en/sql-reference/stored-procedures-overview>

## Domain

Extending Snowflake Functionality

### Question 61Correct

**Which role owns a newly created object?**

**Your answer is correct**

**The role that was used by the user while creating the new object.**

ACCOUNTADMIN

PUBLIC

SYSADMIN

Overall explanation

Snowflake supports discretionary access control (DAC), which means that the role that created an object owns it and can provide access to other roles to that item.

<https://docs.snowflake.com/en/user-guide/security-access-control-overview>

## Domain

Security

### Question 62Correct

Which of the following criteria must be met for Snowflake to reuse the query result cache for a query? Choose all that apply.

**Your selection is correct**

**A new query matches an old query.**

**Your selection is correct**

**The query does not make use of runtime functions.**

**Your selection is correct**

**The table micro-partitions have NOT altered as a result of recluster or consolidation.**

**Your selection is correct**

**The query makes no use of user-defined or external functions.**

**Your selection is correct**

**The underlying data that contributes to the query results has remained unchanged.**

Overall explanation

All of these are correct. Snowflake uses the query result cache if the following conditions are met. A new query matches an old query, and the underlying data contributing to the query results remains unchanged. The table micro-partitions have not changed as a result of clustering or consolidation. The query makes no use of user-defined, external, or runtime functions. Note that queries that use the CURRENT DATE function are eligible for query result caching. <https://docs.snowflake.com/en/user-guide/querying-persisted-results>

## Domain

Architecture

### Question 63Correct

Which of the following statements are true regarding External Tables? Select all that apply.

**Your selection is correct**

**An external table can be joined with other tables.**

An external table can not be joined with other tables.

An external table can not be queried.

**Your selection is correct**

**You can query an external table just like a regular table.**

Overall explanation

Snowflake offers an alternative approach for tables called external tables, which permits the creation of tables with data stored in external cloud storage. External tables remove the need for the data to be loaded into Snowflake. In the case of an External table, the definition of the table is still stored in Snowflake metadata and consists of table structure, file locations, filenames, and other attributes. However, the table's data is saved outside of Snowflake. The external table functionality enables you to query external data like a standard table. External tables may be joined to other tables, and views may be created using them. <https://docs.snowflake.com/en/user-guide/tables-external-intro>

**Domain**

Data Loading and Unloading

**Question 64Correct**

Which of the following Snowflake editions support database failover and failback between Snowflake accounts, thus providing business continuity and disaster recovery? Select all that apply.

**Standard**

**Your selection is correct**

**Business Critical**

**Enterprise**

**Your selection is correct**

**Virtual Private Snowflake**

Overall explanation

Database failover and failback between Snowflake accounts are provided first in the Business Critical edition and are also available in the virtual private Snowflake (VPS) edition. <https://docs.snowflake.com/en/user-guide/intro-editions.html>

**Domain**

Licensing & Features

**Question 65Correct**

A stored procedure can return which type of results?

**Exe files**

**Your selection is correct**

**Tabular Data**

**Binary Executables**

**Your selection is correct**

**Single Value**

Overall explanation

A stored procedure can also return a single value or tabular data if desired; however, it is not a requirement that a stored procedure must return a value.

<https://docs.snowflake.com/en/developer-guide/stored-procedures-vs-udfs>

**Domain**

Extending Snowflake Functionality

**Question 66Correct**

Failsafe is provided as an alternate means to access historical data once the Time Travel retention period has ended.

**Your answer is correct**

**False**

**True**

Overall explanation

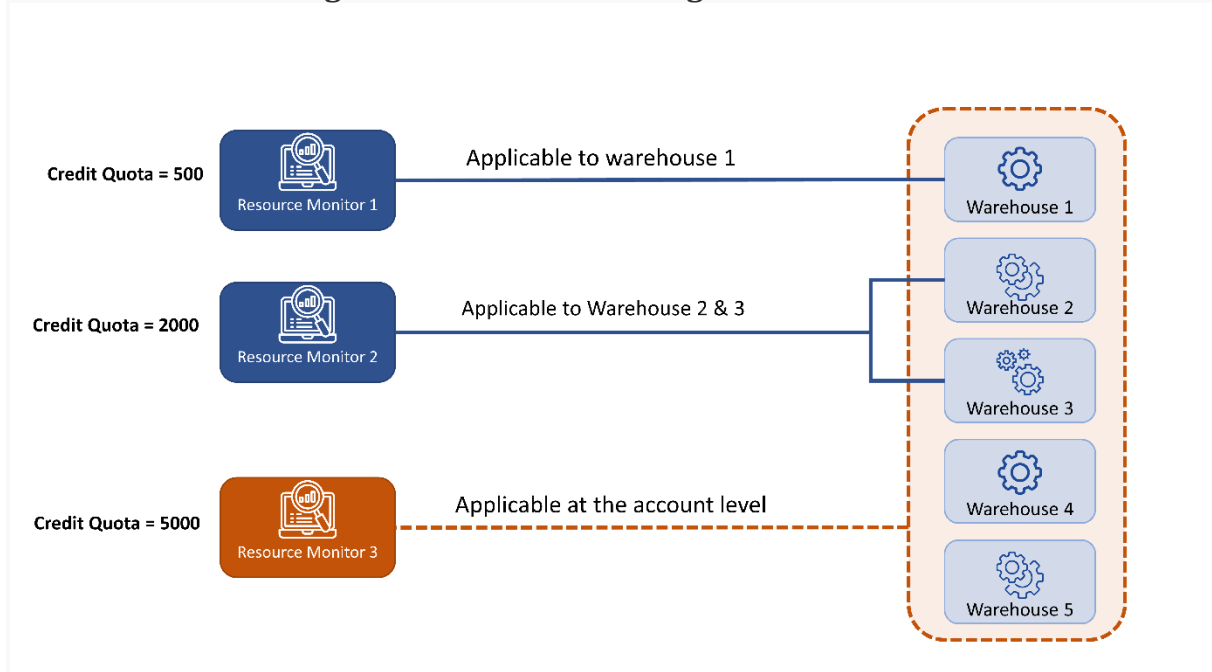
After the Time Travel period has been completed, the fail-safe storage feature stores data for an additional period of seven days. The fail-safe storage offers additional security against data loss; however, only the Snowflake support team can restore data from the fail-safe storage. Fail-safe cannot be used to access historical data but is used to recover from accidental data loss. <https://docs.snowflake.com/en/user-guide/data-failsafe>

**Domain**

Fail-safe

**Question 67 Correct**

Consider the following resource monitor configuration.



Which two of the given virtual warehouses can use a maximum of 5,000 credits?

Warehouse 1

Warehouse 3

**Your selection is correct**

Warehouse 5

**Your selection is correct**

Warehouse 4

Warehouse 2

Overall explanation

Resource monitors can track & manage a single virtual warehouse against a defined quota. Resource monitors can be created to track the credit usage of multiple virtual warehouses together.

Resource Monitors can also be created at the account level, which means that such resource monitors track credit usage at the account level, considering the credit usage of all virtual warehouses.

<https://docs.snowflake.com/en/user-guide/resource-monitors#assignment-of-resource-monitors>

## Domain

Account Usage & Monitoring

### Question 68Correct

You are the database administrator for a large retailer running Snowflake. There is an event table that contains more than 5TB of data. But it does not have a clustering key defined. You need to define a new cluster key on this table. What is the best method to add the cluster key?

The table structure is below, and the clustering key will be created on event\_date.

```
CREATE TABLE events (
```

```
Event_Date DATE,
```

```
Event_Id integer ,
```

```
Event_PayLoad string,
```

```
Event_Origin_Id integer
```

```
);
```

### Your answer is correct

Execute ALTER statement on the table to add the clustering key

```
ALTER TABLE events CLUSTER BY (Event_date);
```

Create a new table called events\_2 with the same structure as events while adding the clustering key.

```
CREATE TABLE events2 (
```

```
Event_Date DATE,
```

```
Event_Id integer ,
```

```
Event_PayLoad string,
```

**Event\_Origin\_Id integer**

**) CLUSTER BY (Event\_Date);**

**Then insert data from the event table into events2, drop the event table, and rename the events2 to be event.**

**Overall explanation**

The easiest way to add a cluster key to an existing table is by running the ALTER statement and using the CLUSTER BY clause to change the clustering key.

**Domain**

Performance Concepts

**Question 69**Correct

**You are the data modeler at a large retail organization that stores transactional data in a Snowflake table called "Transactions." The daily revenue reports are generated using the "Transactions" table, which calculates the revenue for the current day.**

**The "Transactions" table originally had 500GB of data but has now grown to 5TB. You have noticed that over time the performance of the daily revenue reports has degraded.**

**What is the most efficient & cost-effective way of optimizing performance?**

**Increase the size of the virtual warehouse executing the daily reports.**

**Your answer is correct**

**Cluster the Transactions table on the transaction date column**

**Create tables for each year of data, e.g., Transactions\_2017, Transactions\_2018, and Transactions\_2019. Then, insert the relevant data from the Transactions table into these \_year tables. Then, change your report to point to the Transactions\_2019 table.**

**Overall explanation**

Clustering a table on a specific column can optimize queries by eliminating unnecessary partitions from the query processing. A table can be re-clustered by defining a clustering key, which effectively redistributes the data into micro-partitions according to the clustering key, ensuring optimal access to queries that predicate or join on the clustered column. Clustering the table on the transaction date is the most efficient option. The daily report accesses one day at a time and benefits from the partition on the date column. <https://docs.snowflake.com/en/user-guide/tables-clustering-keys>

**Domain**

Performance Concepts

**Question 70**Correct

**Which of the following statements are correct regarding Time Travel & fail-safe storage? Select all that apply.**

**The maximum allowed Time Travel duration for a transient table is 7 days.**

**Your selection is correct**

**There is no fail-safe storage for a transient table.**

**Your selection is correct**

**The maximum allowed Time Travel duration for a transient table is 1 day.**

**A transient table has 7 days of fail-safe storage.**

Overall explanation

Transient and temporary tables don't have fail-safe functionality; therefore, data in such tables goes through zero days of fail-safe storage. Also, Transient and Temporary tables have a maximum of 1 day of Time Travel. <https://docs.snowflake.com/en/user-guide/tables-temp-transient>

**Domain**

Time Travel

**Question 71**Correct

Which of the following are caching mechanisms in Snowflake? Select all that apply.

**Your selection is correct**

**Metadata Caching**

**Memory Caching**

**Your selection is correct**

**Query Result Caching**

**Index Caching**

**Your selection is correct**

**Warehouse Caching**

Overall explanation

Metadata caching is used for queries that can be fulfilled directly from metadata, e.g., the row count of a table. Query Result Caching is for queries that have been executed already. Warehouse caching is within the virtual warehouse instance and is usually based on queries that have already been executed.

**Domain**

Performance Concepts

**Question 72**Correct

As part of a data processing pipeline, you are required to store data in an interim table. The subsequent processes then use the table in the pipeline. The data is deleted and reloaded every time the pipeline is executed. You are required to minimize data storage costs. Which type of table will you create?

**Your answer is correct**

**Transient**

**External**

**Temporary**

**Permanent**

Overall explanation

Based on the requirement, a Transient table is a good option. Transient tables don't have fail-safe storage and have only up to 1 day of Time Travel. Because the data in this table is deleted and reloaded daily, a transient table provides the best solution.

Transient tables are available across sessions, so different processes and sessions can access them. <https://docs.snowflake.com/en/user-guide/tables-temp-transient>

**Domain**

Data Protection

**Question 73**Correct

True or False: When defining a clustering key for a large table, consider using columns frequently used in WHERE clauses.

**Your answer is correct**

**True**

**False**

### Overall explanation

When defining clustering keys, the initial candidate clustering columns are those columns that are frequently used in the WHERE clause or other selective filters. Additionally, columns that are used for joining can also be considered. Furthermore, the columns' cardinality (number of distinct values) is also important. It is crucial to choose a column with a high enough cardinality to allow effective partition pruning while having a low enough cardinality for Snowflake to group data into micro-partitions efficiently. A column with too few distinct values (e.g., gender) will result in minimal partition pruning. On the other hand, a column that has too many distinct values (e.g., customer id) will result in too much overhead when maintaining the partitions. When creating a multi-column cluster key, order the columns from the lowest cardinality to the higher cardinality; otherwise, the effectiveness of clustering will be reduced.

<https://docs.snowflake.com/en/user-guide/tables-clustering-keys>

### Domain

Performance Concepts

#### Question 74 **Correct**

For how long a query remains visible on the query history page in the Snowsight interface?

28 days

**Your answer is correct**

14 days

60 minutes

3 months

### Overall explanation

The query history page lets users view the history of executed and currently executing queries. The query history page can show the history of queries executed in the last 14 days. <https://docs.snowflake.com/en/user-guide/ui-snowsight-activity#query-history>

### Domain

Account Usage & Monitoring

#### Question 75 **Correct**

Which of the following statements is true regarding the ACCOUNTADMIN role? Select all that apply.

**Your selection is correct**

**ACCOUNTADMIN role has full access rights and is the most powerful account.**

**A user with the ACCOUNTADMIN role can NOT create a new reader account.**

**Your selection is correct**

**A user with the ACCOUNTADMIN role can create & manage resource monitors**

**A user with the ACCOUNTADMIN role can NOT view billing information.**

### Overall explanation

ACCOUNTADMIN is the account administrator role with full access rights. As the most powerful role in the organization, access to this role should be rigorously managed. This role encapsulates the SECURITYADMIN and SYSADMIN roles, therefore, has all the privileges of SYSADMIN and SECURITYADMIN too.

<https://docs.snowflake.com/en/user-guide/security-access-control-overview#system-defined-roles>

### Domain

Security

#### Question 76 **Correct**

Which of the following correctly describe Snowpipe?



**Snowpipe is used for backup and recovery.**

**Your answer is correct**

**Snowpipe is used to load a small volume of data that arrives frequently and continuously.**

**Snowpipe is a security mechanism providing a secure pipe for communication. Snowpipe is used to load large volumes of data in a batch manner.**

**Overall explanation**

Snowflake allows continuous data loading using Snowpipe, a serverless service. Snowpipe enables you to load data in a micro-batch manner, loading small volumes of data on each execution. The micro-batch-based data loading is used when a continuous stream of data, such as transactions or events, must be loaded and made available to enterprises quickly. Snowpipe enables continuous data loading and can load data within a few minutes after it arrives in a stage. Snowpipe is serverless and has its own computational capability; therefore, it does not rely on virtual warehouses for processing. Snowflake automatically manages the compute required by a Snowpipe. Snowflake also manages the scaling up and down of a Snowpipe as per the data load requirement. Since a Snowpipe is serverless, its costs are charged separately from virtual warehousing fees. <https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro>

**Domain**

Data Loading and Unloading

**Question 77Correct**

Snowflake database is based on the massively parallel shared nothing architecture used by databases like Teradata and Greenplum and data lakes like Hadoop.

**Your answer is correct**

**No**

**Yes**

**Overall explanation**

Snowflake implements a new hybrid architecture that combines the best features of shared-disk and shared-nothing architectures. Snowflake stores data similarly to a shared-disk architecture, i.e., the data is shared. But it also allows for using several compute engines, each with its own memory and processing capabilities. <https://docs.snowflake.com/en/user-guide/intro-key-concepts#snowflake-architecture>

**Domain**

Architecture

**Question 78Correct**

True or False: Files already loaded from a stage to a table can be loaded again into a cloned table.

**Your answer is correct**

**True**

**False**

**Overall explanation**

Cloning doesn't copy the load metadata of a cloned table. Therefore, the load metadata for a cloned table would be empty. Thus, files already loaded for the source table can be loaded again into the cloned table.

**Domain**

Data Loading and Unloading

**Question 79Correct**

True or False: Snowflake applies new software versions to all Snowflake customers at once?

**Your answer is correct**

**False**

**True**

Overall explanation

Snowflake does not instantly deploy a new version to all Snowflake accounts; rather, customer accounts are moved into the new release over time in a phased manner. Day 1 (early access): Deployed for Enterprise edition (or higher) accounts that have elected for early access. You can enroll an Enterprise edition (or higher) account for early access by contacting Snowflake support. Day 1 or 2 (regular access): Deployment of all Snowflake accounts on the Standard edition. Day 2 (last): All remaining Enterprise edition (or higher) accounts are deployed. Between an early access deployment and a final deployment, a minimum of 24 hours must pass. This staged release strategy enables Snowflake to identify and address any software issues uncovered during early access. <https://docs.snowflake.com/en/user-guide/intro-releases>

**Domain**

Account

**Question 80Correct**

Which of the following operations can be performed on a cloned table?

**Your selection is correct**

**DELETE**

**Your selection is correct**

**DROP**

**Your selection is correct**

**SELECT**

**Your selection is correct**

**CLONE**

Overall explanation

All these operations can be performed on a cloned table because a cloned table is just like any other table. <https://docs.snowflake.com/en/user-guide/tables-storage-considerations#label-cloning-tables>

**Domain**

Cloning

**Question 81Correct**

True or False: In Snowsight (Snowflake web user interface), you can execute only one query at a given time.

**True**

**Your answer is correct**

**False**

Overall explanation

Multiple worksheets can be opened in Snowsight, each with a different query. The queries continue to execute even if the worksheets are inactive; thus, multiple queries can be executed simultaneously. <https://docs.snowflake.com/en/user-guide/ui-snowsight>

**Domain**

Tools & Interfaces

**Question 82Correct**

Cloning a schema will clone which of the following. Select all that apply.

**Your selection is correct**

**All tables in the schema**

**Your selection is correct**

**All other cloneable objects in the schema**

**Your selection is correct**

**The schema itself**

Overall explanation

When a schema is cloned, all child objects within the schema are cloned.

<https://docs.snowflake.com/en/sql-reference/sql/create-clone#additional-rules-that-apply-to-cloning-objects>

**Domain**

Cloning

**Question 83Correct**

Snowflake provides which of the following drivers?

**Your selection is correct**

**JDBC driver for Snowflake**

**Your selection is correct**

**Snowflake driver for the Go language**

**Your selection is correct**

**.NET driver for Snowflake**

**Your selection is correct**

**PHP PDO driver**

**Your selection is correct**

**ODBC driver for Snowflake**

Overall explanation

All of these are valid examples of drivers provided by Snowflake. Snowflake has several drivers and connectors that can be used to connect to your Snowflake instance. These include client tools made by Snowflake, like the web interface and the SnowSQL command-line interface, and drivers and connectors that let different languages and frameworks connect to Snowflake. The following drivers and connectors are currently available · Snowflake Connector for Python · Snowflake Connector for Spark · Snowflake Connector for Kafka · JDBC driver for Snowflake · ODBC driver for Snowflake · .NET driver for Snowflake · Snowflake driver for the Go language · Node.js drivers PHP PDO drivers

**Domain**

Tools & Interfaces

**Question 84Correct**

Assume a share has been granted to a consumer, and the consumer has created a database on the Share. Which of the following correctly describes what occurs if a new object is added to the Share?

**Adding objects to a share is impossible once created and granted to a consumer.**

**The consumer is required to re-create a database from the granted Share object.**

**Your answer is correct**

**The new object becomes accessible to the consumer immediately.**

Overall explanation

Once a share has been granted to a consumer, and the consumer has created a read-only database on the Share, all new objects added to the Share by the data provider automatically become accessible to the consumer as soon as they are added to the Share by the data provider.

## Domain

Data Sharing

### Question 85Correct

The Search Optimization service can be used to improve the performance of which type of queries?

Queries that access all columns & all rows

Any query performed on tables that have greater than 1TB of data

Your answer is correct

Selective point lookup queries

Queries that access a subset of a column in a table

Overall explanation

The search optimization service can be used to improve the performance of

1. Point lookup queries - return only one or a few rows using highly selective filters.
2. Substring & RegEx searches – queries that use LIKE, ILIKE, & RLIKE
3. Queries on fields in VARIANT, OBJECT & ARRAY columns – using equality conditions, IN, ARRAY\_CONTAINS, ARRAY\_OVERLAP, Substring & RegEx and NULL check conditions
4. Queries that use specific geospatial functions with GEOGRAPHY values.

<https://docs.snowflake.com/en/user-guide/search-optimization-service#understanding-the-search-optimization-service>

## Domain

Performance Concepts

### Question 86Correct

What type of virtual warehouse automatically lets you add or remove additional clusters as concurrency and demand change?

Suspended virtual warehouse

Non-virtual warehouse

Your answer is correct

Multi-cluster virtual warehouse

X-Large virtual warehouse

Overall explanation

Multi-cluster virtual warehouses are utilized when the number of concurrent users exceeds a single virtual warehouse's capacity. When the concurrent workload for a virtual warehouse reaches the maximum, new queries are queued. Multi-cluster virtual warehouses address this by adding clusters as needed. When the demand drops, the extra clusters are removed. <https://docs.snowflake.com/en/user-guide/warehouses-considerations#how-are-credits-charged-for-warehouses>

## Domain

Elastic Compute

### Question 87Correct

Which file function allows *any* user or application access to download unstructured data in a Snowflake stage?

GET\_DATA\_FROM\_STAGE

BUILD\_STAGE\_FILE\_URL

BUILD\_SCOPED\_FILE\_URL

**Your answer is correct**

GET\_PREIGNED\_URL

Overall explanation

A pre-signed URL is a simple HTTPS URL for accessing a file using a web browser. A pre-signed URL is generated using a pre-signed access token. Users can temporarily access a file via a pre-signed URL without authorization. The expiry duration of a pre-signed URL is configurable and can be set to the required duration.

<https://docs.snowflake.com/en/user-guide/unstructured-intro#types-of-urls-available-to-access-files>

## Domain

Data Transformation

**Question 88**Correct

Which one of the following is supported by Snowflake for the purpose of auto-provisioning users and group membership?

DAC

ABAC

RBAC

**Your answer is correct**

SCIM

Overall explanation

Snowflake supports SCIM 2.0 and is compatible with Okta and Azure Active Directory. SCIM is an open standard that provides automatic user provisioning and role synchronization based on identity provider information. When a new user is created in the identity provider, the SCIM automatically provisions the user in Snowflake. Additionally, SCIM can sync groups defined in an identity provider with Snowflake roles. <https://docs.snowflake.com/en/user-guide/scim>

## Domain

Security

**Question 89**Correct

An administrator cloned a table called Customer to a new table called Prospects. The administrator then proceeds to load new data into the Customer table. What can you expect to happen to the Prospects table?

Select all that apply.

The new data is loaded into the Customer table and its clone, the Prospects table.

**Your selection is correct**

The new data does not show up in the Prospects table.

**Your selection is correct**

The new data is loaded successfully into the Customer table only.

The data loading fails as it is impossible to load data into a table that was a source for a cloning operation.

Overall explanation

The source and cloned items are independent; thus, modifying data in one will not affect the other. For example, the source table can be dropped altogether, which doesn't affect the cloned table.

<https://docs.snowflake.com/en/user-guide/tables-storage-considerations#label-cloning-tables>

## Domain

Cloning

### Question 90Correct

Which of the following illustration represents the most well-clustered table?

1 A E F L M O P U V Z

2 A E F N M O P U V Z

3 A E F N P X M O V Z

4 A Z  
A Z  
A Z  
A Z  
A Z

3

4

2

Your answer is correct

1

Overall explanation

For a populated table, the clustering depth is the average depth of overlapping micro-partitions for specific columns. The clustering depth starts at 1 (for a well-clustered table) and can be a larger number. If the average depth is smaller, the data for the specified columns are better clustered. <https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions#label-clustering-depth>

## Domain

Performance Concepts

### Question 91Correct

True/False: The compute and storage can be scaled independently in Snowflake architecture.

**Your answer is correct**

**True**

**False**

Overall explanation

True. Snowflake stores data similarly to a shared-disk architecture, i.e., the data is shared. But it also allows for using several compute engines, each with its own memory and processing capabilities. This architecture allows Snowflake to scale compute and storage independently. <https://docs.snowflake.com/en/user-guide/intro-key-concepts#snowflake-architecture>

**Domain**

Architecture

**Question 92****Correct**

Which of the following statements about Materialized Views is correct? Choose all that apply.

**Your selection is correct**

**Materialized views are used to boost query performance. They pre-compute query results and physically store them**

**Your selection is correct**

**A Snowflake service that is invisible to users automatically maintains materialized views in the background.**

**Your selection is correct**

**A materialized view can provide pre-computed answers, enabling some queries to be answered faster.**

Overall explanation

All these statements are correct. A materialized view is a view that pre-computes data based on a SELECT query. The query's results are pre-computed and physically stored to enhance performance for similar queries that are executed in the future. When the underlying table is updated, the materialized view refreshes automatically, requiring no additional maintenance. Snowflake-managed services perform the update in the background transparent to the user without interfering with the user's experience. <https://docs.snowflake.com/en/user-guide/views-materialized>

**Domain**

Performance Concepts

**Question 93****Correct**

True/False: Snowflake customers can control the format using which Snowflake stores the data for a table.

**True**

**Your answer is correct**

**False**

Overall explanation

Snowflake stores data in a proprietary format on cloud object storage, such as AWS S3, Azure Blob Storage, or Google Cloud Storage. Users cannot see the actual files, look at how the data is stored, or access the file directly. Users can not change how Snowflake stores the data behind the scenes.

**Domain**

Architecture

**Question 94****Correct**

True or False. To create Snowflake instances in different regions, you must maintain a separate Snowflake account for each region.

**Your answer is correct**

**True**

**False**

Overall explanation

Each Snowflake account is hosted in a particular Snowflake region. To use Snowflake in multiple regions, a Snowflake customer needs to maintain multiple Snowflake accounts, at least one for each region. <https://docs.snowflake.com/en/user-guide/intro-regions.html>

**Domain**

Licensing & Features

**Question 95****Correct**

Which of the following statement is correct regarding Snowflake billing? Select all that apply.

**Your selection is correct**

**If a virtual warehouse is suspended, it does not contribute to the cost.**

**Snowflake billing is based on the amount of data processed by queries**

**Snowflake billing is based on the number of queries executed.**

**Your selection is correct**

**Snowflake billing is based on the actual used storage.**

Overall explanation

Virtual warehouses in a resumed (active) state contribute to the costs. However, it does not matter if the virtual warehouse is not running a query; if it is resumed, it contributes to the costs. <https://docs.snowflake.com/en/user-guide/cost-understanding-compute> Snowflake charges for data storage in database tables, files staged in internal stages, time travel history, and fail-safe storage. Snowflake doesn't charge on how much data a query processed. <https://docs.snowflake.com/en/user-guide/cost-understanding-overall#how-are-costs-incurred>

**Domain**

Cost & Pricing

**Question 96****Correct**

Snowflake can load data staged in which of the following?

**Your selection is correct**

**Google Cloud Storage**

**Your selection is correct**

**AWS S3**

**Oracle Cloud Storage**

**Your selection is correct**

**Internal Stage**

**VMWare Storage**

**Your selection is correct**

**Azure Blob Storage**

Overall explanation

Snowflake supports loading from Internal Stages and External Stages. External Stages can use AWS S3, Azure Blob, and Google Cloud Storage. Before data can be processed into a Snowflake table, it is typically first made available in a Snowflake stage. This allows Snowflake access to the data to be loaded into a table. Once the data is available



in a stage, the COPY command can be used to copy the data into a table.

<https://docs.snowflake.com/en/user-guide/data-load-overview>

## Domain

Data Loading and Unloading

### Question 97Correct

Which of the following best describes “Bytes spilled to remote storage” shown in a query profile?

“Bytes spilled to remote storage” indicates network issues.

**Your answer is correct**

“Bytes Spilled to remote storage” indicates that the volume of data could not fit in either the memory or the temporary storage of the virtual warehouse and had to be spilled to temporary cloud storage.

“Bytes spilled to remote storage” shows that a file from an on-premise system is being loaded.

“Bytes spilled to remote storage” indicates the amount of data uploaded using the PUT command.

Overall explanation

Snowflake saves data on the warehouse's local disk if it can't fit an operation into memory. Data spilling slows down queries because it requires more IO operations, and disk access is slower than memory access. "Bytes spilled to local storage." indicates local spillage. Snowflake will spill data to remote cloud storage if the local disk becomes full, which is even slower storage than the local disk, making this operation even slower.

"Bytes spilled to remote storage" in the query profile indicates remote spillage.

<https://docs.snowflake.com/en/user-guide/ui-query-profile#queries-too-large-to-fit-in-memory>

## Domain

Performance Concepts

### Question 98Correct

True or False: The data in the views in the ACCOUNT\_USAGE schema can have a latency of up to 3 hours.

False

**Your answer is correct**

True

Overall explanation

The ACCOUNT\_USAGE schema consists of several views that provide usage metrics and metadata information at the account level. Data provided by the ACCOUNT\_USAGE views is NOT real-time and refreshes typically with a lag of 45 minutes to 3 hours, depending on the view. The data in these views are retained for up to 365 days.

<https://docs.snowflake.com/en/sql-reference/account-usage#differences-between-account-usage-and-information-schema>

## Domain

Account Usage & Monitoring

### Question 99Correct

A Snowflake share can only have one consumer account added to it.

True

**Your answer is correct**

False

Overall explanation

A Snowflake share can have zero, one, or multiple consumers added to it

## Domain

Data Sharing

### Question 100Correct

Which one of the following correctly represents the storage hierarchy in Snowflake?

Database->Account->Schema->Table

Account->Table->Database->Schema

Your answer is correct

Account->Database->Schema->Table

Schema->Account->Database->Table

Overall explanation

In Snowflake, the highest level is a Snowflake Account. Customers can have as many accounts as they like. Within an account, you have databases. Each database contains one or more schemas. Schemas contain other Objects. Tables, views, file formats, sequences, UDFs, and stored procedures are all examples of objects available in a schema. An object can be contained in only one schema.

## Domain

Snowflake's catalogue and objects.

### Question 101Correct

Snowflake UDFs can be broadly categorized into which two based on how they return data.

Your selection is correct

Scalar UDFs

Interpreted UDFs

Compiled UDFs

Your selection is correct

Table UDFs

Overall explanation

There are broad types of UDF in terms of the kind of result they can return. Scalar UDFs return one row for each input row, with each output row containing a single column or value. You can also create table UDFs that return zero, one, or several rows for each input, with each result row containing multiple columns. UDTFs are another name for user-defined table functions. <https://docs.snowflake.com/en/sql-reference/udf-overview#scalar-and-tabular-functions>

## Domain

Extending Snowflake Functionality

### Question 102Correct

Which of the following Snowflake Editions support Time Travel? Select all that apply

Your selection is correct

Virtual Private Snowflake

Your selection is correct

Enterprise

Your selection is correct

Business Critical

Your selection is correct

Standard

Overall explanation

Time Travel is supported in all Snowflake editions.

<https://docs.snowflake.com/en/user-guide/data-time-travel#data-retention-period>

## Domain

Time Travel

**Question 103**Correct

Which of the following statements regarding Snowflake's built-in roles are correct?

Select two answers.

USERADMIN is granted to each new user automatically.

Your selection is correct

USERADMIN is dedicated to user and role management only.

Your selection is correct

The USERADMIN role separates the management of users and roles from the management of all grants.

USERADMIN inherits all permissions of SYSADMIN.

Overall explanation

USERADMIN is a role dedicated solely to user and role management. It has the privileges for CREATE USER and CREATE ROLE; therefore, users with this role can create users and roles in the account. USERADMIN can not manage grants, which is the job of SECURITYADMIN role.

<https://docs.snowflake.com/en/user-guide/security-access-control-overview#system-defined-roles>

**Domain**

Security

**Question 104**Correct

True or False: To share data as a provider and consume data as a consumer, you must have two Snowflake accounts, one for sharing data and one for consuming shared data.

Your answer is correct

False

True

Overall explanation

The same Snowflake account can share (as a data provider) and consume data (as a data consumer).

**Domain**

Data Sharing

**Question 105**Correct

Which of the following can be configured for a user profile in Snowsight?

Select two answers.

Your selection is correct

Default Role

**Your selection is correct**

**Default Warehouse**

**Default schema**

**Default database**

**Overall explanation**

Using the profile dialogue in Snowsight, you can enroll in MFA, specify your notification preferences, set your email address, and configure your profile's default role and default warehouse. Other things can also be configured such as name, password, language, etc.

<https://docs.snowflake.com/en/user-guide/ui-snowsight-profile>

**Domain**

Tools & Interfaces

**Question 106Correct**

Which of the following statement is true regarding the Query Processing Layer?

**Your selection is correct**

**The query processing layer is responsible for executing queries.**

The query processing layer is responsible for optimizing query plans

The query processing layer is responsible for generating query plans.

**Your selection is correct**

**The query processing layer can run multiple compute clusters (virtual warehouses) simultaneously.**

**Overall explanation**

The query processing layer is the compute layer through which queries and data processing jobs are executed on the stored data. The compute layer can have multiple clusters for a given Snowflake instance simultaneously. The compute engines in Snowflake are known as virtual warehouses. The cloud services layer performs the query plans and optimization. <https://docs.snowflake.com/en/user-guide/intro-key-concepts>

**Domain**

Architecture

**Question 107Correct**

Snowpark supports which of the following languages?

**Your selection is correct**

**Scala**

**Your selection is correct**

**Python**

**Your selection is correct**

**Java**

**Overall explanation**

Java, Scala & Python are all supported by Snowpark.

<https://docs.snowflake.com/en/developer-guide/snowpark/index>

**Domain**

Extending Snowflake Functionality

**Question 108Correct**

How does a secure UDF differ from a different typical UDF?

**Your selection is correct**

**Secure UDFs don't allow unauthorized users to see the UDF definition.**

**Your selection is correct**

**Secure UDF does not use specific SQL optimizations.**

**Secure UDFs can only be executed by administrators.**

**Secure UDFs are written in machine language and are compiled into Snowflake's code.**

**Overall explanation**

Specific SQL UDF optimizations may allow data that should be hidden from users to be accessed indirectly via different techniques. Secure UDFs do not use these SQL optimizations, ensuring that users have no access to the underlying data, even indirectly. Furthermore, Secure UDFs allow only authorized users to see the definition and information of secure UDFs (i.e., users who are granted the role that owns the UDF). <https://docs.snowflake.com/en/developer-guide/secure-udf-procedure>

**Domain**

Extending Snowflake Functionality

**Question 109Correct**

**Which system function can be used to control access to data in a share and allow specific data only to paying customers?**

**SYSTEM\$DISABLE\_BEHAVIOR\_CHANGE\_BUNDLE**

**SYSTEM\$BLOCK\_INTERNAL\_STAGES\_PUBLIC\_ACCESS**

**SYSTEM\$ALLOWLIST**

**Your answer is correct**

**SYSTEM\$IS\_LISTING\_PURCHASED**

**Overall explanation**

SYSTEM\$IS\_LISTING\_PURCHASED system function can be used to control which data is visible to a paid customer and which to a trial customer.

<https://other-docs.snowflake.com/en/collaboration/provider-listings-preparing#preparing-shares-for-a-paid-listing>

**Domain**

Data Sharing

**Question 110Correct**

**Which Snowflake layer manages the metadata related to micro-partitions, databases, and tables?**

**Your answer is correct**

**Cloud Services**

Query Processing

Client Tools

Database Storage

**Overall explanation**

The cloud services layer contains and manages a variety of metadata, including details regarding how the data is stored, information on the micro-partitions, metadata regarding the databases and tables in your system, the users, roles and security, and so forth.

**Domain**

Architecture

**Question 111**Correct

True/False: If you create more than one virtual warehouse, they will share the memory and CPU resources.

**Your answer is correct**

False

True

Overall explanation

Snowflake stores data in a shared manner, like in shared-disk architecture. But it also allows for using several compute engines, each with its own memory and processing capabilities. The virtual warehouses are independent of each other but access and process the same shared data. <https://docs.snowflake.com/en/user-guide/intro-key-concepts#snowflake-architecture>

**Domain**

Architecture

**Question 112**Correct

Which of the following scenarios is suitable for scaling up a virtual warehouse to a larger size?

The system has many concurrent queries

The system has many concurrent users.

**Your answer is correct**

**Complex queries are executed on the system and are required to finish faster.**

Overall explanation

Based on the complexity of the queries and the desired performance, a virtual warehouse can be scaled up or down. In general, increasing the virtual warehouse size improves query speed for CPU-intensive queries. On the other hand, scaling up is ineffective when dealing with a high number of concurrent users or queries. Instead, a multi-cluster virtual warehouse (scaling out) is utilized to accommodate an increased number of users and queries <https://docs.snowflake.com/en/user-guide/warehouses-considerations>

**Domain**

Performance Concepts

**Question 113**Correct

A data consumer has created a read-only database on a Share object shared by a data provider. The data provider adds an object to the Share. Which of the following statement correctly describe what happens?

Adding objects to a Share is impossible after a consumer creates a read-only database.

The data consumer needs to run "ALTER SHARE <share\_name> REFRESH" to ensure that the added object appears in the read-only database.

**Your answer is correct**

**The data consumer can see and consume the new object immediately.**

Overall explanation

All new objects added to a share object by the data provider automatically become accessible to the consumer <https://docs.snowflake.com/en/user-guide/data-sharing-intro#what-is-a-share>

**Domain**

Data Sharing

**Question 114**Correct

Which of the following roles can manage a Data Exchange share by default?

**USERADMIN**

**Your answer is correct**

**ACCOUNTADMIN**

**SECURITYADMIN**

**SYSADMIN**

Overall explanation

As a default, only the ACCOUNTADMIN role has the privileges to create and manage shares. However, if required, the privileges can be granted to other roles.

<https://docs.snowflake.com/en/user-guide/security-access-privileges-shares>

**Domain**

Security

**Question 115****Correct**

Which of the following transformations are NOT supported by the COPY command?  
Select all that apply.

**Truncate columns**

**Your selection is correct**

**GROUP BY**

**Your selection is correct**

**SUM**

**Your selection is correct**

**JOIN**

Overall explanation

When loading data into a table using the COPY command, Snowflake allows you to do simple transformations on the data as it is being loaded. During the load process, the COPY command allows for modifying the order of columns, omitting one or more columns, casting data into specified data types, and truncating values. While loading the data, complex transformations such as joins, filters, aggregations, and the use of FLATTEN are not supported as they are not essential data transformations. Therefore, joining, filtering, and aggregating the data are supported ONLY after the data has been loaded into a table. <https://docs.snowflake.com/en/user-guide/data-load-overview#id2>

**Domain**

Data Loading and Unloading

**Question 116****Correct**

Which of the following Snowflake edition doesn't support Snowflake Marketplace?

**Business Critical**

**Your answer is correct**

**VPS or Virtual Private Snowflake**

**Enterprise**

**Standard**

Overall explanation

All Snowflake accounts, except VPS Snowflake accounts, can use the Snowflake Marketplace. VPS accounts have isolated metadata and compute and, therefore, can't use the Snowflake marketplace built on the common cloud services and metadata provided by Snowflake. <https://other->

[docs.snowflake.com/en/collaboration/collaboration-marketplace-about.html#about-the-snowflake-marketplace](https://docs.snowflake.com/en/collaboration/collaboration-marketplace-about.html#about-the-snowflake-marketplace)

## Domain

Data Sharing

### Question 117Correct

Dynamic Data Masking provides what sort of security in Snowflake?

Row-level security

Database-level security

Your answer is correct

Column-level security

Object Security

Overall explanation

Snowflake supports masking policies that may be applied to columns and enforced at the column level to provide column-level security. Column-level security is achieved by dynamic data masking or external Tokenization. <https://docs.snowflake.com/en/user-guide/security-column>

## Domain

Security

### Question 118Correct

Which of the following views can be used to view the last 365 days of data loading history for data loaded through Snowpipe?

Select two answers.

ACCOUNT\_USAGE.LOAD\_HISTORY

INFORMATION\_SCHEMA.QUERY\_HISTORY

Your selection is correct

ACCOUNT\_USAGE.COPY\_HISTORY

Your selection is correct

ACCOUNT\_USAGE.PIPE\_USAGE\_HISTORY

Overall explanation

The COPY\_HISTORY view and the PIPE\_USAGE\_HISTORY view in the ACCOUNT\_USAGE schema provide the history of data loading performed through Snowpipe.

[https://docs.snowflake.com/en/sql-reference/account-usage/pipe\\_usage\\_history](https://docs.snowflake.com/en/sql-reference/account-usage/pipe_usage_history)

[https://docs.snowflake.com/en/sql-reference/account-usage/copy\\_history](https://docs.snowflake.com/en/sql-reference/account-usage/copy_history)

## Domain

Account Usage & Monitoring

### Question 119Correct

Which query profile results indicate that a large table may not be well clustered? Select all that apply.



**Your selection is correct**

**The value in the 'Partitions Total' equals 'Partitions Scanned.'**

**The Result node returns many rows.**

**Your selection is correct**

**A significant value for 'Partitions Scanned.'**

**There are many JoinFilter nodes.**

**Overall explanation**

Partition pruning occurs when the number of Partitions scanned is much smaller than Partitions total. If the partitions scanned equal the partition total, the query scanned the complete table. Therefore, no partition pruning happened, and the clustering key should be improved. <https://docs.snowflake.com/en/user-guide/ui-query-profile>

**Domain**

Performance Concepts

**Question 120****Correct**

True or False: When there are more queries than a virtual warehouse can handle, the queries start queuing.

**False**

**Your answer is correct**

**True**

**Overall explanation**

When queries are sent to a warehouse, the warehouse allocates the resources required for each query and begins running the queries. If there aren't enough resources to run all the queries sent to the warehouse, Snowflake queues the extra queries until the resources are available again. Snowflake provides multi-cluster virtual warehouses to overcome this issue. <https://docs.snowflake.com/en/user-guide/warehouses-multicluster>

**Domain**

Performance Concepts

**Question 121****Correct**

Which of the following statement is true regarding key pair authentication in Snowflake? Select all that apply.

**Your selection is correct**

**Key pair authentication is an alternative to simple username/password authentication.**

**Your selection is correct**

**Key pair authentication consists of a private key and one or two public keys.**

**Key pair authentication enables single sign-on (SSO).**

**Key pair authentication requires the user to provide a password.**

**Overall explanation**

Snowflake provides an additional layer of security by supporting key pair authentication in addition to the standard username/password login. This approach comprises private and public keys, with the public key allocated to a user and the private key used for authentication. The user provides a public key during authentication. A user can have up to two public keys, which can be rotated at any point in time. Key pair authentication is supported by all SnowSQL and Snowflake drivers and connectors. All Snowflake editions support Key-pair authentication <https://docs.snowflake.com/en/user-guide/key-pair-auth>

**Domain**

Security