

Practice Exam 4 - Results

Question 1 **Correct**

Which of the following objects do NOT support cloning operations?

Your selection is correct

Database created from a share

Transient Tables

Permanent Tables

Your selection is correct

Named Internal Stages

Overall explanation

A database created from a share object can not be cloned.

Named Internal Snowflake stages are also not supported by Cloning.

Domain

Cloning

Question 2 **Incorrect**

Which statements best describe scaling up a virtual warehouse in Snowflake? Select all that apply.

Your selection is correct

Generally, scaling up is performed to accommodate more complex workloads.

Correct selection

Additional compute resources are added to a virtual warehouse when scaling up.

Your selection is correct

When scaling up, the size of the virtual warehouse is increased to a larger size.

Your selection is incorrect

Scaling up means *decreasing* the size of a virtual warehouse due to improved query processing efficiency.

Overall explanation

It is possible to resize a virtual warehouse to suit changing workloads. Assume a customer began with a small virtual warehouse which may no longer be sufficient to meet increased query complexity. Scaling up the virtual warehouse can accommodate the increasing work complexity. When a virtual warehouse is scaled up, more nodes are added to the compute cluster. <https://docs.snowflake.com/en/user-guide/warehouses-tasks>

Domain

Architecture

Question 3 **Incorrect**

From a cardinality (number of distinct values) perspective, which columns should you consider adding to a clustering key?

Your selection is correct

Columns that have small enough distinct values to enable efficient grouping in micro-partition.

Only columns that have a high cardinality.

Your selection is incorrect

Only columns that have a low cardinality.

Correct selection

Columns that have large enough distinct values to enable efficient partition pruning.

The cardinality of columns doesn't matter when defining a clustering key.

Overall explanation

When defining clustering keys, the initial candidate clustering columns are those columns that are frequently used in the WHERE clause or other selective filters. Additionally, columns that are used for joining can also be considered. Furthermore, the columns' cardinality (number of distinct values) is also important. It is crucial to choose a column with a high enough cardinality to allow effective partition pruning while having a low enough cardinality for Snowflake to group data into micro-partitions efficiently. A column with too few distinct values (e.g., gender) will result in minimal partition pruning. On the other hand, a column that has too many distinct values (e.g., customer id) will result in too much overhead when maintaining the partitions. When creating a multi-column cluster key, order the columns from the lowest cardinality to the higher cardinality; otherwise, the effectiveness of clustering will be reduced.

<https://docs.snowflake.com/en/user-guide/tables-clustering-keys>

Domain

Performance Concepts

Question 4Correct

Which of the following table types allow Time Travel queries?

View-based Tables

Your selection is correct

Permanent Tables

External Tables

Your selection is correct

Temporary Tables

Your selection is correct

Transient Tables

Overall explanation

Transient and temporary tables don't have fail-safe functionality; therefore, data in such tables goes through zero days of fail-safe storage. However, Transient and Temporary tables do have a maximum of 1 day of Time Travel.

Permanent tables also have time travel capability and can go up to 90 days of time travel depending on the configuration.

<https://docs.snowflake.com/en/user-guide/tables-temp-transient>

Domain

Time Travel

Question 5Correct

To implement row-level security, what type of policies should be configured?

Your answer is correct

Row Access Policies

Object Policies

Masking Policies

Table Policies

Overall explanation

Row-level security is implemented by creating row access policies, which include conditions and functions that govern which rows are returned during query execution.

<https://docs.snowflake.com/en/user-guide/security-row-intro>

Domain

Security

Question 6 **Incorrect**

Which of the following use cases are suitable to execute on Snowflake?

Your selection is correct

Data Warehouse Implementation

Use as an Embedded Database

Correct selection

Data Lake Implementation

Your selection is correct

Data Science Processing

Your selection is correct

Data Sharing

Overall explanation

Snowflake is an analytical database for implementing data lakes, data warehouses, data lakehouses, running data science workloads, executing data engineering processes, business intelligence, ad-hoc analysis, data sharing, etc.

Domain

Licensing & Features

Question 7 **Correct**

What is the most appropriate method for a user to execute a set of SQL statements using a task?

Your answer is correct

Create a stored procedure that contains multiple SQL statements and invoke the stored procedure from the task.

CREATE TASK task1 AS call sp_with_multiple_SQLs();

Explanation

A task can only run one SQL statement at a time. To run multiple SQL statements, you have two options: either put all the statements in a stored procedure or create several tasks, each running a single statement, and link them together. However, using a stored procedure is the recommended approach.

See the following link for more information:

[https://docs.snowflake.com/en/user-guide/tasks-](https://docs.snowflake.com/en/user-guide/tasks-intro#:~:text=A%20task%20can%20execute%20any%20one%20of%20the%20following%20types%20of%20functions%3A)

[intro#:~:text=A%20task%20can%20execute%20any%20one%20of%20the%20following%20types%20of%20functions%3A](https://docs.snowflake.com/en/user-guide/tasks-intro#:~:text=A%20task%20can%20execute%20any%20one%20of%20the%20following%20types%20of%20functions%3A)

Create the task as a multi-step task and include all SQL steps sequentially.

```
CREATE MULTISTEP TASK task1
AS statement1; statement2;.
```

Explanation

Snowflake tasks are designed to execute a single SQL statement per task definition. There is no multi-step task in Snowflake, and placing multiple SQL statements sequentially within a task body is not supported.

Run multiple SQL statements directly within a single task without using stored procedures or multiple linked tasks. For example,

```
CREATE TASK task1
AS statement1;
statement2;.
```

Explanation

Snowflake tasks support executing only one SQL statement per task definition. Directly placing multiple SQL statements within a task definition is not supported and will result in an error.

Create a task for each SQL statement (such as task1, task2, etc.), and link them sequentially using a root task that triggers task1, task2, and subsequent tasks in a predefined order. Each task remains idle until its dependencies are fulfilled, ensuring a methodical execution sequence from start to finish.

Explanation

Managing multiple tasks for each individual SQL statement adds extra work and makes things more complicated.

Domain

Tasks

Question 8Correct

Snowflake stores data in fail-safe storage for Temporary tables for what duration?

14 days

7 days

Your answer is correct

0 days

1 day

Overall explanation

Snowflake has transient and temporary tables that don't provide fail-safe capabilities; hence, data in such tables have 0 days of fail-safe storage.

<https://docs.snowflake.com/en/user-guide/tables-temp-transient>

Domain

Fail-safe

Question 9Incorrect

Which of the following contributes toward the costs of a Snowflake system?

Number of queries

Correct selection

Cloud services layer

Your selection is correct

Query processing layer

Your selection is correct

Storage layer

Overall explanation

The query processing layer, storage layer, and cloud services layer all contribute towards the costs of a Snowflake system. <https://docs.snowflake.com/en/user-guide/cost-understanding-overall>

Domain

Cost & Pricing

Question 10Correct

Which of the following factors influence the cost of maintaining a materialized view?

Select two answers.

The number of columns in the base table.

Your selection is correct

The amount of data that changes in the materialized view when updates occur in the base table.

Your selection is correct

The storage required to store results of the query used in the materialized view definition.

The total amount of data in the base table.

Overall explanation

The costs of keeping data in materialized views are impacted by

- 1) The number of materialized views created for each base table.
- 2) The extent of data changes occurring in these materialized views when changes are made to the base table.
- 3) The number of these materialized views with a clustering key is defined.

Each materialized view saves the results of queries, which adds to the amount of storage space your account uses each month.

<https://docs.snowflake.com/en/user-guide/views-materialized#label-materialized-views-maintenance-billing>

Domain

Performance Concepts

Question 11Correct

Which of the following statement best describes a reader's account in Snowflake?
Reader accounts are used to synchronize data across various cloud providers.

Your answer is correct

A reader account is used to share data with a non-Snowflake user (or organization)

A reader account is used for read-only data-loading operations

A reader account is a method for access control in Snowflake

Overall explanation

In Snowflake, sharing data with a non-Snowflake user (or organization) is possible by creating a reader account for that user (or organization). This reader account is created and managed by the data provider. <https://docs.snowflake.com/en/user-guide/data-sharing-reader-create>

Domain

Data Sharing

Question 12Correct

How does Snowflake enforce network policy when an IP address is included in both the block list and the authorized list of the policy?

Snowflake network policy doesn't allow adding an IP address to both the block list and the allowed list

The network policy is invalid

Your answer is correct

Snowflake first applies the blocked list, guaranteeing that the IP address is denied access regardless of whether it is also defined in the allow list.

Snowflake first checks the approved list, guaranteeing that the IP address is permitted to connect, even if it is also defined on the block list.

Overall explanation

A network policy comprises of three components: a name, a list of approved IP addresses, and a list of blocked IP addresses. Snowflake applies the blocked list first if both the authorized and blocked lists are populated.

<https://docs.snowflake.com/en/user-guide/network-policies>

Domain

Security

Question 13Correct

Which Snowflake table types can be used to manage costs for short-lived tables?

Select two answers.

Permanent

Your selection is correct

Temporary

External

Aggregate

Your selection is correct

Transient

Overall explanation

Transient and temporary tables are a good option for short-lived data. They don't have fail-safe storage and have only up to 1 day of Time Travel.

<https://docs.snowflake.com/en/user-guide/tables-temp-transient>

Domain

Snowflake's catalog and objects

Question 14 **Incorrect**

True or False: Multi-factor authentication, or MFA, is enabled by default for all Snowflake accounts.

Your answer is incorrect

False

Correct answer

True

Overall explanation

MFA is enabled by default for all Snowflake accounts and is available in all Snowflake editions. <https://docs.snowflake.com/en/user-guide/security-mfa>

Domain

Security

Question 15 **Incorrect**

How does the Search Optimization service in Snowflake improve query performance?

Select two answers.

Your selection is correct

It uses a persistent data structure as an optimized search access path to speed up point lookups.

It pre-computes the results of common queries.

Correct selection

It improves the performance of queries that use equality predicate.

Your selection is incorrect

It re-distributes data according to the defined clustering key.

It increases the cache time-out so that more queries are answered from the cache.

Overall explanation

The search optimization service can significantly enhance the performance of some lookup and analytical queries that use many predicates for filtering.

The search optimization service uses a persistent data structure as an optimized search access path to speed up point lookups.

<https://docs.snowflake.com/en/user-guide/search-optimization-service>

Domain

Performance Concepts

Question 16 **Incorrect**

A suspended virtual warehouse is increased in size. Which of the following correctly describes what will happen?

Select all that apply.

The virtual warehouse is resumed immediately.

Your selection is correct

The configuration of the virtual warehouse is updated to reflect the new size.

Correct selection

Additional nodes are only provisioned when the virtual warehouse is resumed.

Your selection is correct

The warehouse remains suspended.

Additional nodes are immediately provisioned.

Overall explanation

You can resize a virtual warehouse at any time, even when suspended. When resizing a suspended virtual warehouse, its configuration is updated to reflect the new size. A suspended virtual warehouse has no provisioned nodes; it only has a configuration that tells Snowflake what to provision.

The provisioning of new nodes only occurs when the warehouse is resumed.

<https://docs.snowflake.com/en/user-guide/warehouses-tasks#resizing-a-warehouse>

Domain

Architecture

Question 17 **Correct**

For a non-ACCOUNTADMIN user, what privileges are required to create a share?

SECURITY privileges

MANAGE ACCOUNT privileges

Your answer is correct

CREATE SHARE privileges

CREATE ACCOUNT privileges

Overall explanation

Only the ACCOUNTADMIN role or roles specifically granted the CREATE SHARE privilege can create a share. <https://docs.snowflake.com/en/user-guide/data-sharing-gs>

Domain

Data Sharing

Question 18 **Correct**

True or False. A materialized view is automatically updated if the data in the underlying table is changed.

False

Your answer is correct

True

Overall explanation

A materialized view is a view that pre-computes data based on a SELECT query. The query's results are pre-computed and physically stored to enhance performance for similar queries that are executed in the future. When the underlying table is updated, the materialized view refreshes automatically, requiring no additional maintenance. Snowflake-managed services perform the update in the background transparent to the user without interfering with the user's experience.

<https://docs.snowflake.com/en/user-guide/views-materialized>

Domain

Performance Concepts

Question 19**Incorrect**

The expiry duration can be configured for which of the following URL types?

Correct answer

Pre-signed URL

File URL

Your answer is incorrect

Scoped URL

Overall explanation

A file URL is a permanent Snowflake-hosted URL to a staged file. File URLs don't expire.

A scoped URL is a temporary and encoded URL that allows temporary access to a staged file without requiring any privileges on the stage. A scoped URL expires after 24 hours.

A pre-signed URL is a simple HTTPS URL for accessing a file using a web browser. The expiry duration of a pre-signed URL is configurable and can be set to the required duration.

<https://docs.snowflake.com/en/user-guide/unstructured-intro#types-of-urls-available-to-access-files>

Domain

Data Transformation

Question 20**Incorrect**

Which of the following security features are supported by Snowflake? Select all that apply

MD5 encryption of data at rest

Your selection is correct

AES 256-bit encryption of data at rest

Correct selection

Tri-Secret Secure encryption

Your selection is correct

Key Rotation

Overall explanation

Snowflake uses AES 256-bit encryption as a standard for data at rest to protect all customer data. Snowflake rotates Snowflake-managed keys on an automatic basis every thirty days. When a Snowflake-controlled key and a customer-managed key are combined, the result is a composite master key that protects your data. Tri-Secret Secure is the name used to describe this combination of keys.

<https://docs.snowflake.com/en/user-guide/security-encryption-end-to-end>

Domain

Security

Question 21Incorrect

What is the minimum Snowflake edition that supports Row Access Policies?

Your answer is incorrect

Standard

Correct answer

Enterprise

Virtual Private Snowflake

Business Critical

Overall explanation

The Enterprise edition has several additional capabilities not provided in the Standard edition. These include multi-cluster virtual warehouses, column-level masking, row access policies, materialized views, and search optimization.

<https://docs.snowflake.com/en/user-guide/intro-editions.html>

Domain

Licensing & Features

Question 22Correct

What is the number of nodes in a Medium virtual warehouse?

Your answer is correct

4

6

3

256

Overall explanation

An X-Small virtual warehouse consists of a single node, the smallest possible configuration for a Snowflake virtual warehouse. A Small virtual warehouse consists of two nodes, and a Medium virtual warehouse is composed of four nodes. As the cluster size grows, the number of nodes in that cluster multiplies.

<https://docs.snowflake.com/en/user-guide/warehouses-overview>

Domain

Architecture

Question 23Correct

Which of the following are ways to provision a Snowflake instance?

Select all that apply.

Your selection is correct

Provision a Snowflake-hosted and managed instance (with Azure as the underlying infrastructure)

Install manually on virtual machines on an AWS account.

Your selection is correct

Provision a Snowflake-hosted and managed instance (with GCP as underlying infrastructure)

Install on-premises on commodity hardware.

Your selection is correct

Provision a Snowflake-hosted and managed instance (with AWS as underlying infrastructure)

Perform a hybrid installation using a mix of on-premises technology and cloud infrastructure.

Overall explanation

Snowflake is engineered for the cloud and is available only on AWS, Azure & GCP. A Snowflake instance must be provisioned through Snowflake. Instances are hosted by Snowflake using AWS, Azure, or GCP as the underlying infrastructure.

Domain

Licensing & Features

Question 24Correct

True/False: Snowflake automatically clusters data as it is inserted into a table.

Your answer is correct

True

False

Overall explanation

Snowflake automatically clusters the data as it is inserted into a table.

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions#what-is-data-clustering>

Domain

Architecture

Question 25Correct

Which of the following are true regarding External tables?

Data in external tables can be updated using the DML language.

Data in External tables can be deleted using the DELETE SQL command.

Your selection is correct

An external table points to an external stage.

Your selection is correct

External tables are read-only.

Overall explanation

Since external tables point to an external storage location, data manipulation language (DML) operations cannot be done on them. An external table can only be created against an external stage, which points to a cloud object storage location.

<https://docs.snowflake.com/en/user-guide/tables-external-intro>

Domain

Data Loading and Unloading

Question 26Incorrect

As part of a security incident investigation, you need to extract the history of who logged into Snowflake during the past 1 year. Which one of the following methods should you use?

Use the table function INFORMATION_SCHEMA.LOGIN_HISTORY_BY_USER()

Your answer is incorrect

Use the table function INFORMATION_SCHEMA.LOGIN_HISTORY ()

Correct answer

Query the view SNOWFLAKE.ACCOUNT_USAGE.LOGIN_HISTORY

Ask Snowflake support to provide these details.

Overall explanation

The views in ACCOUNT_USAGE schema provide up to 365 days of history. Therefore they are useful for scenarios where data for a long historical duration is required. The LOGIN_HISTORY view in the ACCOUNT_USAGE schema will therefore fulfill the requirement. Note: The ACCOUNT_USAGE views are not real-time and may have up to 3 hours of latency.

Domain

Account Usage & Monitoring

Question 27Incorrect

Which of the following scaling policies aims to minimize query queuing?

Your answer is incorrect

Economy

Correct answer

Standard

Overall explanation

With the scaling policy set to Standard, Snowflake prefers to spin up extra virtual warehouses almost as soon as it detects that queries are starting to queue up. The Standard scaling policy aims to prevent or minimize queuing. The Economy scaling policy attempts to conserve credits over performance and user experience. It doesn't spin up more virtual warehouses as soon as queuing is observed but instead applies additional criteria to ascertain whether or not to spin up new virtual warehouses. <https://docs.snowflake.com/en/user-guide/warehouses-multicloud#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Domain

Performance Concepts

Question 28Correct

Micro-partitions are small in size and (before compression) are generally of size _____.

500MB-1GB

Your answer is correct

50-500 MB

10-50MB

Overall explanation

Micro-partitions are small and typically store 50 MB to 500 MB of uncompressed data. <https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions.html>

Domain

Architecture

Question 29Correct

You are required to download data from a named internal stage to an on-premises system. Which command should you use?

Your answer is correct

GET

PUT

VALIDATE

COPY

Overall explanation

The GET command is used to download data from an internal stage to an on-premises system. The PUT command uploads data from an on-premises system to an internal stage. To download or upload data to an external stage, cloud provider utilities or other tools are used to interact with data in the cloud storage pointed to by the external stage. <https://docs.snowflake.com/en/user-guide/data-unload-overview#bulk-unloading-process>

Domain

Data Loading and Unloading

Question 30Correct

You need to create a new Reader Account to share a table with a non-Snowflake customer. You are logged in with a SYSADMIN role. When you run the query to create the new Reader Account, you are met with a privilege error. What do you need to do to be able to create the Reader Account?

Use the WebUI to create the Reader Account.

Switch the role to SECURITYADMIN. Only the SECURITYADMIN role is allowed to create new reader accounts.

Your answer is correct

Switch your role to ACCOUNTADMIN. Only the ACCOUNTADMIN role is allowed to create new reader accounts.

Overall explanation

Only the ACCOUNTADMIN role can create new Reader Accounts since it is an account-level activity. <https://docs.snowflake.com/en/user-guide/data-sharing-reader-create>

Domain

Security

Question 31Correct

Which Snowflake objects track DML (Data Manipulation Language) data changes so that operations can be performed on the changed data?

Stage

Explanation

In Snowflake, a Stage is a storage location for loading data into tables and unloading data from tables.

Your answer is correct

Stream

Explanation

Snowflake Streams help you keep track of any changes made to a table, such as new data being added (inserts), existing data being modified (updates), or data being removed (deletes). They allow you to query and process only the changed data since the last offset.

See the link for more details: <https://docs.snowflake.com/en/user-guide/streams-intro>

Task

Explanation

Tasks in Snowflake are automated, time-based processes that run SQL statements, including calling stored procedures, to perform operations such as data loading,

transformation, and scheduled reporting. They can be set to run at specific intervals or triggered by other tasks, enabling complex, dependent workflows.

Snowpipe

Explanation

Snowpipe facilitates continuous and real-time data ingestion into Snowflake, enabling uninterrupted streaming data integration.

Domain

Streams

Question 32Correct

What type of metadata is maintained for each micro-partition? Select all that apply.

Your selection is correct

Maximum & minimum of each column

Your selection is correct

Count of distinct values for each column

Your selection is correct

Range of column values

Mean, median & mode of each column

Overall explanation

Snowflake stores the range of column values in its metadata: the maximum and the minimum value for each column in each micro-partition. Snowflake can intelligently decide which partitions to read when processing a query using this metadata.

Additionally, Snowflake stores the count of distinct values for each column in each partition in the metadata and certain other information to assist in query optimization.

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions>

Domain

Architecture

Question 33Correct

A stored procedure has been configured as a caller's rights stored procedure. A business user executes the stored procedure. Which of the following statement correctly describes what will happen?

The stored procedure executes using the privileges of the PUBLIC role.

The stored procedure executes using the privileges of the role owning the stored procedure.

The stored procedure executes using the privileges of ACCOUNTADMIN.

Your answer is correct

The stored procedure executes using the privileges of the users executing the stored procedure.

Overall explanation

A stored procedure configured to run with callers' rights executes under the permissions of the calling user. <https://docs.snowflake.com/en/sql-reference/stored-procedures-rights>

Domain

Extending Snowflake Functionality

Question 34Correct

Which two system functions are provided by Snowflake to find information about clustering depth? Select all that apply.

Your selection is correct

SYSTEM\$CLUSTERING_INFORMATION

Your selection is correct

SYSTEM\$CLUSTERING_DEPTH

SYSTEM\$CLUSTERING_HISTORY

SYSTEM\$AVERAGE_DEPTH

Overall explanation

Snowflake provides two system functions that can be used to find out clustering information. The two functions are SYSTEM\$CLUSTERING_DEPTH

SYSTEM\$CLUSTERING_INFORMATION <https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions#clustering-depth>

Domain

Performance Concepts

Question 35Correct

You executed a Time Travel query that attempts to access the state of data as it was 120 days ago. What is the likely outcome?

Your answer is correct

The query fails with the following error. "Time travel data is not available for the table. The requested time is either beyond the allowed time travel period or before the object creation time."

The query executes successfully but shows no rows.

The query executes successfully and shows the data as it existed 120 days ago.

The query executes successfully but returns blank values in each row.

Overall explanation

A time travel query will fail, and an error will be returned if the TIMESTAMP, OFFSET, or STATEMENT supplied in the AT | BEFORE clause is outside the Time Travel retention period (90 days max) for the table.

The same error is thrown if the time travel query attempts to access the table data before it was created.

<https://docs.snowflake.com/en/user-guide/data-time-travel#querying-historical-data>

Domain

Time Travel

Question 36Incorrect

You are the database administrator for a large retailer running Snowflake. There is a transactions table that is clustered on the transaction_date column. You need to remove the clustering from this table. What is the correct way to do this?

Your answer is incorrect

Create a new table called transactions_2 with the same structure as transactions without any clustering key defined. Then insert data from the transactions table into transactions_2, drop the transactions table, and rename the transactions_2 as transactions.

ALTER TABLE transactions REMOVE CLUSTER BY;

Correct answer

ALTER TABLE transactions DROP CLUSTERING KEY;

Overall explanation

If you need to drop or alter the clustering keys for a table, you can run the ALTER table command and use the appropriate syntax to drop, add or change clustering keys.
<https://docs.snowflake.com/en/user-guide/tables-clustering-keys#dropping-the-clustering-keys-for-a-table>

Domain

Performance Concepts

Question 37Correct

What is the minimum Snowflake edition which supports automatic encryption of all data?

Your answer is correct

Standard

Virtual Private Snowflake

Enterprise

Business Critical

Overall explanation

Encryption is provided automatically in all Snowflake editions; thus, the minimum edition that supports it is the Standard edition. <https://docs.snowflake.com/en/user-guide/intro-editions.html>

Domain

Licensing & Features

Question 38Correct

Snowflake deploys new releases at what interval?

Daily

Fortnightly

Your answer is correct

Weekly

Monthly

Overall explanation

Snowflake releases new software weekly. The updates are done behind the scenes, so there is no downtime or disruption. To ensure that the customer always has the most up-to-date version of the software, the automatic releases can include bug fixes, new features, and other changes. <https://docs.snowflake.com/en/user-guide/intro-releases>

Domain

Account

Question 39Correct

Which of the statement correctly describes Snowflake architecture?

Shared storage & shared compute engine

Distributed storage and multiple compute engines

Your answer is correct

Shared storage and multiple compute engines

Overall explanation

Snowflake implements a new hybrid architecture that combines the best features of shared-disk and shared-nothing architectures. Snowflake stores data similarly to a shared-disk architecture, i.e., the data is shared. But it also allows for using several compute engines, each with its own memory and processing capabilities.

<https://docs.snowflake.com/en/user-guide/intro-key-concepts#snowflake-architecture>

Domain

Architecture

Question 40 **Incorrect**

True or False: A consumer account has created a read-only database on a shared database. The consumer account can further share this database & its tables with other accounts.

Your answer is incorrect

True

Correct answer

False

Overall explanation

Consumer accounts can only access and query data but cannot add, modify, or create database objects to a shared database.

Consumer accounts cannot clone a shared database, its schemas, or any of its tables.

Consumer accounts cannot use Time Travel on the shared data.

Consumer accounts cannot further share a shared database.

<https://docs.snowflake.com/en/user-guide/data-share-consumers#general-limitations-for-shared-databases>

Domain

Data Sharing

Question 41 **Incorrect**

Which of the following are data-sharing services provided by Snowflake?

Your selection is correct

Snowflake Marketplace

Correct selection

Data Exchange

Your selection is correct

Direct Sharing

Seamless Data Share

Overall explanation

Snowflake offers the following data-sharing services. Direct Sharing Snowflake Marketplace Data Exchange <https://docs.snowflake.com/en/guides-overview-sharing>

Domain

Data Sharing

Question 42 **Incorrect**

An external stage can be enabled during the creation of a stage object or enabled afterward. Which of the following SQL enables the directory table?

Your answer is incorrect

CREATE DIRECTORY TABLE <TABLE_NAME> ON <STAGE_NAME>;

Correct answer

ALTER STAGE <STAGE_NAME> SET DIRECTORY = (ENABLE = TRUE);

GENERATE DIRECTORY TABLE ON <STAGE_NAME>;

ALTER STAGE <STAGE_NAME> ADD DIRECTORY;

Overall explanation

The correct syntax for enabling (or disabling) a directory table for a stage is to use the DIRECTORY = (ENABLE = TRUE | FALSE) syntax while creating or altering a stage.

<https://docs.snowflake.com/en/sql-reference/sql/alter-stage>

Domain

Data Transformation

Question 43Correct

Which of the following Snowflake editions allows using Resource Monitors for monitoring virtual warehouse credit usage? Select all that apply.

Your selection is correct

Business Critical

Your selection is correct

Standard

Your selection is correct

Enterprise

Your selection is correct

Virtual Private Snowflake

Overall explanation

All Snowflake editions support Resource Monitors.

<https://docs.snowflake.com/en/user-guide/intro-editions.html>

Domain

Licensing & Features

Question 44Correct

How does Snowflake improve the performance of queries that exclude a significant amount of data when reading a table?

By using a full table scan

By using MFA

Your answer is correct

By using Partition Pruning

By using materialized views

Overall explanation

The Snowflake metadata allows the query engine to eliminate partitions to optimize query execution. For example, if the query specifies a WHERE condition, partitions NOT containing the value matching that condition will NOT be scanned.

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions#query-pruning>

Domain

Performance Concepts

Question 45Correct

As a data engineer, you are developing jobs to load data into a snowflake table. You have an S3 stage defined, containing a single file containing 100 rows. You have loaded those 100 rows.

You notice that your COPY command is executing successfully but is not loading any row into the target table.

What could be the reason?

Your answer is correct

You have already loaded the single file that was in the Snowflake stage. Snowflake tracks in metadata if a file has already been loaded and will not load it again.

The file format you have defined for your data is incorrect; therefore, no data is getting loaded.

The file in the Snowflake stage is corrupt and, therefore, can't be loaded.

Overall explanation

The load metadata stores a variety of information, such as the name of every file that was loaded into that table and the time stamp corresponding to the time that a file was loaded. By utilizing this load metadata, Snowflake ensures that it will not reprocess a previously loaded file. Since you have already loaded this file, Snowflake will track through metadata which files have been loaded and will not process them again. If you want to retest your copy command, you must place a new file or force the reload by specifying specific parameters during the load process.

<https://docs.snowflake.com/en/user-guide/data-load-considerations-load#load-metadata>

Domain

Data Loading and Unloading

Question 46Correct

Which of the following statements are true for Snowsight?

Select two.

Your selection is correct

Each worksheet in Snowsight can have a different context (role, virtual warehouse, database & schema)

One user can run only one query at a given time.

If a user switches to another worksheet, the Snowflake session ends.

Worksheets share sessions for a given user.

Your selection is correct

Each worksheet in Snowsight is a distinct Snowflake session.

Overall explanation

Each worksheet in Snowsight is a unique session. It can have a different context from other worksheets. A user can run multiple queries using multiple worksheets if required.

Switching to another worksheet does NOT end the session; instead, it stays active and if a query is running in the previous worksheet, it continues.

<https://docs.snowflake.com/en/user-guide/ui-snowsight-worksheets#change-the-session-context-for-a-worksheet>

<https://docs.snowflake.com/en/user-guide/ui-snowsight-worksheets#organizing-worksheets-in-folders>

Domain

Tools & Interfaces

Question 47Correct

Which of the following query profile snippet indicates effective micro-partition pruning?

1

Statistics

| | |
|-------------------------------|---------|
| Scan progress | 97.27% |
| Bytes scanned | 92.62GB |
| Percentage scanned from cache | 11.62% |
| Bytes sent over the network | 0.42MB |
| Partitions scanned | 84182 |
| Partitions total | 86546 |

2

Statistics

| | |
|-------------------------------|---------|
| Scan progress | 19.33% |
| Bytes scanned | 17.02GB |
| Percentage scanned from cache | 72.53% |
| Bytes written to result | 82.87MB |
| Bytes sent over the network | 19.86GB |
| Partitions scanned | 16732 |
| Partitions total | 86547 |

3

Statistics

| | |
|-------------------------------|---------|
| Scan progress | 19.54% |
| Bytes scanned | 18.90GB |
| Percentage scanned from cache | 2.74% |
| Bytes sent over the network | 0.10MB |
| Partitions scanned | 16912 |
| Partitions total | 86546 |

4

QUERY RESULT REUSE [0] 100%

Select all that apply.

Your selection is correct

3

1

4

Your selection is correct

2

Overall explanation

Partition pruning occurs when the number of Partitions scanned is much smaller than Partitions total. If the partitions scanned equal the partition total, the query scanned the complete table. Therefore, no partition pruning happened, and the clustering key should be improved. <https://docs.snowflake.com/en/user-guide/ui-query-profile>

Domain

Performance Concepts

Question 48Correct

Which of the following is true regarding the default role in Snowflake?

Select all that apply.

Your selection is correct

A user can switch their current role if required.

A user can not use any other role other than their default role.

When a user switches to another role in a session, their user's default role automatically reconfigures to the new role.

Your selection is correct

When users log in, their role is automatically set to their default role.

Overall explanation

Each Snowflake user is automatically assigned the PUBLIC role. Each user can be assigned additional roles, one of which can be set as their default role. A user's default role determines the role automatically used in Snowflake sessions that the user initiates; however, this is only a default. Users can switch roles at any point during a session.

<https://docs.snowflake.com/en/user-guide/admin-user-management#user-roles>

Domain

Security

Question 49Correct

True/False: Snowflake credits are billed on a per-second basis.

Your answer is correct

True

False

Overall explanation

Snowflake credits are billed on a per-second usage basis, which means if a virtual warehouse ran for 1 minute 45 seconds, you would be charged for 105 seconds (60 + 45). However, a minimum of 60 seconds of billing applies, so if a virtual warehouse were started and shut down within the first 1st minute, a minimum of 60-second credit usage would apply.

Domain

Architecture

Question 50Correct

You have shared a table with another Snowflake account. A user in the consumer account has executed a query on the shared table. Who will be charged for the query cost?

Your answer is correct

The data consumer

The data provider (the account which shared the data)

Overall explanation

In a data-sharing scenario, Snowflake charges the consumer account for the costs of any compute that the consumer account uses. <https://docs.Snowflake.net/manuals/user-guide/data-sharing-intro.html#how-does-secure-data-sharing-work>

Domain

Data Sharing

Question 51 **Incorrect**

Is it possible to share data with a Snowflake customer whose Snowflake instance exists in a different region than the data provider?

Correct answer

Yes, but you must enable replication first to enable cross-region data sharing.

No, it is not possible to share with customers in other regions.

Your answer is incorrect

Yes. Nothing special needs to be done to enable cross-region data sharing.

Overall explanation

It is possible to share data with other regions, but the provider must enable replication and replicate your existing database to the other region.

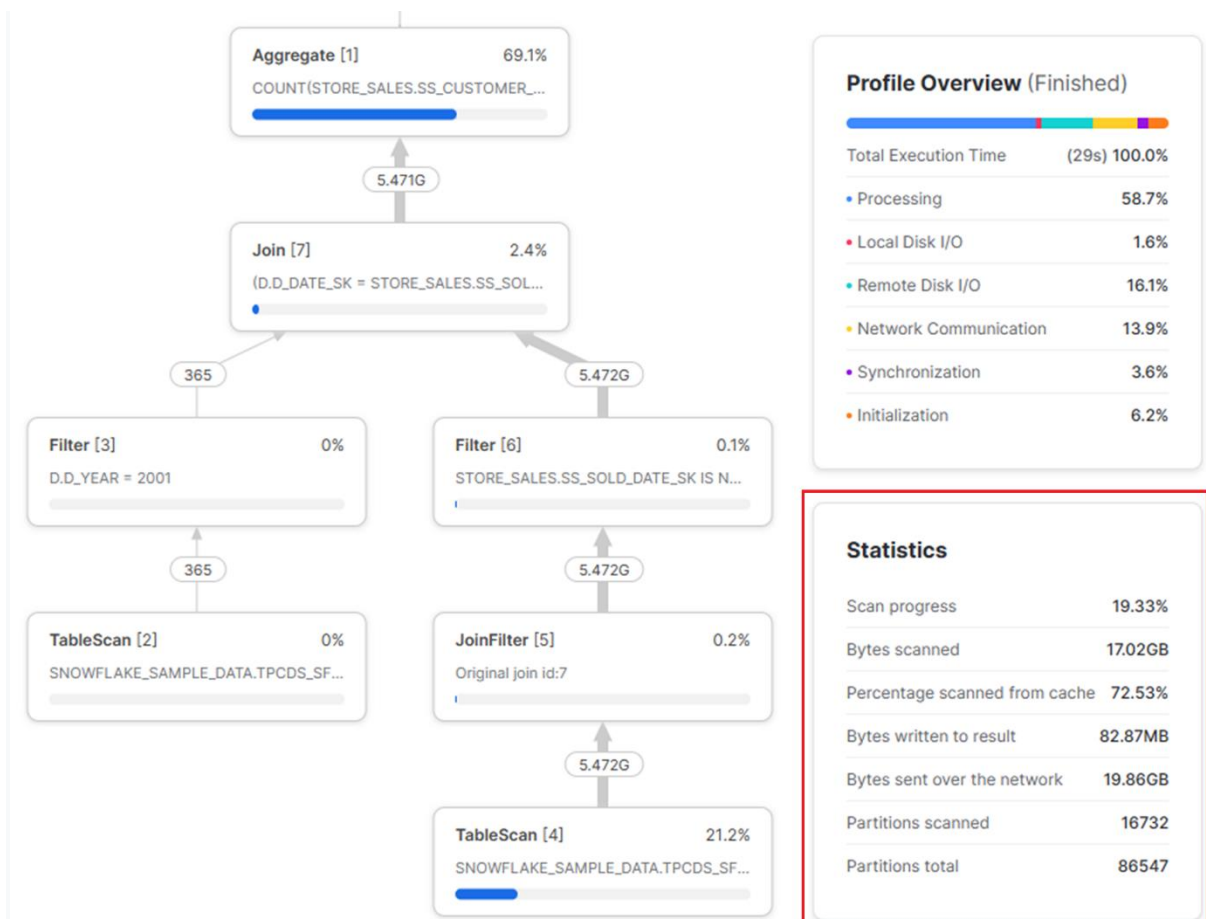
<https://docs.snowflake.com/en/user-guide/secure-data-sharing-across-regions-plaforms>

Domain

Data Sharing

Question 52 **Correct**

Consider the following snippet from the query profile of a finished query.



Which of the following accurately describes the highlighted statistics?

The query profile indicates ineffective partition pruning.

Your selection is correct

The query profile indicates that the virtual warehouse cache was used.

Your selection is correct

The query profile indicates effective partition pruning.

The query profile indicates that the virtual warehouse used is too small for the query.

Overall explanation

Partition pruning occurs when the number of Partitions scanned is much smaller than Partitions total. Snowflake saves data on the warehouse's local disk and sometimes remote cloud storage if it can't fit an operation into memory. Data spilling slows down queries because it requires more IO operations, and disk access is slower than memory access. Each virtual warehouse builds its cache by periodically copying relevant micro-partitions from cloud storage to local SSD storage. A virtual warehouse's cache may contain data for multiple queries with similar criteria, columns, and rows. The virtual warehouse can retrieve the data from its local cache instead of the cloud storage, speeding up queries. <https://docs.snowflake.com/en/user-guide/ui-query-profile>

Domain

Performance Concepts

Question 53 Incorrect

True/False: Snowflake creates new micro-partitions using the order of the data in which it is inserted.

Correct answer

True

Your answer is incorrect

False

Overall explanation

Micro partitions are added to a table in the order of how the data arrived in the table, so if new data is added to a table, new micro-partition(s) will be created to accommodate that data. Snowflake partitions are immutable; they can not be changed once created. Therefore, any update to existing data or loading new data into a table will create new micro-partitions. The micro-partitions are created and populated with data in the order of how the data arrived. <https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions>

Domain

Architecture

Question 54 **Correct**

Which of the following is not an architecture layer in Snowflake?

Database Storage

Your answer is correct

Virtual Machine

Virtual Warehouse

Cloud Services

Overall explanation

Snowflake architecture has three distinct layers: Database Storage - Cheap cloud storage on AWS, Azure, or Google Cloud Query Processing - Primarily composed of virtual warehouses Cloud Services - The brain of the whole operation <https://docs.snowflake.com/en/user-guide/intro-key-concepts#snowflake-architecture>

Domain

Architecture

Question 55 **Incorrect**

True or False: The INFORMATION_SCHEMA contains information on objects that have been deleted.

Your answer is incorrect

True

Correct answer

False

Overall explanation

INFORMATION_SCHEMA does NOT include information for any dropped objects. However, if the information on deleted objects is required, you must query the ACCOUNT_USAGE views. <https://docs.snowflake.com/en/sql-reference/account-usage#differences-between-account-usage-and-information-schema>

Domain

Account Usage & Monitoring

Question 56 **Correct**

Which of the following statements about Query Result Cache are correct? Choose all that apply.

Your selection is correct

The query result cache is used if an identical query is run within 24 hours of the original query.

The query result cache is kept in the virtual warehouse SSD storage.

Your selection is correct

If a query result cache is not reused within 24 hours, it is purged.

Your selection is correct

The cloud services layer stores the query result cache.

Your selection is correct

The maximum number of days a query result cache can be retained is 31.

Overall explanation

The query result cache is valid for 24 hours. If a new query matching a previous query is made within those 24 hours, the results are retrieved from the query result cache. The query result cache is maintained in the cloud services layer and can be utilized by any user to return query results. The result cache of a query is initially valid for 24 hours, but it is extended for another 24 hours when a new query uses it. This extension can last up to 31 days before being removed. It is purged if the query result cache is not used for 24 hours. <https://docs.snowflake.com/en/user-guide/querying-persisted-results>

Domain

Architecture

Question 57 **Incorrect**

Consider a database with the name MARKETING. The database has a table called CUSTOMER in the PUBLIC schema. The CUSTOMER table has 10,000 rows. The PUBLIC schema also has a view called CUSTOMER_COUNT with the following definition.

"SELECT COUNT(*) FROM MARKETING.PUBLIC.CUSTOMER;"

You create a temporary table with the same name, i.e., CUSTOMER, in the PUBLIC schema of the MARKETING database. The temporary table has zero rows.

Which of the following correctly describes the results when the view is queried?

Select two answers.

Your selection is correct

A "SELECT * FROM CUSTOMER_COUNT;" query will return 10,000 as the result when executed in a new session.

Your selection is incorrect

A "SELECT * FROM CUSTOMER_COUNT;" query will return 10,000 as the result when executed in the same session in which the temporary table was created.

Correct selection

A "SELECT * FROM CUSTOMER_COUNT;" query will return zero as the result when executed in the same session in which the temporary table was created.

A "SELECT * FROM CUSTOMER_COUNT;" query will return zero rows as the result executed in a new session.

Overall explanation

If a temporary table is created in a schema with the same name as a permanent (or transient) table, the temporary table effectively hides the permanent table in that session. Queries and other operations during the session will affect only the temporary table.

This behavior can affect the views as well. In the same session where a temporary table was created, a temporary table can hide the permanent table used by a view, resulting in unexpected results. However, the view is unaffected when queried from a different session because temporary tables are limited to the session where they are created.

<https://docs.snowflake.com/en/user-guide/tables-temp-transient#potential-naming-conflicts-with-other-table-types>

Domain

Snowflake's catalog and objects

Question 58Correct

When creating a new table, if no type is specified, the table is created automatically as which type?

Your answer is correct

Permanent

Transient

Temporary

Overall explanation

The default & implicit type for a new table is Permanent. To create other types of tables, you must explicitly specify the type as temporary, transient, or external.

Domain

Data Protection

Question 59Correct

Which one of the following is functionality provided by Snowpark?

Your answer is correct

Snowpark automatically converts the data-processing programming constructs to SQL and pushes them down to Snowflake for execution

Snowpark is a diagnostic tool designed to troubleshoot connectivity to Snowflake.

Snowpark is the name for the web-based user interface in Snowflake.

Snowpark is the command line client that allows you to connect to Snowflake and execute queries.

Overall explanation

Snowpark is a library created by Snowflake that provides APIs for accessing and processing data in applications written in a programming language other than SQL. Snowpark allows programmers to utilize common programming languages such as Java, Scala, and Python to construct apps that handle data using standard programming structures. Snowpark automatically converts the data-processing programming constructs to SQL and pushes it down to Snowflake for execution. As a result,

developers may utilize a familiar language while benefiting from Snowflake's scale and execution engine. <https://docs.snowflake.com/en/developer-guide/snowpark/index>

Domain

Extending Snowflake Functionality

Question 60Correct

Query processing in Snowflake is performed by compute engines called _____.

Spark Machines

Hadoop Clusters

Compute Nodes

Your answer is correct

Virtual Warehouses

Overall explanation

The compute engines in Snowflake are known as virtual warehouses.

<https://docs.snowflake.com/en/user-guide/warehouses-overview>

Domain

Architecture

Question 61Correct

True/False: As the size of the virtual warehouse increases, the amount of Snowflake credit consumed per second increases.

False

Your answer is correct

True

Overall explanation

The number of Snowflake credits consumed depends on the size of the virtual warehouse, the duration while it is in a running state, and how many virtual warehouses you have created. As the size of a virtual warehouse increases, the number of credits consumed per hour or second also increases.

<https://docs.snowflake.com/en/user-guide/warehouses-considerations#how-are-credits-charged-for-warehouses>

Domain

Architecture

Question 62Incorrect

True or False. You can use expressions in the clustering key definition to reduce a column's cardinality.

Correct answer

True

Your answer is incorrect

False

Overall explanation

If you need to use a column with high cardinality in a clustering key, you can define the key as an expression, which can help lower the cardinality by lowering the number of distinct values. For example, if a table has a timestamp column you want to use in a clustering key, you can convert the timestamp to a date to reduce the cardinality. This approach will help improve partition pruning.

<https://docs.snowflake.com/en/user-guide/tables-clustering-keys#strategies-for-selecting-clustering-keys>

Domain

Performance Concepts

Question 63Correct

A virtual warehouse is in a running state and is executing a complex query. The administrator increases the size of the virtual warehouse. What best describes the impact on the running query and any future queries? Select all that apply.

Your selection is correct

Only future queries can take advantage of the increased size.

Your selection is correct

There is no impact on the running query.

The current query immediately utilizes the additional nodes available after the resize operation.

The existing running queries are stopped and re-submitted after the virtual warehouse is resized.

Overall explanation

When a virtual warehouse is resized, any currently executing queries are not impacted—only new queries are affected by the new size.

<https://docs.snowflake.com/en/user-guide/warehouses-tasks#resizing-a-warehouse>

Domain

Architecture

Question 64Correct

What is the maximum period for Time Travel for Temporary tables?

7 days

Your answer is correct

1 day

21 days

14 days

Overall explanation

Transient and Temporary tables in Snowflake support Time Travel for up to 1 day, irrespective of the Snowflake edition used. <https://docs.snowflake.com/en/user-guide/tables-temp-transient>

Domain

Time Travel

Question 65Correct

True or False: When a security administrator creates a new custom role, the custom role is automatically assigned to all existing users.

True

Your answer is correct

False

Overall explanation

A new custom role is not automatically assigned to any user, but all new roles must be assigned manually to users. <https://docs.snowflake.com/en/user-guide/security-access-control-overview>

Domain

Security

Question 66Incorrect

Which of the following can be chosen when creating a new Snowflake account?

Select all that apply.

Your selection is correct

Region

Account Locator

Organization Name

Your selection is correct

Administrator user details

Correct selection

Account Name

Your selection is correct

Snowflake Edition

Overall explanation

Using the CREATE ACCOUNT statement, you can specify the account name, the Snowflake edition, the region (which contains the cloud platform information), the region group, and details about the administrative user, including name, password, email, etc.

<https://docs.snowflake.com/en/sql-reference/sql/create-account>

Domain

Licensing & Features

Question 67 **Correct**

Which information is displayed in the Statistics box in the Query Profile?

Select two answers.

Your selection is correct

Bytes spilled to remote storage

Your selection is correct

Bytes spilled to local storage

Query Steps

Operator tree

Overall explanation

Please see the link for details.

<https://docs.snowflake.com/en/user-guide/ui-query-profile#statistics>

Domain

Performance Concepts

Question 68 **Incorrect**

A Table UDF will return what type of result? Select all that apply

Your selection is correct

For each input, it can return multiple rows containing multiple columns

Your selection is correct

For each input, it can return one row containing a single column

Your selection is correct

For each input, it can return multiple rows containing a single columns

Your selection is correct

For each input, it can return one-row containing multiple columns

Correct selection

For each input, it can return zero rows

Overall explanation

All of these are valid answers. Table UDFs can return zero, one, or several rows for each input, with each result row containing one or multiple columns. UDTFs are another name for user-defined table functions. An example of such a function can be FLATTEN which returns several rows and columns for a single input.

<https://docs.snowflake.com/en/sql-reference/udf-overview#scalar-and-tabular-functions>

Domain

Extending Snowflake Functionality

Question 69Correct

Which of the following editions support user-defined functions?

Your selection is correct

Business Critical

Your selection is correct

Enterprise

Your selection is correct

Standard

Your selection is correct

Virtual Private Snowflake

Overall explanation

UDFs are a foundational feature and are supported by all Snowflake editions.

<https://docs.snowflake.com/en/user-guide/intro-editions.html>

Domain

Licensing & Features

Question 70Correct

A resource monitor can NOT control the cost of which of the following? Select all that apply.

Your selection is correct

Serverless compute

A virtual warehouse created by a user

Your selection is correct

Cloud Services costs

Overall explanation

Resource monitors cannot control the credit usage of serverless features such as Snowpipe, Automatic Reclustering, and Materialized View maintenance. So, resource monitors can only manage virtual warehouses created by the customer, and their resource monitors can only suspend virtual warehouses operated by the user. Any Snowflake-managed compute resource, such as Snowpipe, Serverless tasks, etc., can not be tracked or managed by resource monitors. In addition, a Resource Monitor can not control the costs of cloud services. A warehouse-level resource monitor can monitor

credit usage by Cloud Services, but the resource monitor can not suspend the cloud services. <https://docs.snowflake.com/en/user-guide/resource-monitors#assignment-of-resource-monitors>

Domain

Account Usage & Monitoring

Question 71 **Incorrect**

Which of the following statements is true regarding Resource Monitors in Snowflake?

Select all that apply.

Your selection is correct

Resource monitors are not meant to tightly control how credit much is used every hour.

Your selection is incorrect

Resource monitors can control credit usage to a high degree of precision.

Resource monitors can monitor credit usage on a per-minute basis.

Correct selection

Resource monitors are not meant to set highly precise limits on how much credit can be used.

Overall explanation

Resource monitors are not meant to precisely control how much is used every hour (or minute). Instead, they are meant to track and control how much is used daily, weekly, monthly, etc. They are not meant to set exact limits on how many credits can be used. When a resource monitor's credit limits are reached, it may take a while for the assigned warehouses to be put on hold, even if the action is "Suspend Immediately.", which may result in more credit than the credit limit to be used.

<https://docs.snowflake.com/en/user-guide/resource-monitors>

Domain

Account Usage & Monitoring

Question 72 **Correct**

How many Snowflake accounts are required if you need to create two Snowflake instances, one in Europe/Middle East region and one in Asia Pacific Region?

Zero

One

Your answer is correct

Two

Four

Overall explanation

You will need to create two different accounts, one for each region. Each Snowflake account is hosted in a particular Snowflake region. To use Snowflake in multiple regions, a Snowflake customer needs to maintain multiple Snowflake accounts, at least one for each region. <https://docs.snowflake.com/en/user-guide/intro-regions.html>

Domain

Licensing & Features

Question 73Correct

Which of the following schemas are automatically created with a new database?

Select all that apply.

PRIVATE

INTERNAL

Your selection is correct

INFORMATION_SCHEMA

Your selection is correct

PUBLIC

CATALOGUE

Overall explanation

The INFORMATION_SCHEMA provides metadata on the objects in the parent database of the INFORMATION_SCHEMA. It is automatically created with every database and can not be dropped, renamed, or moved.

The PUBLIC schema is also automatically created with every database, but it is just like an ordinary schema. It can be dropped, renamed, or moved. If required, additional schemas may be created under a database.

<https://docs.snowflake.com/en/sql-reference/info-schema#information-schema-views-and-table-functions>

Domain

Snowflake's catalog and objects

Question 74Incorrect

When a database or a schema is cloned, which of the following objects will NOT be cloned?

Permanent Tables

Transient Tables

Correct selection

Internal Snowflake Stages

Your selection is correct

External Tables

Your selection is correct

Snowpipes that point to a Named Internal Snowflake Stage

Overall explanation

Cloning for an external table is not supported. The data for external tables sits outside Snowflake; therefore, cloning an external table does not work.

Named Internal Stages cannot be cloned.

When a database or schema is cloned, any Snowpipe that points to a Named Internal Stage is not cloned.

<https://docs.snowflake.com/en/user-guide/object-clone>

Domain

Cloning

Question 75Correct

Which of the following layers is responsible for query optimization and dispatching?

Database Storage

Client Optimization Library

Your answer is correct

Cloud Services

Query Processing

Overall explanation

The cloud services layer is responsible for query planning and optimization. The actual query processing is performed by virtual warehouses.

Domain

Architecture

Question 76Correct

True or False: Combining Cloning and Time Travel to create clones of data at a certain point in time is possible.

False

Your answer is correct

True

Overall explanation

Combining Cloning and Time Travel can generate a clone of a table, database, or schema as it existed at a specific point in time. Because both Time Travel & Cloning are metadata operations, they can easily be combined.

<https://docs.snowflake.com/en/user-guide/data-time-travel#cloning-historical-objects>

Domain

Cloning

Question 77Correct

True or False: In the Snowsight worksheet view, you can share a worksheet with another user in the same Snowflake account.

False

Your answer is correct

True

Overall explanation

Snowsight lets you share worksheets and folders with other Snowflake users in your account, allowing others to view and execute SQL in your worksheets and folders.

<https://docs.snowflake.com/en/user-guide/ui-snowsight>

Domain

Tools & Interfaces

Question 78Correct

Which of the following query profile results indicate that a query used a virtual warehouse sized too small for the query?

Your selection is correct

The query profile shows a significant value for "Bytes spilled to local storage."

Your selection is correct

The query profile shows a significant value for "Bytes spilled to remote storage."

There are many Join nodes.

The query performs a full table scan, and no partition pruning occurs.

The Result node returns many rows.

Overall explanation

Snowflake saves data on the warehouse's local disk if it can't fit an operation into memory. Data spilling slows down queries because it requires more IO operations, and disk access is slower than memory access. "Bytes spilled to local storage." indicates local spillage. Snowflake will spill data to remote cloud storage if the local disk becomes full, which is even slower storage than the local disk, making this operation even slower. "Bytes spilled to remote storage" in the query profile indicates remote spillage. One of the ways to avoid spilling is to use a larger warehouse, which will increase the overall available RAM, local storage, and parallelism and might be able to fit the query in memory. <https://docs.snowflake.com/en/user-guide/ui-query-profile#queries-too-large-to-fit-in-memory>

Domain

Performance Concepts

Question 79**Incorrect**

Which of the following objects can be cloned? Select all that apply.

Your selection is correct

File Format

Your selection is correct

Stage

Virtual Warehouse

Correct selection

Sequence

Your selection is correct

Task

Share

Overall explanation

Virtual warehouses & Share objects cannot be cloned. Tables, Schemas & Databases can be cloned. Other objects that can be cloned include Stages, File Formats, Tasks, Sequences, and Streams. <https://docs.snowflake.com/en/user-guide/object-clone>

Domain

Cloning

Question 80**Correct**

Which of the following statements correctly describe Search Optimization in Snowflake? Select all that apply.

The search optimization service requires an active virtual warehouse.

Your selection is correct

The search optimization works in the background and updates the search access paths when the underlying table data changes.

Your selection is correct

The search optimization service is a Snowflake-managed service.

Overall explanation

The search optimization service is a Snowflake-managed service; it executes in the background and doesn't require a virtual warehouse. However, note that credit and storage costs are associated with Search optimization. When the data in the table is changed (for example, by loading new data sets or performing their DML operations), the maintenance service updates the search access path to reflect the changes. The search optimization configuration on a table and the maintenance service are transparent to the users. <https://docs.snowflake.com/en/user-guide/search-optimization-service>

Domain

Performance Concepts

Question 81Correct

Which of the following can be configured for a user profile in Snowsight?

Select two answers.

Your selection is correct

Notifications

Tri Secret Secure

Single Signon

Your selection is correct

Multi-factor Authentication

Overall explanation

Using the profile dialogue in Snowsight, you can enroll in MFA, specify your notification preferences, set your email address, and configure your profile's default role and default warehouse. Other things can also be configured such as name, password, language, etc.

<https://docs.snowflake.com/en/user-guide/ui-snowsight-profile>

Domain

Tools & Interfaces

Question 82Correct

Which of the following objects can be shared? Select all that apply.

Your selection is correct

Secure UDFs

Your selection is correct

Secure Materialized Views

Views

Materialized Views

Your selection is correct

Tables

Your selection is correct

Secure Views

Overall explanation

Standard views cannot be shared. Direct data sharing enables sharing of the following types of objects: Tables, External tables, Secure views, Secure materialized views, Secure UDFs. <https://docs.snowflake.com/en/user-guide/data-sharing-intro>

Domain

Data Sharing

Question 83Correct

True or False: Storage costs are billed to the customer for data stored in Time Travel and fail-safe storage.

False

Your answer is correct

True

Overall explanation

Snowflake charges for Time Travel and failsafe data storage. The cost of maintaining data for Time Travel and fail-safe is calculated every 24 hours based on the number of days it is maintained and the time since it was last modified.

<https://docs.snowflake.com/en/user-guide/data-cdp-storage-costs>

Domain

Time Travel

Question 84Correct

True or False: An external function can only be written in C++.

Your answer is correct

False

True

Overall explanation

One of the benefits of creating an external function is that it can be written in languages that standard UDFs cannot. Many languages can be used to create external functions since they exist outside of Snowflake. <https://docs.snowflake.com/en/sql-reference/external-functions-introduction#advantages-of-external-functions>

Domain

Extending Snowflake Functionality

Question 85Correct

If the query result cache for a query is not used by any future query, what is the duration after which it will be purged?

365 days

3600 seconds

Your answer is correct

24 hours

31 days

Overall explanation

Once a result cache is generated for a query stays valid for 24 hours. If another query that reuses the query result cache is executed within that 24-hour window, the result cache expiry is extended for another 24 hours from that point onwards. If the result cache for a query keeps getting used, it will stay valid for up to 31 days. After 31 days, the result cache for a query will be purged regardless of any other condition.

<https://docs.snowflake.com/en/user-guide/querying-persisted-results>

Domain

Performance Concepts

Question 86Incorrect

True or False: A resource monitor can track cloud services costs associated with a virtual warehouse but can not stop cloud services costs from occurring.

Your answer is incorrect

False

Correct answer

True

Overall explanation

A resource monitor can not control the costs of cloud services. A warehouse-level resource monitor can monitor credit usage by Cloud Services, but the resource monitor can not suspend the cloud services. <https://docs.snowflake.com/en/user-guide/resource-monitors#assignment-of-resource-monitors>

Domain

Account Usage & Monitoring

Question 87Correct

True or False: Snowflake allows single sign-on (SSO) with a SAML 2.0-compliant external identity provider for federated authentication (IdP).

False

Your answer is correct

True

Overall explanation

Snowflake supports federated authentication, allowing for single sign-on (SSO). Users authenticate using a SAML 2.0-compliant external identity provider (IdP). After IdP authentication, users can access Snowflake without logging in. Snowflake natively supports the majority of SAML 2.0 compliant identity providers, including Okta, ADFS, OneLogin, and Ping Identity PingOne. <https://docs.snowflake.com/en/user-guide/admin-security-fed-auth>

Domain

Security

Question 88Incorrect

Which of the following transformation is supported by the COPY command?

Correct answer

CAST

AVG

SUM

Your answer is incorrect

FLATTEN

Overall explanation

When loading data into a table using the COPY command, Snowflake allows you to do simple transformations on the data as it is being loaded. During the load process, the COPY command allows for modifying the order of columns, omitting one or more columns, casting data into specified data types, and truncating values. While loading the data, complex transformations such as joins, filters, aggregations, and the use of FLATTEN are not supported as they are not essential data transformations. Therefore, joining, filtering, and aggregating the data are supported ONLY after the data has been loaded into a table. <https://docs.snowflake.com/en/user-guide/data-load-overview#id2>

Domain

Data Loading and Unloading

Question 89Correct

Which of the following correctly describes the behavior of a materialized view when its base table is dropped?

Select two answers.

The materialized view is automatically dropped.

Your selection is correct

The materialized view is suspended.

Your selection is correct

The materialized view is NOT dropped automatically.

A table with a materialized view on top can not be dropped. The materialized view must be dropped first.

Overall explanation

If the base table of a materialized view is dropped, the materialized view is NOT automatically dropped; however, it is suspended. You must manually drop the materialized view.

<https://docs.snowflake.com/en/user-guide/views-materialized#dropping-the-base-table>

Domain

Performance Concepts

Question 90Correct

Database and table cloning can be used to satisfy which of the following scenarios?

Select all that apply.

Improve JOIN performance.

Improve performance for SELECT queries.

Your selection is correct

Create point-in-time backups by cloning with Time Travel.

Your selection is correct

Create backups of critical databases.

Your selection is correct

Quickly create non-prod databases by creating clones of production objects.

Load data incrementally.

Overall explanation

Cloning in Snowflake lets you quickly spin up development and testing environments by creating copies of databases, schemas, or tables.

Snowflake's Time Travel feature allows you to access and clone data at specific points in time within a retention period. By leveraging Time Travel with cloning, you can create testing environments that reflect specific scenarios or historical data states.

Time Travel automatically provides some back and recovery capability; however, it is limited to the maximum data retention period, which can only be up to 90 days. Therefore, performing regular backups in Snowflake is still a requirement. One way of backing up data in Snowflake is to export data from individual tables to cloud storage; however, you can also utilize the cloning capability in Snowflake to perform backups of critical databases.

Domain

Cloning

Question 91Incorrect

Which of the following statements best describe Snowflake's multi-factor authentication (MFA) capability? Select all that apply

Only Snowflake's Web Interface and SnowSQL allow multi-factor authentication.

Your selection is correct

An administrator can turn off a user's MFA.

Correct selection

Multi-factor authentication (MFA) is enabled by default for all users; however, users must enroll themselves in MFA manually.

Your selection is correct

SnowSQL, Snowflake ODBC, JDBC, and the Python Connector are all MFA-compatible.

Overall explanation

Multi-factor authentication (MFA) is enabled by default for all Snowflake accounts, and any Snowflake user can enroll themselves in MFA through the Snowflake web interface. An administrator can disable a user's MFA enrolment; in this case, the user must re-enroll to access the MFA features and functionality. An administrator with the SECURITYADMIN or above role can disable MFA for a user. All Snowflake client tools, including the web interface, SnowSQL, and the various connectors and drivers, support MFA. <https://docs.snowflake.com/en/user-guide/security-mfa>

Domain

Security

Question 92Correct

Which one of the following is true regarding Snowflake Marketplace?

A Snowflake account can only consume data from the Snowflake Marketplace.

A separate Snowflake account is needed to publish data to the Snowflake Marketplace.

Only Snowflake can publish new data sets to Snowflake Marketplace.

Your answer is correct

Any Snowflake account can consume or publish data to the Snowflake Marketplace.

Overall explanation

Except for Virtual private Snowflake accounts, the Snowflake Marketplace is available to all Snowflake accounts hosted on Amazon Web Services, Google Cloud Platform, and

Microsoft Azure. Any Snowflake account (again, except for VPS accounts) can become a data provider and publish datasets to the Marketplace for a cost or for free. In addition, you are required to sign up as a partner first and become an approved data provider. <https://other-docs.snowflake.com/en/collaboration/collaboration-marketplace-about.html>

Domain

Data Sharing

Question 93 **Correct**

You are the performance DBA at a large airline company with a Snowflake Data warehouse. A large table (>5TB) contains telemetry data from airplane sensors.

The table is usually accessed by the "event_date" on which the data was generated, but often there are queries that access the table through the column "airplane_id".

How can you optimize the table so that queries which use either event_date or the airplane_id in the WHERE clause run faster?

Do nothing; Snowflake will ensure the efficiency of querying itself.

Your answer is correct

Introduce a cluster key on the combination of the two columns i.e.

CLUSTER BY(event_date,airplane_id)

Concatenate the two columns and generate a new column in the table. Cluster the table on that column

Increase the size of the virtual warehouse so that queries run faster.

Overall explanation

Clustering a table on a specific column can optimize queries by eliminating unnecessary partitions from the query processing. A table can be re-clustered by defining a clustering key, which effectively redistributes the data into micro-partitions according to the clustering key, ensuring optimal access to queries that predicate or join on the clustered column. Introducing a cluster key on both the event_date & airplane_id columns creates an optimal partitioning scheme for access via either column.

<https://docs.snowflake.com/en/user-guide/tables-clustering-keys>

Domain

Performance Concepts

Question 94 **Incorrect**

Which of the following data types can optimally store semi-structured data?

Select three answers.

Your selection is incorrect

VARCHAR

CHAR

Your selection is correct

VARIANT

Your selection is incorrect

STRING

Correct selection

OBJECT

Correct selection

ARRAY

Overall explanation

OBJECT, ARRAY and VARIANT are suitable data types for storing and processing semi-structured data.

<https://docs.snowflake.com/en/sql-reference/data-types-semistructured>

Domain

Data Transformation

Question 95**Incorrect**

In a VARIANT column, the NULL values are stored as a literal string "null"

Correct answer

True

Your answer is incorrect

False

Overall explanation

NULL values are stored as a literal string "null" in a VARIANT column.

<https://docs.snowflake.com/en/user-guide/semistructured-considerations#null-values>

Domain

Data Loading and Unloading

Question 96**Correct**

PowerBI, Looker, MicroStrategy, Tableau, Qlik are what type of partners in the Snowflake partner ecosystem?

Machine Learning & Data Science

Your answer is correct

Business Intelligence

Data Integration

Security, Governance & Observability

Overall explanation

All of these are Business Intelligence partners of Snowflake. Please see

<https://docs.Snowflake.com/en/user-guide/ecosystem.html>

Domain

Partners

Question 97**Incorrect**

An administrator must grant explicit privileges on a cloned database as the cloning does not copy the privileges of the source database.

Correct answer

True

Your answer is incorrect

False

Overall explanation

A cloned object does not inherit any privileges from its source object; for instance, a cloned table does not inherit any privileges from its source table.

However, if a database or schema is cloned, privileges are inherited by the child objects in the database. The database doesn't inherit any privileges.

<https://docs.snowflake.com/en/user-guide/object-clone#access-control-privileges-for-cloned-objects>

Domain

Cloning

Question 98Correct

The PUT command can be used in which of the following scenarios?

Load data from a Snowflake internal stage into a Snowflake table.

Load data from a Snowflake external stage into a Snowflake table.

Load data from the Snowflake Web UI.

Your answer is correct

Transfer data from an on-premises system to a Snowflake stage.

Overall explanation

The PUT command uploads data from an on-premises system to an internal stage. The GET command is used to download data from an internal stage to an on-premises system. To download or upload data to an external stage, cloud provider utilities or other tools are used to interact with data in the cloud storage pointed to by the external stage. <https://docs.snowflake.com/en/user-guide/data-unload-overview#bulk-unloading-process>

Domain

Data Loading and Unloading

Question 99Correct

True or False: Using the views in the INFORMATION_SCHEMA, you can access the history of usage that occurred 5 minutes ago.

Your answer is correct

True

False

Overall explanation

The data provided via the INFORMATION_SCHEMA views is real-time, and there is no latency in the information provided. So, if you are asked which schema should be used if there is a requirement to view real-time data, then the views in INFORMATION_SCHEMA should be used as they contain real-time information.

<https://docs.snowflake.com/en/sql-reference/account-usage#differences-between-account-usage-and-information-schema>

Domain

Account Usage & Monitoring

Question 100Correct

Time Travel can help with which of the following scenarios?

(Assume a time travel duration of 90 days).

Your selection is correct

A new data pipeline rolled out in production yesterday has deleted all rows from a production table.

A data corruption issue that corrupted three production tables 93 days ago was discovered.

Your selection is correct

An administrator accidentally dropped a production table last week.

Overall explanation

You can use Time Travel AT and BEFORE extensions to recover all deleted rows from the production table.

You can use the UNDROP statement, which can be used to recover tables, schemas, or even complete databases after they have been dropped.

Using Time Travel, you can not recover from a data issue before the maximum time travel period, i.e., 90 days. You must request Snowflake support to recover data from the fail-safe storage.

<https://docs.snowflake.com/en/user-guide/data-time-travel#time-travel-sql-extensions>

Domain

Time Travel

Question 101Correct

You are unloading data from a table into multiple files. Which COPY command parameter will let you control the maximum size of each file?

Your answer is correct

MAX_FILE_SIZE

MAX_FILE_BYTES

FILE_SIZE_LIMIT

Overall explanation

When data is unloaded from Snowflake, it is automatically compressed using gzip compression. This is the default behavior; however, you can specify alternate compression methods or turn off compression entirely. The unloading process automatically exports to multiple files so that it can take advantage of the parallelism offered by Snowflake. However, if needed, you can set the SINGLE parameter to true to ensure the export goes to a single file. The default size of each output file is 16 MB but can be changed using the MAX_FILE_SIZE parameter. The maximum allowed size per file is 5GB if you export data to cloud storage. <https://docs.snowflake.com/en/user-guide/data-unload-considerations#unloading-to-a-single-file>

Domain

Data Loading and Unloading

Question 102Correct

True or False: The data provider is responsible for and pays for the compute costs of a reader account.

Your answer is correct

True

False

Overall explanation

Since the data provider creates and administers the reader account, all the reader account's compute expenses are invoiced to the provider account. Therefore, the reader account's use of the virtual warehouse compute is added to the provider account compute charges.

Domain

Data Sharing

Question 103**Correct**

Which of the following is true regarding Query Result Cache for a given query? Select all that apply.

Your selection is correct

If a query reuses the query result cache, it is extended by another 24 hours.

Your selection is correct

If no new query reuses the query result cache, the cache is purged after 24 hours.

Your selection is correct

The query result cache may be extended for up to a maximum of 31 days

The query result cache can be manually purged.

Overall explanation

The query result cache for a query has an initial validity period of twenty-four hours. The cache is purged if a new query doesn't reuse the previously generated cache within 24 hours. If a new query uses the result cache, the validity period for the query result cache is reset to another 24 hours. It is now valid for another 24 hours from when it was reused. This extension of the first query result cache can continue for up to a maximum of 31 days from the point in time when a query result cache was initially produced.

After 31 days, the query result cache for a query is purged altogether.

<https://docs.snowflake.com/en/user-guide/querying-persisted-results>

Domain

Performance Concepts

Question 104**Correct**

As an alternative to data loading, Snowflake allows the creation of external tables through which you can query external data without loading it first.

False

Your answer is correct

True

Overall explanation

Snowflake provides external tables, which permit the creation of tables with data stored in external cloud storage. Using external tables, you can query data without loading it into Snowflake. <https://docs.snowflake.com/en/user-guide/data-load-overview#external-tables-data-lake>

Domain

Data Loading and Unloading

Question 105**Correct**

Which Snowflake edition has a dedicated metadata storage, dedicated pool of compute resources and supports protected health information (PHI) data according to HIPAA and HITRUST CSF regulations?

Business Critical Edition

Enterprise Edition

Your answer is correct

Virtual Private Snowflake (VPS) edition

Standard Edition

Overall explanation

The VPS edition is meant to provide isolation from other customers; thus, each instance has its own metadata store and compute resources. It also supports HIPPA & HITRUST CSF regulations.

<https://docs.snowflake.com/en/user-guide/intro-editions.html>

Domain

Licensing & Features

Question 106**Incorrect**

Which of the following statements are true when a new trial account is created for a partner using Partner Connect?

Select all that apply.

Your selection is correct

A new user is created.

Your selection is incorrect

A new reader account is created.

Correct selection

A virtual warehouse is created.

Correct selection

A new empty database is created.

Data sharing is disabled.

Your selection is correct

A new custom role is created.

Overall explanation

While connecting to a Partner application, Snowflake automatically creates several objects, such as an empty database, virtual warehouse, default user, and custom role. When the partner app reads or writes to your account, it uses these objects.

<https://docs.snowflake.com/en/user-guide/ecosystem-partner-connect#connecting-with-a-snowflake-partner>

Domain

Partners

Question 107Correct

Which of the following statements correctly describe a share?

Your selection is correct

A share object has information about the consumer(s)

Share objects are available only to enterprise customers.

Your selection is correct

A share object is a container that contains objects to be shared

Overall explanation

A share acts as a container for objects that need to be shared. A share object encapsulates the database & the schema to be shared, the tables and secure views which will be shared, and the consumer account(s) to which the Share will be available.

Sharing and consequently share objects are available to all Snowflake editions

Domain

Data Sharing

Question 108Incorrect

Which performance problems can easily be identified through the query profile?

Select two answers.

Correct selection

A UNION without the ALL clause

Too many fields in the GROUP BY clause

Incorrect use of Windowing functions

Your selection is correct

Exploding joins

Your selection is incorrect

Too many fields in the SELECT clause

Overall explanation

Typical performance issues you can identify using the query profile include Exploding joins, UNION without ALL, queries that don't fit in the memory, and partition pruning issues.

<https://docs.snowflake.com/en/user-guide/ui-query-profile#common-query-problems-identified-by-query-profile>

Domain

Performance Concepts

Question 109Correct

True or False: A cloned object does not contribute to overall storage until DML operations on the source or target object are done.

False

Your answer is correct

True

Overall explanation

When tables, schemas, or databases are cloned, the cloned item does not contribute to total storage until data manipulation language (DML) operations are performed on the source or target, which modify or delete existing data or add additional data.

<https://docs.snowflake.com/en/user-guide/tables-storage-considerations#label-cloning-tables>

Domain

Cloning

Question 110Correct

When a multi-cluster virtual warehouse is suspended, which one of the following caches will be purged?

Metadata Cache

Your answer is correct

Warehouse Cache (local disk cache)

Query Result Cache

Browser Cache

Overall explanation

Every time a virtual warehouse accesses data from a table, it caches that data locally. This data cache can improve the performance of subsequent queries if those queries can reuse the data in the cache instead of reading from the table in the cloud storage. Reading from a local cache is a much more efficient operation than reading data from the cloud storage; therefore, it improves performance for queries that can take advantage of it.

The warehouse cache is purged if the virtual warehouse is suspended. When the virtual house is resumed, the warehouse cache is rebuilt over time as queries are processed.

<https://docs.snowflake.com/en/user-guide/warehouses-considerations#how-does-warehouse-caching-impact-queries>

Domain

Performance Concepts

Question 111Correct

Which of the following ACCOUNT_USAGE view can be used to view roles granted to users?

GRANTS_TO_ROLES

ACCESS_HISTORY

Your answer is correct

GRANTS_TO_USERS

OBJECT_DEPENDENCIES

Overall explanation

The GRANTS_TO_USERS view can be used to view information about roles that have been granted to users. This view also contains historical information (up to 365 days), so roles that have been granted and revoked in the last 365 days will also be shown.

https://docs.snowflake.com/en/sql-reference/account-usage/grants_to_users

Domain

Account Usage & Monitoring

Question 112Correct

Which of the following are serverless features provided by Snowflake?

Multi-cluster virtual warehouses

Virtual warehouses

Your selection is correct

Automatic Clustering

Your selection is correct

Snowpipe

Overall explanation

Snowpipe is a serverless service for continuous data loads. Automatic clustering is a serverless Snowflake-managed service that maintains micro-partitions for tables according to the defined clustering key.

Domain

Data Loading and Unloading

Question 113Correct

Which of the following statements about a multi-cluster virtual warehouse in auto-scale mode is true?

Your selection is correct

When Snowflake detects that queries are beginning to queue, additional warehouses are started.

Your selection is correct

Multi-cluster virtual warehouses support the same properties and actions as regular virtual warehouses. They can be suspended, resumed, and modified to a different size.

Multi-cluster virtual warehouses do not support Suspend, Resume, or Resize operations.

Your selection is correct

When the query demand decreases, warehouses are progressively shut down.

Overall explanation

When you start a multi-cluster virtual warehouse in auto-scaling mode, the number of active virtual warehouses equals the minimum warehouse count. Snowflake creates new warehouses based on demand, up to the maximum warehouse count. Snowflake shuts down virtual warehouses when demand decreases until the number equals the minimum warehouse count. Multi-cluster virtual warehouses support the standard virtual warehouse attributes and actions, such as specifying and altering warehouse size, suspending or auto-suspending a suspended multi-cluster virtual warehouse, resuming or auto-resuming a suspended multi-cluster virtual warehouse.

<https://docs.snowflake.com/en/user-guide/warehouses-multicluster>

Domain

Performance Concepts

Question 114Correct

Which of the following does Snowflake Scripting support?

Your selection is correct

Branching constructs such as IF-ELSE and CASE statements

Your selection is correct

Looping Constructs such as FOR, WHILE, REPEAT, and LOOP.

Your selection is correct

Declaring & using variables

Your selection is correct

Handling exceptions

Overall explanation

All of these are supported in Snowflake Scripting. Snowflake Scripting allows you to use variables, if-else expressions, looping, cursors, manage result sets, and allows you to handle errors. Snowflake scripting is typically used to create stored procedures, but it may also be used to create procedural code outside of a stored procedure.

<https://docs.snowflake.com/en/developer-guide/snowflake-scripting/index>

Domain

Extending Snowflake Functionality

Question 115Correct

You have a medium-sized virtual warehouse for executing ad-hoc queries. Due to many queries concurrently executing at peak times, the queries start queuing. What is the best way to reduce queuing?

Request users to spread their queries throughout the day.

Your answer is correct

Configure the virtual warehouse to be a multi-cluster virtual warehouse. The virtual warehouse scales out and back depending on the query demands.

Create several Medium-sized virtual warehouses and redirect users to use different virtual warehouses.

Kill all long-running queries

Overall explanation

Multi-cluster virtual warehouses are frequently used in scenarios where the number of concurrent queries exceeds the capacity of a single virtual warehouse. When a virtual warehouse's concurrent workload exceeds its maximum capacity, additional queries are placed in the queue. Multi-cluster virtual warehouses dynamically add additional clusters based on demand to solve the queueing issue. When demand decreases, the additional clusters are decommissioned. This process is also known as scaling out or auto-scaling. <https://docs.snowflake.com/en/user-guide/warehouses-multicuster>

Domain

Performance Concepts

Question 116Incorrect

What value does the CURRENT_CLIENT function return when called from an application using a JDBC driver to connect?

Correct answer

The version of the JDBC driver

The string 'JDBC.'

Your answer is incorrect

The name of the application

The string 'External Application.'

Overall explanation

The CURRENT_CLIENT function returns the client's version from where the query was executed. When called from a query executed by an application using JDBC or ODBC driver, the version of the driver is returned. For example, Calling CURRENT_CLIENT in Snowsight returns 'Go 1.1.5' since it uses the Go Driver. CURRENT_CLIENT in Classic Web UI returns 'Snowflake UI 20230324175929' since the Classic Web UI connects directly without a driver. Calling CURRENT C_CLIENT in SnowSQL returns 'SnowSQL 1.2.24'. https://docs.snowflake.com/en/sql-reference/functions/current_client

Domain

Tools & Interfaces

Question 117Incorrect

Which of the following actions occur when a directory table's metadata is refreshed?

Select all that apply.

Correct selection

Files that have been deleted are removed from the directory table's metadata.

Your selection is correct

The directory table's metadata reflects changes to the file paths.

Your selection is correct

Files that are new in the external stage are added to the directory table's metadata.

Overall explanation

All of these are valid answers.

A directory table's metadata can be manually refreshed or configured to auto-refresh. When the metadata is refreshed, it synchronizes it with the current state of files in the stage. It will synchronize files that are new in the external stage are added to the directory table's metadata, and deleted files are removed from the directory table's metadata. Additionally, any changes to the file paths are reflected in the directory table's metadata.

<https://docs.snowflake.com/en/user-guide/data-load-dirtables-intro#automatically-refreshing-directory-table-metadata>

Domain

Data Transformation

Question 118Correct

What is the impact on the existing & new queries when a virtual warehouse is scaled up or down? Select all that apply

Existing queries will be impacted, and will start using the new size immediately.

Your selection is correct

Existing queries are not impacted.

Your selection is correct

The changed size will impact only new queries.

Overall explanation

Only new queries are affected by the changed size; existing queries on the virtual warehouse remain unaffected. When a virtual warehouse is scaled down, nodes are removed from the virtual warehouse only when there is no active query on the virtual warehouse. Thus, existing queries are not impacted by resizing.

<https://docs.snowflake.com/en/user-guide/warehouses-considerations#warehouse-resizing-improves-performance>

Domain

Performance Concepts

Question 119Incorrect

Which of the following roles are available out of the box in Snowflake?

Your selection is correct

ORGADMIN

Your selection is correct

PUBLIC

Correct selection

SYSADMIN

SUPERUSER

Overall explanation

Built-in Snowflake roles include ORGADMIN, ACCOUNTADMIN, USERADMIN, SECURITYADMIN, SYSADMIN, and PUBLIC. <https://docs.snowflake.com/en/user-guide/security-access-control-overview#system-defined-roles>.

Domain

Security

Question 120Correct

What is the minimum Snowflake built-in role that can be used to enforce a network policy on a Snowflake account?

ACCOUNTADMIN

Your answer is correct

SECURITYADMIN

ORGADMIN

SYSADMIN

Overall explanation

The SECURITYADMIN role is the minimum Snowflake built-in role allowed to enforce a network policy at an account level.

<https://docs.snowflake.com/en/sql-reference/sql/desc-user#usage-notes>

Domain

Security

Question 121Correct

True or False: The data in the views in the ACCOUNT_USAGE schema is real-time.

Your answer is correct

False

True

Overall explanation

The ACCOUNT_USAGE schema consists of several views that provide usage metrics and metadata information at the account level. Data provided by the ACCOUNT_USAGE views is NOT real-time and refreshes typically with a lag of 45 minutes to 3 hours, depending on the view. The data in these views are retained for up to 365 days.

<https://docs.snowflake.com/en/sql-reference/account-usage#differences-between-account-usage-and-information-schema>

Domain

Account Usage & Monitoring

Question 122Correct

True or False: When clustering keys are defined for a table, there is a cost for maintaining and re-organizing micro-partitions according to the clustering keys.

Your answer is correct

True

False

Overall explanation

For tables with a clustering key defined, Automatic Clustering, a Snowflake service, re-clusters the micro-partitions as needed, distributing data according to the clustering key to achieve appropriate partition pruning. Snowflake internally maintains the clustered tables and any resource requirements with Automatic Clustering. Automatic Clustering only adjusts those micro-partitions which benefit from the re-clustering process. Automatic Clustering does not need a virtual warehouse but uses Snowflake-managed CPU, RAM, etc. Therefore, it has a cost attached, which should appear under serverless costs. Clustering a table uses credits like any other data modification (DML) action in Snowflake. Re-clustering also adds extra storage when data is physically redistributed and new micro-partitions are created. The original micro-partitions are kept for Time Travel and Fail-safe purposes, resulting in increased storage.
<https://docs.snowflake.com/en/user-guide/tables-auto-reclustering#credit-usage-and-warehouses-for-automatic-clustering>

Domain

Performance Concepts

Question 123Correct

A permanent table can be cloned to which other table types?

Select all that apply.

Your selection is correct

Temporary Table

Your selection is correct

Transient Table

External Table

Your selection is correct

Permanent Table

Overall explanation

A permanent table may be cloned into another permanent table, a transient or temporary table.

However, a transient or temporary table can not be cloned into a permanent table.

Domain

Cloning

Question 124Correct

True or False: Snowflake customers can retrieve data from fail-safe storage using Time Travel SQL extensions.

True

Your answer is correct

False

Overall explanation

Once the data is in fail-safe storage, only Snowflake support can help retrieve the data. The customer cannot access fail-safe storage.

Domain

Fail-safe

Question 125Incorrect

Which of the following approaches will give you the credit used by each virtual warehouse for the last 12 months?

Query the `METERING_DAILY_HISTORY` view in the `ACCOUNT_USAGE` schema

Your answer is incorrect

Query the `METERING_HISTORY` view in the `ACCOUNT_USAGE` schema

Request Snowflake support to provide you with a detailed virtual warehouse credit usage report.

Correct answer

Query the `WAREHOUSE_METERING_HISTORY` view in the `ACCOUNT_USAGE` schema

Overall explanation

The `WAREHOUSE_METERING_HISTORY` view in the `ACCOUNT_USAGE` schema provides credit usage for each virtual warehouse in your system, broken down by hour. The view provides information from the last 365 days.

https://docs.snowflake.com/en/sql-reference/account-usage/warehouse_metering_history

Domain

Account Usage & Monitoring

Question 126Correct

The `COPY` command can load data from which of the following?

Tape drives

Your selection is correct

External Stages

Detachable hard disks

Your selection is correct

Internal Stages

Overall explanation

Similar to how data warehouses use staging, Snowflake uses a Stage object. Snowflake uses stages to aid in the loading and unloading of data. The data must first be available in a Snowflake stage to load data into a Snowflake table. `COPY` command can be used to load data into a table after the data is loaded in a stage. Data unloading or exporting is also performed via a Stage object; the data can only be extracted to a stage, internal or external. <https://docs.snowflake.com/en/user-guide/data-load-overview>

Domain

Data Loading and Unloading