Practice Exam 6 - Results

Question 1 Skipped

In the Query Profile, what does the TableScan operator represent?

Correct answer

Access to a single table

Combining two inputs on a given condition

Access to data stored in stage objects

Adding records to a table

Overall explanation

The TableScan operator represents access to a single table. Please see

https://docs.snowflake.com/en/user-guide/ui-query-profile#data-access-and-generation-operators

Domain

Performance Concepts

Question 2 Skipped

What privilege is required to resize a virtual warehouse?

ALTER

MONITOR

USAGE

Correct answer

MODIFY

Overall explanation

The MODIFY privilege allows a role to alter the size of a virtual warehouse.

https://docs.snowflake.com/en/sql-reference/sql/alter-warehouse#access-control-requirements

Domain

Security

Question 3 Skipped

What of the following is true about the objects created by Partner Connect during the process of connecting to a partner?

Select all that apply.

Correct selection

The partner uses the PC_<partner>_USER to connect to your Snowflake instance.

Correct selection

The PC_<partner>_ROLE allows the partner application to access objects granted to the PUBLIC role.

The PC_<partner>_DB database is only created if there is no other database in your system.

Correct selection

The PC_<partner>_DB database is empty and can be used to load or store data for querying.

The PC_<partner>_WH virtual warehouse is only created if there is no other virtual warehouse in your system.

Overall explanation

The PC_<partner>_DB is created empty. However, you can configure it to use an existing database if required.

The PC_<partner>_ROLE inherits the PUBLIC role privileges. You can grant additional privileges to this role to allow the partner application to access objects in your system. This role is also granted to SYSADMIN to ensure that SYSADMINs can access the objects and data created by the partner application.

https://docs.snowflake.com/en/user-guide/ecosystem-partner-connect#connecting-with-a-snowflake-partner

Domain

Partners

Question 4 Skipped

An administrator needs to grant access to a role so that they can create a materialized view in the database MARKETING and schema PUBLIC. Which statement will provide the required privileges to the role?

GRANT MATERIALIZED VIEW ON SCHEMA MARKETING.PUBLIC TO <role_name>;

Correct answer

GRANT CREATE MATERIALIZED VIEW ON SCHEMA MARKETING.PUBLIC TO ROLE <role_name>;

GRANT CREATE MATERIALIZED VIEW ON SCHEMA MARKETING.PUBLIC TO <role_name>;

GRANT MATERIALIZED VIEW ON SCHEMA MARKETING.PUBLIC TO ROLE <role_name>;

Overall explanation

The correct syntax is GRANT CREATE MATERIALIZED VIEW ON SCHEMA <schema_name> TO ROLE <role_name>;

https://docs.snowflake.com/en/user-guide/views-materialized#privileges-on-a-materialized-view-s-schema

Domain

Security

Question 5 Skipped

Which of the following can be fulfilled through the ACCESS_HISTORY view in the ACCOUNT_USAGE schema?

Select two answers.

Identify which roles were used by the logged-in user.

Correct selection

Help identify data that is unused and is not being queried.

Correct selection

Identify what data was accessed, when, and who accessed it.

Identify who logged into the system.

Overall explanation

Using the ACCESS_HISTORY view, you can identify what data was accessed, when, and who accessed it. Using this information, you can also identify what data is not being accessed at all.

There are other benefits of using ACCESS_HISTORY data, which can be found at the following link.

https://docs.snowflake.com/en/user-guide/access-history#benefits

Domain

Account Usage & Monitoring

Question 6 **Skipped**

True or False: To create and execute a stored procedure, you must have a minimum of Enterprise edition.

True

Correct answer

False

Overall explanation

All Snowflake editions stored procedures, so the minimum edition required to create and execute a stored procedure is the Standard edition.

https://docs.snowflake.com/en/user-guide/intro-editions.html

Domain

Licensing & Features

Question 7 **Skipped**

Which of the following can be used to find the query ID of the most recent query executed in the current session?

Select two answers.

Correct selection
SELECT LAST_QUERY_ID();
SELECT LAST_QUERY_ID(-2);
Correct selection
SELECT LAST_QUERY_ID(-1);
SELECT LAST_QUERY_ID(2);
SELECT LAST_QUERY_ID(1);
Overall explanation
The LAST_QUERY_ID function returns the query ID of a specified query in the current session. The function takes a number as the parameter, which specifies the position of the query in the session.

The parameter can take positive or negative values. A negative value means you are attempting to fetch the most recent query in the session, where

-1 = most recent query

-2 = 2nd most recent query

, and so on. The function defaults to -1, so if no value is provided, it will return the query id of the most recent query.

A positive number returns the earliest queries in the session. i.e.

1 = first query

2 = 2nd query

https://docs.snowflake.com/en/sql-reference/functions/last_query_id

Domain
General
Question 8Skipped
Which function is used to determine the fully qualified URL and port for Snowight when configuring Snowight for access through a proxy or a firewall?
SYSTEM$REFERENCE
SYSTEM$GET_TAG
SYSTEM$GENERATE_SAML_CSR
Correct answer
SYSTEM$ALLOWLIST
Overall explanation
You need to add the fully qualified URL and port values to the proxy servers or firewall settings to use a proxy or firewall to connect to Snowsight.

Use the SNOWSIGHT_DEPLOYMENT item in the return value of the SYSTEM$ALLOWLIST function to find out the fully qualified Snowsight URL and port.

https://docs.snowflake.com/en/user-guide/ui-snowsight-gs#accessing-sf-web-interface-through-a-proxy-or-firewall

Domain
Security

Question 9 Skipped
When data is shared between Snowflake accounts, what type of database is created on the consumer side for consuming the shared data?

Permanent
Correct answer
Read-only
Writable
Temporary
Open Access

Overall explanation
The correct answer is read-only. The consumer creates a database from the Share object as a read-only database. https://docs.snowflake.com/en/user-guide/data-sharing-intro#how-does-secure-data-sharing-work

Domain
Data Sharing

Question 10 Skipped
What role is required to execute the following command successfully?

SHOW ORGANIZATION ACCOUNTS;

Correct answer
ORGADMIN
ACCOUNTADMIN
SECURITYADMIN
SYSADMIN

Overall explanation
Only users with the ORGADMIN role can execute the "SHOW ORGANIZATION ACCOUNTS" command.

https://docs.snowflake.com/en/sql-reference/sql/show-organization-accounts

Domain
Security

Question 11 Skipped
The compute layer in Snowflake architecture performs which one of the following?

Query Planning
Query Optimization
Correct answer
Query Processing
Cache Query Results
Overall explanation
The compute layer is responsible for query processing or query execution. Query planning & optimization is performed by the cloud services layer. The query result cache is stored and managed in the cloud services layer.
https://docs.snowflake.com/en/user-guide/intro-key-concepts
Domain
Architecture
Question 12 Skipped
Which of the following best describes “Bytes spilled to local storage” shown in a query profile?
Correct answer
“Bytes spilled to local storage” indicates the volume of data that couldn’t fit in memory and had to be spilled to a virtual warehouse temporary storage.
“Bytes spilled to local storage” indicates the amount of data downloaded using Snowsight.
“Bytes spilled to local storage” is the number of micro-partitions pruned during query execution.
“Bytes spilled to local storage” indicates the amount of data downloaded using the GET command.
Overall explanation
Snowflake saves data on the warehouse's local disk if it can't fit an operation into memory. Data spilling slows down queries because it requires more IO operations, and disk access is slower than memory access. "Bytes spilled to local storage." indicates local spillage. Snowflake will spill data to remote cloud storage if the local disk becomes full, which is even slower storage than the local disk, making this operation even slower. "Bytes spilled to remote storage" in the query profile indicates remote spillage.
https://docs.snowflake.com/en/user-guide/ui-query-profile#queries-too-large-to-fit-in-memory
Domain
Performance Concepts
Question 13 Skipped
Which of the following contributes towards the storage costs in Snowflake?

Select all that apply.

Correct selection
Permanent Tables
Correct selection
Temporary Tables
Warehouse Cache
Correct selection
Transient Tables
Query Result Cache
Overall explanation

Data stored in permanent tables counts towards the storage costs.

Data stored in temporary & transient tables also contribute towards the storage costs until they are dropped or data is cleared.

Data in Fail-safe storage and Time Travel storage also contribute to the storage costs.

Transient and temporary tables, however, do not contribute towards Fail-safe storage costs and have a maximum of 1-day Time Travel costs.

Caching is NOT considered for determining storage costs. The query result cache & metadata cache are part of the cloud services layer.

The warehouse cache (local disk cache) is part of a virtual warehouse and does NOT contribute to storage costs.

https://docs.snowflake.com/en/user-guide/cost-understanding-overall

Domain
Cost & Pricing

Question 14 Skipped

A user ran a complex query that took 2 hours to complete on a medium size virtual warehouse. The user re-ran the query 6 hours later in a new session. The query returned results almost immediately. What is the reason for this quick execution?

The results were retrieved from the browser cache.

The results were retrieved from the user stage; therefore, the execution was fast.

Correct answer

The results were returned immediately by using the query result cache.

Data in the underlying tables has been deleted. The query ran fast because it had nothing to process.

Overall explanation

When Snowflake runs a query, it caches the results of that query for a predetermined amount of time. The stored query results are referred to as the Query Result Cache. The Query Result Cache can be used to fulfill future queries if they are similar to a previously executed query & there have been no changes to the data in the tables being queried. When Snowflake returns query results using a query result cache, the procedure is exceptionally quick since it does not include the actual execution of the query. Because the query is not being executed, there is no need for any virtual warehouses, reducing computing costs. https://docs.snowflake.com/en/user-guide/querying-persisted-results

Domain
Performance Concepts
Question 15 Skipped
In the Snowsight worksheet view, which of the following can be selected to set the context under which a query executes?
User
Correct selection
Database
Correct selection
Role
Correct selection
Schema
Table
Correct selection
Virtual Warehouse
Overall explanation
You can choose the primary role under which the query is executed, and the virtual warehouse used to run the query. You can also choose the database and schema to which the worksheet view defaults, so you don't need to prefix tables in the specified database & schema. https://docs.snowflake.com/en/user-guide/ui-snowsight
Domain
Tools & Interfaces
Question 16 Skipped
The query history in the Snowflake Web Interface is kept for how many days?
30 days
365 days
Correct answer
14 days
6 months
Overall explanation
The query history page lets users view the history of executed and currently executing queries. The query history page can show the history of queries executed in the last 14 days. https://docs.snowflake.com/en/user-guide/ui-snowsight-activity#query-history
Domain
Account Usage & Monitoring
Question 17 Skipped
Which of the following transformations are supported by the COPY command? Select all that apply.
Correct selection
Omit columns
Join with other tables
Correct selection
Reorder columns
Correct selection
Cast columns
Overall explanation
When loading data into a table using the COPY command, Snowflake allows you to do simple transformations on the data as it is being loaded. During the load process, the COPY command allows for modifying the order of columns, omitting one or more

columns, casting data into specified data types, and truncating values. While loading the data, complex transformations such as joins, filters, aggregations, and the use of FLATTEN are not supported as they are not essential data transformations. Therefore, joining, filtering, and aggregating the data are supported ONLY after the data has been loaded into a table. https://docs.snowflake.com/en/user-guide/data-load-overview#id2

Domain

Data Loading and Unloading

Question 18 Skipped

Which of the following is true regarding privileges and roles in Snowflake? Select all that apply

Privileges can be assigned to users directly.

Correct selection

Roles can be assigned to users.

Correct selection

Privileges can be assigned to roles.

New custom roles cannot be created.

Overall explanation

Snowflake's access control is built on the role-based access control (RBAC) approach, which assigns rights to roles and roles to users. The privileges given to a role are inherited by all users in that role. https://docs.snowflake.com/en/user-guide/security-access-control-overview

Domain

Security

Question 19 Skipped

Which of the following are valid statements when loading data into a table (called EMPLOYEE) using the table stage?

Select two answers.

COPY INTO EMPLOYEE FROM TABLE_STAGE;

COPY INTO EMPLOYEE SELECT * FROM TABLE_STAGE;

Correct selection

COPY INTO EMPLOYEE;

Correct selection

COPY INTO EMPLOYEE FROM @%EMPLOYEE;

Overall explanation

The FROM clause can be omitted when loading data from a table stage. In such a case, Snowflake automatically assumes data is being loaded from the table stage.

So both **COPY INTO EMPLOYEE;** and **COPY INTO EMPLOYEE FROM @%EMPLOYEE;** are correct.

https://docs.snowflake.com/en/user-guide/data-load-local-file-system-copy#table-stage

Domain
Data Loading and Unloading

Question 20 Skipped

Which of the following can view and modify an existing resource monitor? Select all that apply.

A system administrator (i.e., a person with the SYSADMIN role).

Correct selection

A user who has MONITOR and MODIFY privilege on the resource monitor.

Any user of the system

Correct selection

An account administrator (i.e., a person with the ACCOUNTADMIN role)

Overall explanation

From a privilege perspective, only Account Administrators (users with ACCOUNTADMIN role) can create resource monitors. However, account administrators can grant privileges to the resource monitor to allow other users to view and modify the resource monitor configuration. The MONITOR and MODIFY privileges on a resource monitor allow other users to view and modify a specific resource monitor.
https://docs.snowflake.com/en/user-guide/resource-monitors#access-control-privileges-for-resource-monitors

Domain
Account Usage & Monitoring

Question 21 Skipped

Which of the following terms refers to the same layer in Snowflake architecture? Select all that apply.

Virtual Warehouses

Correct selection

Cloud Services Layer

Storage Layer

Correct selection

Services Layer

Overall explanation

The terms Cloud Services Layer or Services Layer are used interchangeably.

https://docs.snowflake.com/en/user-guide/intro-key-concepts#cloud-services

Domain
Architecture

Question 22 Skipped

Which view type doesn't use some of the internal Snowflake optimizations?

Permanent Views

Correct answer

Secure views

External Views

Materialized Views

Overall explanation
For typical views, internal optimizations can indirectly expose data to users.

Secure views hide the underlying data by removing some of the internal Snowflake optimizations.

https://docs.snowflake.com/en/user-guide/views-secure

Domain
Security
Question 23 Skipped
What term is used for a pre-computed dataset obtained from a SELECT query specification and stored for future use?
Secure View
View
Output View
Correct answer
Materialized View
Overall explanation
A materialized view is a view that pre-computes data based on a SELECT query. The query's results are pre-computed and physically stored to enhance performance for similar queries that are executed in the future. When the underlying table is updated, the materialized view refreshes automatically, requiring no additional maintenance. Snowflake-managed services perform the update in the background transparent to the user without interfering with the user's experience.

https://docs.snowflake.com/en/user-guide/views-materialized

Domain
Performance Concepts
Question 24 Skipped
Which of the following correctly describes a reader account in Snowflake?
A reader account is required to access external tables.
A reader account is exclusively used for testing reasons.
A reader account is used to distribute query costs between departments.
Correct answer
A reader account can be used to share data with non-Snowflake users.
Overall explanation
Sharing data with a non-Snowflake user or organization is possible by creating a reader account. This reader account is created by the data provider solely for sharing purposes.
https://docs.snowflake.com/en/user-guide/data-sharing-reader-create
Domain
Data Sharing
Question 25 Skipped

Which statements are correct regarding the costs of using event notifications to refresh a directory table's metadata?

Select all that apply.

Correct selection
An additional cost is charged for the event notifications.
The refresh operation is free.
The event notifications are free.
Correct selection
A small maintenance cost is charged for the refresh operation.
Overall explanation
A small maintenance cost is charged for refreshing a directory table's metadata, whether through notifications or manually (through ALTER STAGE <stage-name> REFRESH). This small maintenance cost is accounted for under the cloud services costs.

Additionally, when using cloud platform notifications, an additional cost is charged, which appears as Snowpipe charges in your billing statement. The Snowpipe cost is charged because Snowpipe is used for event notifications to trigger the automatic refresh.

https://docs.snowflake.com/en/user-guide/data-load-dirtables-intro#billing-for-directory-tables

Domain
Data Transformation
Question 26Skipped
True or False: The data in the views in the INFORMATION_SCHEMA is real-time.
Correct answer
True
False
Overall explanation
The data provided via the INFORMATION_SCHEMA views is real-time, and there is no latency in the information provided. So, if you are asked which schema should be used if there is a requirement to view real-time data, then the views in INFORMATION SCHEMA should be used as they contain real-time information.
https://docs.snowflake.com/en/sql-reference/account-usage#differences-between-account-usage-and-information-schema
Domain
Account Usage & Monitoring
Question 27Skipped
Which of the following statements are true regarding Snowflake's architecture? Select all that apply

Snowflake uses a monolithic architecture in which compute and storage are tightly coupled.

Correct selection

Snowflake storage & compute are independent of each other.

In Snowflake, the compute must be increased whenever the storage is increased.

Correct selection

You can increase or decrease the compute resources in Snowflake without changing the storage.

Overall explanation

Snowflake implements a new hybrid architecture that decouples compute and storage. Snowflake stores data similarly to a shared-disk architecture, i.e., the data is shared. But it also allows for using several compute engines on the same shared data, each with its own memory and processing capabilities. This hybrid architecture allows Snowflake to increase or decrease compute without requiring storage changes and vice versa.

Domain

Architecture

Question 28 **Skipped**

True or False: Network policies can be used to allow or deny access to specific IP addresses.

Correct answer

True

False

Overall explanation

Administrators can configure the system to allow or deny access to specific IP addresses through network policies. A network policy consists of the policy name, a comma-separated list of allowed IP addresses, and a list of blocked IP addresses
https://docs.snowflake.com/en/user-guide/network-policies

Domain

Security

Question 29 **Skipped**

The cloud services layer in Snowflake provides which of the following functions? Select all that apply.

Correct selection

Transaction control / ACID compliance

Correct selection

Cloning

Storage for data

Correct selection

Data Sharing

Overall explanation

Snowflake's data sharing, cloning, and data exchange features are all managed through the cloud services layer using metadata. The cloud services layer also provides ACID compliance. ACID means a database system must allow several transactions to run in isolation and commit or roll back a transaction as a unit, assuring system consistency.

Domain

Architecture

Question 30 **Skipped**

For an unpopulated table, the clustering depth is _____?

-1

1000
1
Correct answer
Zero
Overall explanation
For a populated table, the clustering depth is the average depth of overlapping micro-partitions for specific columns. The clustering depth starts at 1 (for a well-clustered table) and can be a larger number. For an unpopulated table, the clustering depth is zero. https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions#label-clustering-depth
Domain
Performance Concepts
Question 31 Skipped
Which of the following can create a new share? Select all that apply.
Correct selection
A user with a role having the CREATE SHARE privileges
A user with the SYSADMIN role
A user with the SECURITY ADMIN role
Correct selection
A user with the ACCOUNTADMIN role
Overall explanation
Only users with ACCOUNTADMIN roles or with CREATE SHARE permission can create a share. https://docs.snowflake.com/en/user-guide/data-sharing-gs
Domain
Data Sharing
Question 32 Skipped
Which of the following statements are true regarding Reader Accounts in Snowflake?

Select all that apply.

A reader account can raise support requests.
A single reader account can access data shared by different providers.
Correct selection
A reader account can only access data shared by the provider account that created the reader account.
Correct selection
Any support requests for the reader account must be raised through the provider account.
Overall explanation
Reader accounts do not have a Snowflake licensing agreement, so it does not have access to support. Instead, the provider account (that created the reader account) manages the support requests.

Only data from the provider account that created the reader account can be consumed by the reader account.

https://docs.snowflake.com/en/user-guide/data-sharing-reader-create

Domain
Data Sharing

Question 33 Skipped
How many days of historical data can you access through the views in the ACCOUNT_USAGE schema?

Correct answer
365 days
90 days
1 day
7 days

Overall explanation
The ACCOUNT USAGE schema consists of several views that provide usage metrics and metadata information at the account level. Data provided by the ACCOUNT_USAGE views is NOT real-time and refreshes typically with a lag of 45 minutes to 3 hours, depending on the view. The data in these views are retained for up to 365 days.
https://docs.snowflake.com/en/sql-reference/account-usage#differences-between-account-usage-and-information-schema

Domain
Account Usage & Monitoring

Question 34 Skipped
Which of the following requirement is NOT a good reason to create a Stored Procedure?

Correct answer
Calculate the average of two values
Execute a series of SQL statements to create a new user and a new database for that user.
Execute a series of SQLs to clean up shared temporary data every day.
Find and drop tables that have not been used in a long time

Overall explanation
Even though stored procedures and UDFs look similar, they serve different purposes. The job of a UDF is to take in input, perform computations, and return a value, whereas the job of a stored procedure is to execute one or more SQL queries. Thus in the given options calculating the average of two values is best fulfilled by a UDF.
https://docs.snowflake.com/en/sql-reference/stored-procedures-overview

Domain
Extending Snowflake Functionality

Question 35 Skipped
True or False: When a Snowflake data provider shares data with a data consumer, the data consumer is not charged for any additional storage costs.

False

Correct answer
True

Overall explanation
Metadata operations in the cloud services layer allow data sharing without physically copying it. Since the provider account stores and pays for the data storage, the data consumer doesn't have to pay anything extra for storage. However, the data consumer pays for the compute used to run queries on shared data. When queries are run on shared data, the compute of the data consumer is used.

https://docs.snowflake.com/en/user-guide/data-sharing-intro#how-does-secure-data-sharing-work

Domain

Data Sharing

Question 36 Skipped

Which of the following semi-structured file formats are supported by Snowflake? Select all that apply

Correct selection

ORC

YAML

PDF

Correct selection

JSON

Correct selection

AVRO

Overall explanation

Snowflake includes built-in support for several semi-structured data formats. Snowflake supports JSON Avro ORC Parquet XML https://docs.snowflake.com/en/user-guide/semistructured-intro.html

Domain

Data Loading and Unloading

Question 37 Skipped

Which of the following statements are correct regarding Time Travel & fail-safe storage? Select all that apply.

The maximum allowed Time Travel duration for a temporary table is 7 days.

Correct selection

There is no fail-safe storage for a temporary table.

A temporary table has 7 days of fail-safe storage.

Correct selection

The maximum allowed Time Travel duration for a temporary table is 1 day.

Overall explanation

Transient and temporary tables don't have fail-safe functionality; therefore, data in such tables goes through zero days of fail-safe storage. Also, Transient and Temporary tables have a maximum of 1 day of Time Travel. https://docs.snowflake.com/en/user-guide/tables-temp-transient

Domain

Time Travel

Question 38 Skipped

Which of the following activities are not required to be performed by a Snowflake customer?

Select all that apply.

Correct selection

Ensure the physical security of a data center.

Correct selection

Installation of Snowflake database Software

Correct selection

Provision hardware for installing the Snowflake database

Management of user access & privileges

Correct selection

Configuration and Testing of High availability of hardware at the data center level

Overall explanation

Snowflake, a software-as-a-service product, doesn't require a customer to manage the data center and its physical security, hardware install hardware or software or manage high availability.

Domain

Licensing & Features

Question 39 **Skipped**

Which of the following statements regarding a Stream object in Snowflake are true? (select all that apply)

A Stream object can keep track of DML changes for a table for several years.

Explanation

A Stream object monitors and records the DML changes to a source table within the specified data retention period. After that, the stream becomes stale, and the DML changes are not available. Snowflake extends table data retention to prevent stale streams up to the maximum defined extension defaulting to 14 and maximum up to 90.

See the following links for more details: https://docs.snowflake.com/en/user-guide/streams-intro#data-retention-period-and-staleness and https://docs.snowflake.com/en/sql-reference/parameters#max-data-extension-time-in-days

Correct selection

Materialized views do NOT support Streams.

Explanation

Streams do not support materialized view currently.

Correct selection

A stream object records DML changes (inserts, updates, and deletes) made to a table at the row level.

Explanation

Snowflake Streams help you keep track of any changes made to a table, such as new data being added (inserts), existing data being modified (updates), or data being removed (deletes). They allow you to query and process only the changed data since the last offset.

See the link for more details: https://docs.snowflake.com/en/user-guide/streams-intro

External tables do NOT support Streams.

Explanation

Streams support external tables. Insert-only streams in Snowflake track changes to external tables by capturing both inserts and updates as new rows, but they ignore delete operations. This means they only reflect additions and modifications to data,

treating updates as new inserts.
See the link for more details: https://docs.snowflake.com/en/user-guide/streams-intro#types-of-streams

Domain
Streams

Question 40 Skipped
Through Snowflake sharing, a data provider can share data with which of the following? Select all that apply.

One Drive Users

Correct selection

A non-Snowflake customer

Google Drive users

Correct selection

Multiple Snowflake customers

Correct selection

Another Snowflake customer

Overall explanation
You can share data with multiple consumers: Snowflake customers, non-Snowflake customers, or a mix of both.

Domain
Data Sharing

Question 41 Skipped
A virtual warehouse was started, used for 5 minutes and 15 seconds, and shut down afterward. The customer will be charged for how many seconds?

3600 seconds

360 seconds

0 seconds

Correct answer

315 seconds

Overall explanation
Snowflake credits are billed on a per-second usage basis, which means if a virtual warehouse ran for 5 minutes and 15 seconds, you would be charged for 315 seconds (5*60 + 15). However, note that a minimum of 60 seconds of billing applies, so if a virtual warehouse were started and shut down within the first 1st minute, a minimum of 60-second credit usage would apply.

Domain
Architecture

Question 42 Skipped
Which of the following statements are true regarding scaling up / down and scaling out virtual warehouses?

Select all that apply.

Scaling up & scaling down is managed automatically by Snowflake.

Scaling out & scaling back a virtual warehouse is a manual process.

Correct selection

Scaling out & scaling back a virtual warehouse is performed automatically by Snowflake.

Correct selection

Scaling up & scaling down a virtual warehouse is a manual process.

Overall explanation

Multi-cluster virtual warehouses dynamically (and automatically) add additional clusters based on demand to solve the queueing issue. When demand decreases, the additional clusters are decommissioned. This process is also known as scaling out (and scaling back) or auto-scaling.

Scaling up and down is a manual process, requiring someone to run a statement to increase or decrease the size of the virtual warehouse.

https://docs.snowflake.com/en/user-guide/warehouses-multicluster

Domain

Performance Concepts

Question 43 Skipped

True or False: Snowflake Scripting be used to create stored procedures.

Correct answer

True

False

Overall explanation

Snowflake scripting is typically used to create stored procedures, but it may also be used to create procedural code outside of a stored procedure.
https://docs.snowflake.com/en/sql-reference/stored-procedures-snowflake-scripting

Domain

Extending Snowflake Functionality

Question 44 Skipped

A large table in Snowflake may contain millions or hundreds of millions of micro-partitions.

Correct answer

True

False

Overall explanation

The number of micro-partitions for a given table depends mainly on the amount of data in that table. For a very large table, the number of micro-partitions can run into millions or hundreds of millions of micro-partitions. https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions

Domain

Architecture

Question 45 Skipped

Which of the following is true regarding Query Profile? Select all that apply.

Correct selection

It shows a graphical representation of all steps.

Correct selection

A query profile is available for all queries, whether running, completed, or failed.

Correct selection

It shows the query plan for a query.

A query profile is not available for queries that have not been completed.

Overall explanation

Query Profile provides query execution details. It displays a graphical representation of the main components of the processing plan for the specified query, as well as statistics for each component and overall query information and statistics. Query Profile is available for all queries, whether running, completed, or failed. Query Profile is a valuable tool for learning how queries work. It can be used if you want or need to know more about how a query executes. It is designed to assist you in identifying typical errors in SQL query expressions so that you may identify potential performance bottlenecks and devise strategies to improve. https://docs.snowflake.com/en/user-guide/ui-query-profile

Domain

Performance Concepts

Question 46 **Skipped**

The GET command is used for which of the following purposes?

Download data from an external stage on an on-premises system.

Download data from an internal stage to the cloud storage.

Correct answer

Download data from an internal stage to an on-premises system.

Download data from a Snowflake table to any type of stage.

Overall explanation

The GET command is used to download data from an internal stage to an on-premises system. The PUT command uploads data from an on-premises system to an internal stage. To download or upload data to an external stage, cloud provider utilities or other tools are used to interact with data in the cloud storage pointed to by the external stage. https://docs.snowflake.com/en/user-guide/data-unload-overview#bulk-unloading-process

Question 47 **Skipped**

Which of the following multi-cluster virtual warehouse configurations indicate a multi-cluster virtual warehouse running in auto-scale mode?

MIN_CLUSTER_COUNT = 3

MAX_CLUSTER_COUNT = 3

Correct selection

MIN_CLUSTER_COUNT = 1

MAX_CLUSTER_COUNT = 5

MIN_CLUSTER_COUNT = 1

MAX_CLUSTER_COUNT = 1

MIN_CLUSTER_COUNT = 0

MAX_CLUSTER_COUNT = 0

Correct selection

MIN_CLUSTER_COUNT = 2

MAX_CLUSTER_COUNT = 4

Overall explanation
A multi-cluster virtual warehouse can be created in maximized or auto-scaling modes. Auto-Scaling mode is enabled by selecting different values for the multi-minimum clusters and maximum warehouse count. As a result, Snowflake starts and stops warehouses dynamically based on the workload needs.

The maximized mode is enabled by setting the minimum and maximum warehouse count of the multi-cluster to the same value. Therefore, as soon as the multi-cluster virtual warehouse is established, all warehouses in the multi-cluster are started up.

https://docs.snowflake.com/en/user-guide/warehouses-multicluster#setting-the-scaling-policy-for-a-multi-cluster-warehouse

Domain
Performance Concepts

Question 48 Skipped

As the database administrator, you defined new clustering keys for a large table. So, while Snowflake re-clusters data, what should you expect? Select all that apply.

Correct selection

SELECT queries continue to execute as normal while Snowflake redistributes data in micro-partitions.

DML queries are disallowed while Snowflake redistributes data in micro-partitions.

Correct selection

DML queries continue to execute as normal while Snowflake redistributes data in micro-partitions.

SELECT queries are blocked from execution while Snowflake redistributes data in micro-partitions.

Overall explanation
Snowflake's re-clustering operation is transparent to the user and does not block any DML or SELECT queries. A table that is being re-clustered will behave exactly like any other table when being queried, updated, or changed.
https://docs.snowflake.com/en/user-guide/tables-auto-reclustering#non-blocking-dml

Domain
Performance Concepts

Question 49 Skipped

True or False: A stored procedure must return a value.

True

Correct answer

False

Overall explanation

A stored procedure can also return a single value or tabular data if desired; however, it is not a requirement that a stored procedure must return a value.
https://docs.snowflake.com/en/developer-guide/stored-procedures-vs-udfs
Domain
Extending Snowflake Functionality
Question 50 Skipped
You want to load a specific list of files from an S3 stage. What is the correct syntax for achieving this?

Assume the file names are

delta1.csv

delta2.csv

delta3.csv

and the stage is called my_stage, and the table is called my_table

COPY INTO my_table FROM @my_stage/delta1.csv,@my_stage/delta2.csv,@my_stage/delta3.csv
COPY delta1.csv,delta2.csv,delta3.csv INTO my_table FROM @my_stage
Correct answer
COPY INTO my_table FROM @my_stage files=('delta1.csv', 'delta2.csv', 'delta3.csv')
Overall explanation
You can specify the exact file names in the COPY command so that only those files are accessed and loaded. The syntax for that can be found on the link below.
https://docs.snowflake.com/en/user-guide/data-load-considerations-load#lists-of-files
Domain
Data Loading and Unloading
Question 51 Skipped
Which of the following statements regarding multi-cluster virtual warehouses are true? Select all that apply.
Correct selection
Multi-cluster virtual warehouses automatically remove additional clusters when query demand decreases.
Correct selection
A minimum of Enterprise edition is required for multi-cluster virtual warehouse capability.
Correct selection
Multi-cluster virtual warehouses automatically add additional clusters when simultaneous queries increase to a number that existing virtual warehouses can not handle.
Multi-cluster virtual warehouses cannot be set to auto-suspend or auto-resume.
Overall explanation
Multi-cluster virtual warehouses are utilized when the number of concurrent users exceeds a single virtual warehouse's capacity. When the concurrent workload for a virtual warehouse reaches the maximum, new queries are queued. Multi-cluster virtual warehouses address this by adding clusters as needed. When the demand drops, the

extra clusters are removed. Enterprise edition is required to use the multi-cluster virtual warehouse feature. Besides the automatic addition and removal of compute clusters, multi-cluster virtual warehouses behave the same as typical virtual warehouses so that they can be suspended or resumed and auto-suspended or auto-resumed. https://docs.snowflake.com/en/user-guide/warehouses-multicluster

Domain

Architecture

Question 52 Skipped

True/False: Snowflake is based on existing database technology, which has been retrofitted to run on the cloud.

True

Correct answer

False

Overall explanation

Snowflake has been designed for the cloud and has been designed from scratch. Snowflake implements a new hybrid architecture that decouples compute and storage.

Domain

Architecture

Question 53 Skipped

Snowflake allows which two of the following approaches for loading data?

Correct selection

Bulk Data Loading

Correct selection

Continuous Data Loading

Intermittent Data Loading

Overall explanation

Snowflake supports data loading in two primary ways. The COPY command can be used to load bulk data or huge files. To load data into a table, the COPY command requires the usage of a virtual warehouse. The other method of loading data into Snowflake is via the Snowpipe. Snowpipe is the ideal technique for loading data when the data is arriving continuously in a messaging or streaming manner.
https://docs.snowflake.com/en/user-guide/data-load-overview#bulk-vs-continuous-loading

Domain

Data Loading and Unloading

Question 54 Skipped

True or False: When setting up replication for cross-cloud or cross-region data sharing, the data provider must replicate data once for each data consumer.

True

Correct answer

False

Overall explanation

Only one instance of data per cloud or region must be replicated. Once the instance is replicated, more than one consumer can use this data.
https://docs.snowflake.com/en/user-guide/secure-data-sharing-across-regions-plaforms https://docs.snowflake.com/en/user-guide/secure-data-sharing-across-regions-plaforms#data-sharing-considerations

Domain

Data Sharing

Question 55 Skipped

Which of the following is true regarding Snowflake Marketplace?

Correct selection

The Snowflake Marketplace is an online marketplace where you can purchase and sell datasets.

All data sets on the Snowflake Marketplace are provided for a cost.

Correct selection

Using the Snowflake Marketplace, a customer can import data from outside your company into their Snowflake instance and utilize it to enrich their data.

Using the Snowflake Marketplace, A customer can bid on different data sets; only the highest bidder can access the data set.

Overall explanation

Snowflake Marketplace is a marketplace for finding and accessing third-party datasets made accessible by various organizations. These third-party datasets are generally supplied for a fee but are occasionally made accessible for free. There is no bidding, and the data sets are available to everyone (free or for a cost). https://other-docs.snowflake.com/en/collaboration/collaboration-marketplace-about.html

Domain

Data Sharing

Question 56 Skipped

Which of the following is a scenario where using Snowpipe is the optimal choice?

You are loading large files that arrive monthly

Your users are running massive data science workloads.

Correct answer

You need to load small volumes of frequently arriving data

Your users are running frequent tactical queries.

Overall explanation

Snowpipe is the ideal method for loading data when the data is arriving continuously in a messaging or streaming manner, and there is a requirement for data to be loaded almost immediately. https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro

Domain

Data Loading and Unloading

Question 57 Skipped

Snowflake supports which of the following authentication mechanisms? Select all that apply

Correct selection

Key Pair Authentication

MD5 authentication

Correct selection

Multi-factor authentication

Plain Text Password authentication

Overall explanation

Multi-factor authentication adds additional protection to the login process in Snowflake. Snowflake provides key pair authentication as a more secure alternative to the traditional username/password authentication approach. Additionally, Snowflake offers federated authentication, enabling users to access their accounts via a single sign-on (SSO). Users authenticate using SAML 2.0-compliant single sign-on (SSO) via an external identity provider (IdP).

Domain
Security
Question 58 Skipped
Which mechanism allow a Snowflake customer to query data without loading it first?
Snowpipe
Virtual Table
Correct answer
External Table
COPY
Overall explanation
Snowflake offers an alternative approach for tables called external tables, which permits the creation of tables with data stored in external cloud storage. External tables remove the need for the data to be loaded into Snowflake. In the case of an External table, the definition of the table is still stored in Snowflake metadata and consists of table structure, file locations, filenames, and other attributes. However, the table's data is saved outside of Snowflake. The external table functionality enables you to query external data like a standard table. External tables may be joined to other tables, and views may be created using them. https://docs.snowflake.com/en/user-guide/tables-external-intro
Domain
Data Loading and Unloading
Question 59 Skipped
Which of the following aspects should you consider when defining a clustering key for a large table?
Cluster all numeric columns
Correct selection
Cluster columns that are frequently used in join statements
Cluster all character columns
Correct selection
Cluster columns that are used frequently in WHERE clauses
Overall explanation
When defining clustering keys, the initial candidate clustering columns are those columns that are frequently used in the WHERE clause or other selective filters. Additionally, columns that are used for joining can also be considered. Furthermore, the columns' cardinality (number of distinct values) is also important. It is crucial to choose a column with a high enough cardinality to allow effective partition pruning while having a low enough cardinality for Snowflake to group data into micro-partitions efficiently. A column with too few distinct values (e.g., gender) will result in minimal partition pruning. On the other hand, a column that has too many distinct values (e.g., customer id) will result in too much overhead when maintaining the partitions. When creating a multi-column cluster key, order the columns from the lowest cardinality to the higher cardinality; otherwise, the effectiveness of clustering will be reduced. https://docs.snowflake.com/en/user-guide/tables-clustering-keys
Domain
Performance Concepts
Question 60 Skipped
Once created, which of the following cannot be converted to any other table type?

Select two answers.

Correct selection
Transient
Correct selection
Temporary
Permanent
Overall explanation
Once created, temporary and transient tables cannot be changed into any other table type.

https://docs.snowflake.com/en/user-guide/tables-temp-transient#creating-a-temporary-table

https://docs.snowflake.com/en/user-guide/tables-temp-transient#creating-a-transient-table-schema-or-database

Domain
Snowflake's Catalog and objects
Question 61 **Skipped**
True or False: To prevent performance issues, an error is returned if a query using INFORMATION_SCHEMA is not sufficiently selective.
Correct answer
True
False
Overall explanation
If the filters supplied in an INFORMATION SCHEMA query are not sufficiently selective, the following error is returned. Information schema query returned too much data. Please repeat the query with more selective predicates.
https://docs.snowflake.com/en/sql-reference/info-schema#general-usage-notes
Domain
Account Usage & Monitoring
Question 62 **Skipped**
Multi-factor authentication can be enabled for which of the following? Select all that apply.
Correct selection
Python
Correct selection
Snowflake WebUI
Snowpipe
Correct selection
ODBC
Overall explanation

MFA is enabled by default for all Snowflake accounts and is available in all Snowflake editions. All Snowflake client tools, including the web interface, SnowSQL, and the various connectors and drivers, support MFA. Snowpipe is a snowflake-managed serverless service. A Snowflake user can not log into it; therefore, it doesn't require MFA. https://docs.snowflake.com/en/user-guide/security-mfa

Domain

Security

Question 63 Skipped

True or False: Snowflake recommends a maximum of 3 or 4 columns in a clustering key.

Correct answer

True

False

Overall explanation

Snowflake recommends using a maximum of 3 or 4 columns in a clustering key. Any more columns in the clustering key result in more maintenance costs and do not provide enough benefits to justify the clustering costs.
https://docs.snowflake.com/en/user-guide/tables-clustering-keys

Domain

Performance Concepts

Question 64 Skipped

Snowflake stores table data in which format?

JSON

CSV & TSV files

Correct answer

A proprietary format

Parquet format

Overall explanation

Snowflake stores data in a proprietary format on cloud object storage, such as AWS S3, Azure Blob Storage, or Google Cloud Storage. Users cannot see the actual files, or look at how the data is stored, or access the file directly.

Domain

Architecture

Question 65 Skipped

When a database or a schema is cloned, which of the following statements are regarding Snowpipes in that database?

Correct selection

Any Snowpipes that reference an internal stage are NOT cloned

Correct selection

Any Snowpipes that reference an external stage are cloned

Any Snowpipes that reference an internal stage are cloned

Any Snowpipes that reference an external stage are NOT cloned

Overall explanation

When a database or schema is cloned, any Snowpipe that points to a Named Internal Stage is not cloned. Snowpipe referencing an external stage is cloned.
https://docs.snowflake.com/en/user-guide/object-clone#cloning-and-pipes

Domain

Cloning

Question 66 Skipped

What layer in Snowflake's architecture is responsible for user authentication and authorization?

Correct answer

Cloud Services

Database Storage

Query Processing

Client Tools

Overall explanation

The cloud services layer manages authentication and authorization. When a user logs in, the cloud services layer validates their credentials. When a user submits a query, the cloud services layer parses and optimizes the query plan.

https://docs.snowflake.com/en/user-guide/intro-key-concepts

Domain

Architecture

Question 67 **Skipped**

Which of the following correctly describes how materialized views are refreshed?

Additional SQL statements need to be scheduled to refresh a materialized view.

Correct selection

A materialized view is automatically updated if the data in the underlying table is changed.

A materialized view can be set to auto-refresh using the REFRESH_ON_BASE_TABLE_UPDATE parameter.

Materialized can ONLY be refreshed manually.

Correct selection

Materialized views are automatically refreshed by Snowflake managed service.

Overall explanation

When the underlying table is updated, the materialized view refreshes automatically, requiring no additional maintenance. Snowflake-managed services perform the update in the background transparent to the user without interfering with the user's experience.

https://docs.snowflake.com/en/user-guide/views-materialized

Domain

Performance Concepts

Question 68 **Skipped**

True or False: Snowflake can transform data after a partner software has loaded it.

False

Correct answer

True

Overall explanation

After data has been loaded into Snowflake through partner software, Snowflake SQL or other mechanisms can transform data within Snowflake.

Domain

Data Loading and Unloading

Question 69 **Skipped**

Which of the following are the correct privileges to allow a role (named MKT_USERS) to add, configure, or remove Search Optimization for a table called CUSTOMER in a schema called MARKETING?

Correct answer

GRANT ADD SEARCH OPTIMIZATION ON SCHEMA MARKETING TO ROLE MKT_USERS;

GRANT SEARCH OPTIMIZATION ON SCHEMA MARKETING TO ROLE MKT_USERS;
GRANT ADD SEARCH OPTIMIZATION ON TABLE CUSTOMER TO ROLE MKT_USERS;
GRANT SEARCH OPTIMIZATION ON TABLE CUSTOMER TO ROLE MKT_USERS;

Overall explanation

To add, configure, or remove search optimization for a table, you must have

a) OWNERSHIP privileges on the table.

b) ADD SEARCH optimization privileges on the schema that contains the table. The syntax is GRANT ADD SEARCH OPTIMIZATION ON SCHEMA <schema_name> TO ROLE <role>

https://docs.snowflake.com/en/user-guide/search-optimization-service#what-access-control-privileges-are-needed-for-the-search-optimization-service

Domain

Security

Question 70 **Skipped**

Suppose a multi-cluster virtual warehouse is configured with a minimum cluster count of 1 and maximum cluster count of 3. Is the virtual warehouse in Maximized mode?

Yes

Correct answer

No

Overall explanation

A multi-cluster virtual warehouse can be created in maximized or auto-scaling modes. The maximized mode is enabled by setting the minimum and maximum warehouse count of the multi-cluster to the same value. Therefore, as soon as the multi-cluster virtual warehouse is established, all warehouses in the multi-cluster are started up. Auto-Scaling mode is enabled by selecting different values for the multi-minimum clusters and maximum warehouse count. As a result, Snowflake starts and stops warehouses dynamically based on the workload needs.
https://docs.snowflake.com/en/user-guide/warehouses-multicluster#setting-the-scaling-policy-for-a-multi-cluster-warehouse

Domain

Performance Concepts

Question 71 **Skipped**

What is the minimum Snowflake edition that supports materialized views?

Business Critical

Correct answer

Enterprise
Virtual Private Snowflake
Standard
Overall explanation
The minimum Snowflake edition that supports Materialized Views is Enterprise. All editions above the enterprise edition also support martialized views.
https://docs.snowflake.com/en/user-guide/intro-editions.html
Domain
Licensing & Features
Question 72 **Skipped**
Which of the following languages are supported for creating UDFs in Snowflake?
Correct selection
SQL
Correct selection
Java
Correct selection
JavaScript
C++
C#
Correct selection
Python
Overall explanation
SQL, Java, JavaScript, and Python can be used to create UDFs in Snowflake.
https://docs.snowflake.com/en/sql-reference/udf-overview#supported-languages
Domain
Extending Snowflake Functionality
Question 73 **Skipped**
Snowflake releases a behaviour change release at what frequency?
Once every week
Correct answer
Once a Month
Once every fortnight
Once a year
Overall explanation
Once a month, Snowflake also releases a "behavior change release." These changes to existing behaviors may affect customers who already use the service. Over two months, the new behavior is adopted by everyone. The behavior change is not enabled during the first month unless the customer opts in. The behavior modification is enabled automatically in the second month, but a customer can opt-out if desired.
https://docs.snowflake.com/en/user-guide/intro-releases
Domain
Account
Question 74 **Skipped**
True or False: You can increase or decrease the size of a virtual warehouse as required.
False
Correct answer
True
Overall explanation

A virtual warehouse can be scaled up to ensure satisfactory performance when the complexity of the queries has increased. Scaling up allows the virtual warehouse to expand in size to keep up with the increasing complexity of the tasks. When a virtual warehouse is scaled up, the number of nodes in the compute cluster increases. Scaling down a virtual warehouse is generally done in response to decreased query complexity, where a smaller virtual warehouse may be sufficient to meet the performance requirements. When a virtual warehouse is scaled down, nodes are removed from the virtual warehouse. https://docs.snowflake.com/en/user-guide/warehouses-considerations

Domain

Performance Concepts

Question 75 Skipped

What role is required to enable replication for multiple accounts?

Correct answer

ORGADMIN

ACCOUNTADMIN

SECURITYADMIN

SYSADMIN

Overall explanation

The ORGADMIN must be used to enable replication for source and target databases.

https://docs.snowflake.com/en/user-guide/database-replication-config#prerequisite-enable-replication-for-accounts-in-the-organization

Domain

Security

Question 76 Skipped

Which of the following Snowflake Editions support private connectivity to Snowflake's internal stages?

Correct selection

Virtual Private Snowflake (VPS) edition

Standard Edition

Enterprise Edition

Correct selection

Business Critical Edition

Overall explanation

Private connectivity enables you to ensure that access to your Snowflake instance is via a secure connection and, potentially, to block internet-based access completely. Private connectivity to Snowflake requires the Business-Critical edition as a minimum.

Domain

Licensing & Features

Question 77 Skipped

Which of the following chart types are supported by Snowsight?

Select all that apply.

Correct selection
Scorecards
Pareto Charts
Pie Charts
Correct selection
Bar Charts
Area charts
Overall explanation
Snowsight supports

Bar charts,

Line charts,

Scatterplots,

Heat grids and

Scorecards

https://docs.snowflake.com/en/user-guide/ui-snowsight-visualizations

Domain
Tools & Interfaces
Question 78 **Skipped**
True or False: A share must have at least one consumer added.
True
Correct answer
False
Overall explanation
A Snowflake share can be defined without a consumer added to it. One or more consumers can be added to the Share afterward
Domain
Data Sharing
Question 79 **Skipped**
True or False: The COPY command can unload data using a SELECT query.
Correct answer
True
False
Overall explanation
The COPY command allows unloading or exporting data from a table or a view and also allows using queries (SELECT) to unload data. https://docs.snowflake.com/en/user-guide/data-unload-overview#bulk-unloading-using-queries
Domain
Data Loading and Unloading
Question 80 **Skipped**

Snowflake stores which of the following metadata about data in a micro-partition. Select all that apply.

Correct selection

The range of values for each of the columns in the micro-partition.

Correct selection

Additional properties for optimization and efficient processing.

Correct selection

The number of distinct values.

Overall explanation

All of these are valid examples of the metadata that Snowflake stores for micro-partition. Snowflake stores the range of column values in its metadata: the maximum and the minimum value for each column in each micro-partition. Snowflake can intelligently decide which partitions to read when processing a query using this metadata. Additionally, Snowflake stores the count of distinct values for each column in each partition in the metadata and certain other information to assist in query optimization. https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions

Domain

Architecture

Question 81 **Skipped**

For a multi-cluster virtual warehouse, what is the maximum number of clusters?

1

Correct answer

10

50

24

Overall explanation

A multi-cluster virtual warehouse supports anywhere from one and ten different clusters simultaneously. The minimum number of clusters supported is one, and the maximum number of allowed clusters is ten.

https://docs.snowflake.com/en/user-guide/warehouses-multicluster#what-is-a-multi-cluster-warehouse

Domain

Performance Concepts

Question 82 **Skipped**

Which of the following correctly describes a materialized view? Select all that apply.

Correct selection

A materialized view is automatically updated if the data in the underlying table is changed.

Correct selection

A Snowflake-managed service keeps a materialized view in sync with the base table.

A materialized view must be manually updated if the underlying table's data is changed.

An active virtual warehouse is required to sync a materialized view with its base table.

Overall explanation

A materialized view is a view that pre-computes data based on a SELECT query. The query's results are pre-computed and physically stored to enhance performance for similar queries that are executed in the future. When the underlying table is updated, the materialized view refreshes automatically, requiring no additional maintenance. Snowflake-managed services perform the update in the background transparent to the user without interfering with the user's experience.
https://docs.snowflake.com/en/user-guide/views-materialized

Domain

Performance Concepts

Question 83 Skipped

Snowflake allows which of the following access control methods? Select all that apply.

Correct selection

Discretionary access control (DAC)

Correct selection

Role-based access control (RBAC)

Management access control (MAC)

Global access control (GAC)

Overall explanation

Snowflake's access control system is built on the RBAC idea, which means that privileges are issued to roles and roles to users. The privileges associated with a role are given to all users assigned to it. Snowflake also supports discretionary access control (DAC), which means that the role that created an object owns it and can provide access to other roles to that item. https://docs.snowflake.com/en/user-guide/security-access-control-overview

Domain

Security

Question 84 Skipped

Which of the following statements regarding multi-cluster virtual warehouses are true? Select all that apply.

Correct selection

Multi-cluster virtual warehouses can be set to auto-suspend or auto-resume, just like a standard virtual warehouse

Correct selection

Multi-cluster virtual warehouses automatically add or remove virtual warehouses in response to changing workload demands.

You must manually add or remove clusters in a multi-cluster virtual warehouse.

Multi-cluster virtual warehouses cannot be set to auto-suspend or auto-resume.

Overall explanation

When the concurrent workload for a virtual warehouse reaches the maximum, new queries are queued. Multi-cluster virtual warehouses address this by adding clusters as needed. When demand drops, the extra clusters are removed. Multi-cluster virtual warehouses can be suspended or resumed or set to auto-suspend/auto-resume, just like any other virtual warehouse. https://docs.snowflake.com/en/user-guide/warehouses-multicluster

Domain

Architecture

Question 85 Skipped
What is the minimum Snowflake edition that supports private connectivity to Snowflake?
Standard
Correct answer
Business Critical
Virtual Private Snowflake
Enterprise
Overall explanation
Private connectivity enables you to ensure that access to your Snowflake instance is via a secure connection and, potentially, to block internet-based access completely. Private connectivity to Snowflake requires the Business-Critical edition as a minimum.
Domain
Licensing & Features
Question 86 Skipped
Which of the following are examples of virtual warehouse sizes?
Low
High
Correct selection
Medium
Correct selection
6X-Large
Correct selection
X-Small
Overall explanation
Snowflake has made it easy and quick for a user to choose a virtual warehouse by labeling the configuration in T-shirt sizes: X-Small, Small, Medium, Large, and so on https://docs.snowflake.com/en/user-guide/warehouses-overview
Domain
Architecture
Question 87 Skipped
True or False: After a table has been cloned, dropping the original table will also drop the clone.
Correct answer
False
True
Overall explanation
Micro-partitions and metadata in the cloud services layer enable rapid and efficient zero-copy cloning because the cloned table's metadata references the existing micro-partitions. The source and cloned items are independent; thus, modifying data in one will not affect the other. For example, the source table can be dropped altogether, which doesn't affect the cloned table. https://docs.snowflake.com/en/user-guide/tables-storage-considerations#label-cloning-tables
Domain
Cloning
Question 88 Skipped
A share can only be created by which of the following? Select all that apply.
Correct selection
A role with CREATE SHARE privileges.

Correct selection

ACCOUNTADMIN role

SYSADMIN role

SECURITYADMIN role

Overall explanation

A share can be created only by the ACCOUNTADMIN role or roles that have been explicitly granted CREATE SHARE privilege. https://docs.snowflake.com/en/user-guide/data-sharing-gs

Domain

Data Sharing

Question 89 **Skipped**

Which of the following is true when a virtual warehouse is scaled up to a larger size? Select all that apply.

The virtual warehouse cannot be scaled down to a smaller size.

Correct selection

The charging for the new size is not started until all new nodes in the larger virtual warehouse are provisioned.

All existing queries are terminated and must be re-submitted.

Correct selection

The increased size does not affect any queries already executing on the virtual warehouse.

Correct selection

Only new queries benefit from the larger virtual warehouse size.

Overall explanation

When a virtual warehouse is scaled up, the charging for the new size does not begin until all the new nodes in the larger virtual warehouse have been provisioned. Only new queries are affected by the changed size; existing queries on the virtual warehouse remain unaffected. https://docs.snowflake.com/en/user-guide/warehouses-

Domain

Performance Concepts

Question 90 **Skipped**

What is the minimum Snowflake edition required to share data with other Snowflake accounts?

Correct answer

Standard

Business Critical

Virtual Private Snowflake

Enterprise

Overall explanation

Data sharing is supported in all Snowflake editions; thus, the minimum edition that supports it is the Standard edition. https://docs.snowflake.com/en/user-guide/intro-editions.html

Domain

Licensing & Features

Question 91 **Skipped**

Which method can you use to retrieve the history of data loaded into tables through Snowpipe and the COPY INTO command?

Query the PIPE_USAGE_HISTORY view in the ACCOUNT_USAGE schema

Query the LOAD_HISTORY view in the ACCOUNT_USAGE schema

Use QUERY_HISTORY table function in INFORMATION_SCHEMA

Correct answer

Query the COPY_HISTORY view in the ACCOUNT_USAGE schema

Overall explanation

The COPY_HISTORY view in the ACCOUNT_USAGE schema can be used to view history for data loaded through either the COPY command or continuous data loaded through Snowpipe. The COPY_HISTORY view shows the history for the last 365 days.

The LOAD_HISTORY view shows data only for the COPY command. The PIPE_USAGE_HISTORY view shows only the Snowpipe history.

https://docs.snowflake.com/en/sql-reference/account-usage/copy_history

Domain

Account Usage & Monitoring

Question 92 Skipped

Which of the following sentences accurately describes scaling out in Snowflake? Choose all that apply.

Scaling out is accomplished by increasing or decreasing the size of a virtual warehouse.

Correct selection

Scaling out is accomplished through the usage of multi-cluster virtual warehouses.

Correct selection

Scaling out can assist in reducing query queuing.

Overall explanation

Typically, a virtual warehouse has a defined set of computing resources that it can use to execute queries. When queries are sent to a warehouse, the warehouse allocates the resources required for each query and begins running the queries. If there aren't enough resources to run all the queries sent to the warehouse, Snowflake queues the extra queries until the resources are available again. Snowflake provides multi-cluster virtual warehouses to overcome this issue. Multi-cluster virtual warehouses are frequently used in scenarios where the number of concurrent queries exceeds the capacity of a single virtual warehouse. When a virtual warehouse's concurrent workload exceeds its maximum capacity, additional queries are placed in the queue. Multi-cluster virtual warehouses dynamically add additional clusters based on demand to solve the queueing issue. When demand decreases, the additional clusters are decommissioned. This process is also known as scaling out or auto-scaling.

Domain

Performance Concepts

Question 93 Skipped

Which of the following can retrieve data from fail-safe storage?

Correct answer

Snowflake support

Customer

Anyone
Cloud provider
Overall explanation
Once the data is in fail-safe storage, only Snowflake support can help retrieve the data. The customer cannot access fail-safe storage. The cloud provider cannot access any of the data stored by Snowflake, whether in fail-safe storage or otherwise.
https://docs.snowflake.com/en/user-guide/data-failsafe
Domain
Fail-safe
Question 94Skipped
Which of the following is true regarding Pre-signed URLs?

Select all that apply.

A pre-signed URL expires after 24 hours.
Correct selection
The expiry time for a pre-signed URL can be configured.
Correct selection
Anyone with a pre-signed URL can use the URL to access the referenced file.
Correct selection
Pre-signed URLs are suitable to provide access to users & applications without needing authentication or authorization.
Correct selection
A pre-signed URL is generated with an access token and can be accessed without requiring authentication.
Overall explanation
A pre-signed URL is a simple HTTPS URL for accessing a file using a web browser. A pre-signed URL is generated using a pre-signed access token. Users can temporarily access a file via a pre-signed URL without authorization. The expiry duration of a pre-signed URL is configurable and can be set to the required duration.

https://docs.snowflake.com/en/user-guide/unstructured-intro#types-of-urls-available-to-access-files

Domain
Data Transformation
Question 95Skipped
Consider the CUSTOMER table in the SNOWFLAKE_SAMPLE_DATA.TPCH_SF1 schema. Your virtual warehouse is in a suspended state but is set to auto-resume. Which of the following queries will result in the virtual warehouse being resumed? Select all that apply.
USE SNOWFLAKE_SAMPLE_DATA.TPCH_SF1;

SHOW TABLES LIKE '%CUSTOMER%';

SELECT COUNT(*) FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.CUSTOMER;
Correct selection

SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.CUSTOMER;
DESCRIBE TABLE SNOWFLAKE_SAMPLE_DATA.TPCH_SF1;
Correct selection
SELECT C_MKTSEGMENT,SUM(C_ACCTBAL) FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.CUSTOMER GROUP BY C_MKTSEGMENT;
Overall explanation
Metadata cache or cloud services operations do not require an active virtual warehouse. Other queries will need an active virtual warehouse.

Statistics are kept in the metadata cache in the cloud services layer for each table, micro-partition, and column. The metadata cache can return results if the query simply counts the number of rows.

Similarly, the cloud services layer can provide table definitions (i.e., DESCRIBE) and a list of tables in a schema (i.e., SHOW TABLES LIKE).

Domain
Architecture
Question 96Skipped
Which of the following statements is true regarding the PUBLIC role? Select all that apply.
Correct selection
The PUBLIC role is the least privileged role in a Snowflake system.
The PUBLIC role is not pre-defined and must be created by an account administrator.
The PUBLIC role is the most privileged role in a Snowflake system.
Correct selection
The PUBLIC role is automatically assigned to every user in Snowflake.
Overall explanation
The PUBLIC role is one of the out-of-the-box roles in Snowflake. The PUBLIC role has the fewest privileges and is assigned automatically to all users.
https://docs.snowflake.com/en/user-guide/security-access-control-overview#system-defined-roles.
Domain
Security
Question 97Skipped
You are required to implement column-level security in Snowflake. Which techniques can you use? Select two.
Correct selection
External Tokenization
Object Security
Row-level security
Correct selection
Dynamic Data Masking
Overall explanation

Snowflake supports masking policies that may be applied to columns and enforced at the column level to provide column-level security. Column-level security is achieved by dynamic data masking or external Tokenization. https://docs.snowflake.com/en/user-guide/security-column

Domain

Security

Question 98 Skipped

Which of the following statement describe micro-partitions correctly?

Correct selection

Micro-partitions are immutable.

Data in micro-partitions is organized in a row storage format.

Micro-partitions are mutable and can be updated.

Correct selection

Data in micro-partitions is organized in a columnar format.

Overall explanation

Snowflake partitions are immutable, which means they cannot be changed once created. Table data is mapped to individual micro-partitions and is further organized using a columnar format. https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions.html

Domain

Architecture

Question 99 Skipped

Consider the following snippet from the query profile of a finished query.

Which of the following accurately describes the highlighted statistics?

The query profile indicates ineffective partition pruning.

Correct selection

The query profile indicates effective partition pruning.

The query profile indicates that the virtual warehouse cache was used.

Correct selection

The query profile indicates that the virtual warehouse used is too small for the query.

Overall explanation

Partition pruning occurs when the number of Partitions scanned is much smaller than Partitions total. Snowflake saves data on the warehouse's local disk if it can't fit an operation into memory. Data spilling slows down queries because it requires more IO operations, and disk access is slower than memory access
https://docs.snowflake.com/en/user-guide/ui-query-profile

Domain

Performance Concepts

Question 100 **Skipped**

True/False: Once the Time Travel retention period has ended for a temporary table, Historical data in Temporary tables can not be recovered by Snowflake support

False

Correct answer

True

Overall explanation

Transient and temporary tables don't have fail-safe functionality; therefore, data in such tables goes through zero days of fail-safe storage. Also, Transient and Temporary tables have a maximum of 1 day of Time Travel. Therefore, once the Time Travel period for these tables is complete, there is no way to recover historical data.
https://docs.snowflake.com/en/user-guide/tables-temp-transient

Domain

Data Protection

Question 101 **Skipped**

Which of the following statement regarding Tasks in Snowflake is correct? Select all that apply.

A task can be created without any specific privileges.

Explanation

Creating a task in Snowflake requires the user to have the "CREATE TASK" privilege on the schema where the task will reside.

Correct selection

A task can be manually executed using an EXECUTE TASK command.

Explanation

Using the EXECUTE TASK command it is possible to manually trigger an asynchronous single run of a job, which can be either a standalone task or the root task in a task graph. This run is independent of the scheduled executions defined for the task.
For more information, see the following link: https://docs.snowflake.com/en/sql-reference/sql/execute-task

A single task permits a user to run multiple SQL statements on a predefined schedule.

Explanation

Each task in Snowflake is limited to executing a single SQL command or statement. It cannot execute multiple SQL statements within the same task definition.

Correct selection

A task enables users to run a single SQL command or statement on a defined schedule.

Explanation

Snowflake tasks are designed to execute single SQL commands or statements on a predefined schedule. Tasks are used to automate the execution of individual SQL commands or statements at specified intervals. They can also use procedural logic using stored procedures or Snowflake Scripting.
For more information, see the following link:
https://docs.snowflake.com/en/sql-reference/sql/create-task#required-parameters

Domain

Tasks

Question 102 Skipped

You are the owner of a table called CUSTOMER. The table is in the PUBLIC schema in the database called MARKETING. Which permissions do you need to provide on which objects to grant SELECT access to a role called ANALYST on the CUSTOMER table?

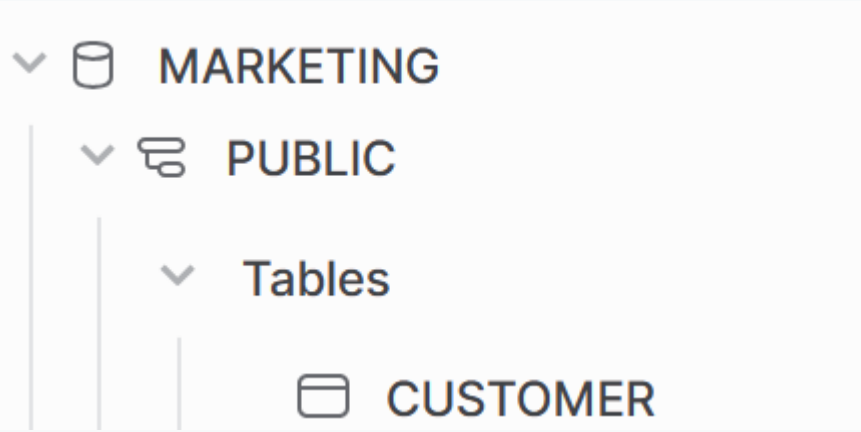

Select all that apply.

grant READ on schema PUBLIC to role ANALYST;

Correct selection

grant USAGE on schema PUBLIC to role ANALYST;

grant SELECT on schema PUBLIC to role ANALYST;

Correct selection

grant USAGE on database MARKETING to role ANALYST;

grant SELECT on database MARKETING to role ANALYST;

Correct selection

grant SELECT on table CUSTOMER to role ANALYST;

Overall explanation

SELECT privilege on the table is required for the ANALYST role to read data from the table. However, users (in the ANALYST role) cannot run their queries if the ANALYST role doesn't have the privileges to use the schema and the database containing the table. The USAGE privilege allows users to run the USE DATABASE or USE SCHEMA command or access the table using the full namespace i.e. MARKETING.PUBLIC.CUSTOMER.

Therefore, USAGE privileges on both database and the schema are required.

Domain
Security

Question 103 Skipped
Which of the following are Time Travel SQL extensions?

Correct selection
UNDROP

Correct selection
BEFORE

TIMETRAVEL

TIME

Correct selection
AT

Overall explanation
To support Time Travel queries, Snowflake supports special SQL extensions. It supports the AT and BEFORE statements which can be used with SELECT statements or while cloning tables, schemas, and databases. Snowflake also supports the UNDROP statement, which can be used to recover tables, schemas, or even complete databases after they have been dropped. https://docs.snowflake.com/en/user-guide/data-time-travel#time-travel-sql-extensions

Domain
Time Travel

Question 104 Skipped
How are privileges inherited in a Snowflake role hierarchy?

Only the direct child role inherits privileges.

Only roles at the same level of the hierarchy inherit the privileges.

Only the direct parent role inherits privileges.

Correct answer
The privileges of a role are inherited by all roles above it in the hierarchy.

Overall explanation
Roles may also be granted to other roles, creating a role hierarchy. The privileges associated with a role are inherited by all roles in the hierarchy above that role.

https://docs.snowflake.com/en/user-guide/security-access-control-overview#:~:text=Roles%20can%20be%20also%20granted%20to%20other%20roles%2C%20creating%20a%20hierarchy%20of%20roles.%20The%20privileges%20associated%20with%20a%20role%20are%20inherited%20by%20any%20roles%20above%20that%20role%20in%20the%20hierarchy.

Domain
Security
Question 105 Skipped
Query Result Cache reuse can be turned off using which parameter?
ENABLE_QUERY_RESULT_CACHE
Correct answer
USE_CACHED_RESULT
PURGE_CACHED_RESULTS
DISABLE_QUERY_RESULT_CACHE
Overall explanation
Query result cache is enabled by default but can be turned off at a session, user, or account level using the USE_CACHED_RESULT parameter.
https://docs.snowflake.com/en/user-guide/querying-persisted-results
Domain
Performance Concepts
Question 106 Skipped
Which of the following scenario requires you to have replication configured to recover?

Select all that apply.

A new data pipeline rolled out in production yesterday has deleted all rows from a production table.
Correct selection
The cloud region hosting your primary Snowflake instance has gone down due to a catastrophic failure, and data is unavailable.
An administrator accidentally dropped a production table last week.
Correct selection
The cloud platform hosting your primary Snowflake instance failed and has become unavailable.
A data corruption issue that corrupted three production tables 93 days ago was discovered.
Overall explanation
If a cloud provider or a region goes down, Snowflake users may be affected. To ensure the least impact, you must be ready for cloud provider outages to keep Snowflake available to your users.

Snowflake account-level replication & database replication synchronizes critical account objects and data from the primary account to one or more secondary accounts in a different region or cloud platform. Database replication allows read-only copies of databases from a primary Snowflake account to a new region or cloud provider. In the event of a failure on the primary site, switch your workloads from the primary to one of the secondary locations.

https://docs.snowflake.com/en/user-guide/account-replication-intro

Domain
Data Protection
Question 107 Skipped
Which of the following can be cloned?
Internal Named Stages
Correct selection
Schemas
Correct selection
Databases
Correct selection
Tables
Overall explanation
Named Internal Stages cannot be cloned. When a database or schema is cloned, any Snowpipe that points to a Named Internal Stage is not cloned. Named External Stages can be cloned. Since a table stage is associated with a table, it is automatically cloned when the table is cloned. Additionally, external tables cannot be cloned either. Databases, Schema, Tables, etc., can be cloned. https://docs.snowflake.com/en/user-guide/object-clone#cloning-and-stages
Domain
Cloning
Question 108 Skipped
Your organization's security policies require that certain rows in tables are unavailable to users for querying. Which of the following Snowflake features can you use to meet these requirements? Select two.
Correct selection
Row Access Policies
Column Level Security
External Views
Correct selection
Secure Views
Overall explanation
Secure views can be used to return only certain rows from a table. Additionally, secure views hide the underlying data by removing some of the internal Snowflake optimizations. Alternatively, Row-level security (RLS) can be used to return only certain rows. RLS is implemented by creating row access policies, which include conditions and functions that govern which rows are returned during query execution.
https://docs.snowflake.com/en/user-guide/views-secure
https://docs.snowflake.com/en/user-guide/security-row-intro
Domain
Security
Question 109 Skipped
Why is Snowflake considered a SaaS (Software-as-a-Service) product?

Select all that apply.

Correct selection
The customer is not required to procure, install and manage any hardware.
Correct selection

Snowflake regularly updates the software, and all accounts receive these updates automatically, eliminating the need for manual installations, maintenance, and patches.

Correct selection

It provides Pay as you Go licensing, allowing users to pay only for the resources and features they use.

Correct selection

Snowflake runs in the cloud and is available over the Internet.

Overall explanation

All of these are characteristics of a Software-as-a-Service product.

Domain

Licensing & Features

Question 110 Skipped

Which of the following Snowflake edition doesn't support data sharing?

Standard

Enterprise

Business Critical

Correct answer

Virtual Private Snowflake (VPS)

Overall explanation

Virtual Private Snowflake (VPS) cannot use secure data sharing, Marketplace, etc., because VPS accounts have isolated metadata, compute, and storage and therefore don't have sharing capabilities.

Domain

Data Sharing

Question 111 Skipped

What is the maximum allowed duration for Time Travel in the Snowflake Enterprise edition?

45 days

0 days

1 day

Correct answer

90 days

Overall explanation

Depending on the Snowflake edition, the Time Travel duration might range from 1 to 90 days. The Standard edition allows for one day of Time Travel. Time Travel is possible for up to 90 days in the Enterprise version and above.

https://docs.snowflake.com/en/user-guide/data-time-travel#data-retention-period

Domain

Time Travel

Question 112 Skipped

Which of the following contributes toward the storage costs in a Snowflake system?

External tables

Correct selection

Temporary Tables

Correct selection

Transient tables

Correct selection

Permanent tables

Overall explanation
External tables do not contribute towards storage costs because the data for External tables are stored outside of Snowflake on cloud storage. However, permanent, temporary, and transient tables contribute to varying storage costs. https://docs.snowflake.com/en/user-guide/cost-understanding-data-storage#temporary-and-transient-tables-costs

Domain
Cost & Pricing

Question 113 Skipped
In Snowflake, which of the following can be used to load continuously arriving data?

SnowStorm
SnowTrickle
SnowFast
Correct answer
Snowpipe

Overall explanation
Snowflake allows continuous data loading using Snowpipe, a serverless service. Snowpipe enables you to load data in a micro-batch manner, loading small volumes of data on each execution. The micro-batch-based data loading is used when a continuous stream of data, such as transactions or events, must be loaded and made available to enterprises quickly. Snowpipe enables continuous data loading and can load data within a few minutes after it arrives in a stage. Snowpipe is serverless and has its own computational capability; therefore, it does not rely on virtual warehouses for processing. Snowflake automatically manages the compute required by a Snowpipe. Snowflake also manages the scaling up and down of a Snowpipe as per the data load requirement. Since a Snowpipe is serverless, its costs are charged separately from virtual warehousing fees. https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro

Domain
Data Loading and Unloading

Question 114 Skipped
When cloning a database, your current role must have which privilege (as a minimum) on the source database?

WRITE
SELECT
Correct answer
USAGE

Overall explanation
To clone a table, you need SELECT privileges on the source table. For cloning Pipes, Streams & Tasks, you require OWNERSHIP privileges; for all other objects that can be cloned, you need the USAGE privilege. https://docs.snowflake.com/en/sql-reference/sql/create-clone#general-usage-notes

Domain
Cloning

Question 115 Skipped
SQL Workbench, DBeaver & Erwin are what type of partners in the Snowflake partner ecosystem?

Security, Governance & Observability
Data Integration

Machine Learning & Data Science

Correct answer

SQL Development and Management

Overall explanation

All of these are SQL Development and Management partners of Snowflake. Please see https://docs.Snowflake.com/en/user-guide/ecosystem.html

Domain

Partners

Question 116 Skipped

Which of the following keys are combined in Tri-Secret Secure encryption? Choose two.

Correct selection

Customer-managed

Public key

Hash key

Correct selection

Snowflake-managed

Overall explanation

Tri-Secret Secure refers to the combination of a Snowflake-managed key and a customer-managed key, which results in the creation of a composite master key to protect your data. Tri-Secret Secure requires the Business Critical edition as a minimum and can be activated by contacting Snowflake support.
https://docs.snowflake.com/en/user-guide/security-encryption-manage

Domain

Security

Question 117 Skipped

Which of the following is true of Internal Stages in Snowflake?

Select two answers.

Files are automatically uploaded to a default named internal stage.

Named internal stages are automatically created for new files.

Correct selection

Tables are automatically assigned an internal stage.

Correct selection

Users are automatically assigned an internal stage.

Overall explanation

Table stages are internal stages automatically created for each table and can be used to load data into that table.

Each user also automatically gets an internal stage object, created as soon as a user is created.

https://docs.snowflake.com/en/user-guide/data-load-local-file-system-create-stage#types-of-internal-stages

Domain
Data Loading and Unloading

Question 118 Skipped
Which of the following is a way to improve query performance in Snowflake?

Join Indices

Correct answer

Clustering Keys

Secondary Indices

Query Hints

Overall explanation
Clustering a table on a specific column can optimize queries by eliminating unnecessary partitions from the query processing. A table can be re-clustered by defining a clustering key, which effectively redistributes the data into micro-partitions according to the clustering key, ensuring optimal access to queries that predicate or join on the clustered column https://docs.snowflake.com/en/user-guide/tables-clustering-keys

Domain
Performance Concepts

Question 119 Skipped
How does defining a clustering key help with improving query performance?

Snowflake distributes data on different compute clusters when clustering keys are defined

Defining clustering keys results in Snowflake pre-computing query results.

Correct answer

Optimal partition pruning occurs if the queries use predicates on columns that are part of the clustering key

Overall explanation
Clustering a table on a specific column can optimize queries by eliminating unnecessary partitions from the query processing. A table can be re-clustered by defining a clustering key, which effectively redistributes the data into micro-partitions according to the clustering key, ensuring optimal access to queries that predicate or join on the clustered column https://docs.snowflake.com/en/user-guide/tables-clustering-keys

Domain
Performance Concepts

Question 120 Skipped
Which statements are true regarding Worksheets in Snowsight?

Correct selection

Each worksheet acts as a different session.

Correct selection

You can set a different context (role, warehouse, schema, and database) for each worksheet.

All worksheets share the same session.

You can not change the context (role, warehouse, schema, and database) for any worksheet.

Overall explanation
In the worksheet view, you can choose the primary role under which the query is executed, and the virtual warehouse used to run the query. You can also choose the

database and schema to which the worksheet view defaults, so you don't need to prefix tables in the specified database & schema. Each worksheet is an independent session.
https://docs.snowflake.com/en/user-guide/ui-snowsight

Domain
Tools & Interfaces

Question 121 Skipped

John, who has the SYSADMIN role, ran a query. Another user, Jane, attempts to view the result of the query executed by John using the Snowsight query history.

Which of the following correctly describes the outcome?

Correct answer
Jane can NOT see the result set of queries executed by John.
Jane can see the result set of queries executed by John only if she has the ACCOUNTADMIN role.
Jane can only see the result set of queries executed by John if John grants her permission.
Jane can see the result set of queries executed by John only if she is also part of the SYSADMIN role.

Overall explanation
You can only see the results of historical queries that you have run. For privacy reasons, the Query Detail page doesn't show the query results for queries run by other users, even if you have the privilege to see the query details for those queries.

https://docs.snowflake.com/en/user-guide/ui-history#viewing-query-details-and-results

Domain
Security

Question 122 Skipped

Which columns will be part of the result set when a directory table is queried?

Select all that apply.

IS_ENCRYPTED
Correct selection
LAST_MODIFIED
Correct selection
MD5
Correct selection
ETAG
IS_DELETED

Overall explanation
When a directory table is queried, the result set contains the FILE_URL for each file in the stage object. The result set also contains additional metadata, such as the file's

relative path, which shows the file's path relative to the stage. The result set also has metadata such as the size of the file in bytes and the timestamp of when a file was last modified, the MD5 checksum for the file, and an ETAG file, which changes if the contents of the file change. When querying a directory table, you can filter the result set using the WHERE clause on any of these fields. For example, you can use the size column to limit your results to only those files that are greater than 10MB.

https://docs.snowflake.com/en/user-guide/data-load-dirtables-manage#output

Domain
Data Transformation

Question 123 Skipped

What are resource monitors used for?

Monitor the resource allocation for each virtual warehouse

Scale virtual warehouses up and down automatically

Monitor how many queries are queued

Correct answer

Control costs and credit use by virtual warehouses

Overall explanation

Resource monitors help manage virtual warehouse costs and avoid unexpected credit usage. Resource monitors can control credit usage by monitoring credit usage against a defined upper limit, notifying administrators when a certain percentage of the limit is reached, and even suspending virtual warehouses if necessary.

https://docs.snowflake.com/en/user-guide/resource-monitors

Domain
Account Usage & Monitoring

Question 124 Skipped

What is the minimum Snowflake edition that supports fail-safe?

Enterprise

Correct answer

Standard

Virtual Private Snowflake

Business Critical

Overall explanation

Fail-safe is supported in all Snowflake editions; therefore, the minimum edition with fail-safe support is the Standard edition. https://docs.snowflake.com/en/user-guide/data-failsafe

Domain
Fail-safe

v